# Pivotal



#### **Use of Open Source**

This product may be distributed with open source code, licensed to you in accordance with the applicable open source license. If you would like a copy of any such source code, EMC will provide a copy of the source code that is required to be made available in accordance with the applicable open source license. EMC may charge reasonable shipping and handling charges for such distribution. Please direct requests in writing to EMC Legal, 176 South St., Hopkinton, MA 01748, ATTN: Open Source Program Office.

#### Copyright © 2013 EMC Corporation. All rights reserved.

EMC believes the information in this publication is accurate as of its publication date. The information is subject to change without notice.

THE INFORMATION IN THIS PUBLICATION IS PROVIDED "AS IS." EMC CORPORATION MAKES NO REPRESENTATIONS OR WARRANTIES OF ANY KIND WITH RESPECT TO THE INFORMATION IN THIS PUBLICATION, AND SPECIFICALLY DISCLAIMS IMPLIED WARRANTIES OF MERCHANTABILITY OR FITNESS FOR A PARTICULAR PURPOSE.

Use, copying, and distribution of any EMC software described in this publication requires an applicable software license

For the most up-to-date listing of EMC product names, see EMC Corporation Trademarks on EMC.com

All other trademarks used herein are the property of their respective owners.

Date: 7/1/13

## **Table of Contents**

About Pive	otal, Inc	1
About this	Guide	1
	Conventions	
	onventions	
	and Syntax Conventions	
	•	
	Overview	
	mmand Center Overview	
	Command Center UI	
Pivota	HD Manager	4
Perfor	mance Monitor (nmon)	5
Postgr	eSQL Database	5
Architectu	ral Overview	6
Chanter 2:	Installing Pivotal Command Center 2.0.x	-
	Platforms	
	ownloads	
	tes	
•	ccessibility	
_	necks	
	n Instructions	
	g, Stopping, and Restarting Command Center Services	
	ning Command Center	
	teps	
	g Pivotal Command Center	
	Pivotal Command Center	
<b>Chapter 3:</b>	Using Pivotal Command Center UI	.12
Overview		.12
Logging I	1	.12
Browser S	upport	.13
Login	Screen	.13
Select	ng a Cluster	.13
Settin	ys Menu	.14
Dashboar	1	.14
Cluster Ar	alysis	.15
	e Job Monitor	
•	tails	
	Monitor	
	ery Monitor	
	· Pivotal Command Center Performance Monitor	
	Pivotai Command Center Performance Monitor	
Appendix A	Creating a YUM EPEL Repository	.23
• •	Best Practises -Troubleshooting	
	Issues Initor Page	
1ah M.		

## **Preface**

This preface includes four sections:

- About Pivotal, Inc.
- About this Guide
- Document Conventions

### **About Pivotal, Inc.**

Greenplum is currently transitioning to a new corporate identity (Pivotal, Inc.). We estimate that this transition will be completed in 2013. During this transition, there will be some legacy instances of our former corporate identity (Greenplum) appearing in our products and documentation. If you have any questions or concerns, please do not hesitate to contact us through our web site,

http://gopivotal.com/about-pivotal/support.

#### **About this Guide**

This guide includes the following topics:

- Chapter 1, "Overview"— An overview of the functionality and architecture of Pivotal Command Center.
- Chapter 2, "Installing Pivotal Command Center 2.0.x" Basic instructions for installing and configuring Pivotal Command Center to monitor a running instance of Pivotal HD.
- Chapter 3, "Using Pivotal Command Center UI"— An overview of the functionality of the Pivotal Command Center console (user interface).
- Chapter 4, "Pivotal Command Center Performance Monitor"— An overview of the Pivotal Command Center performance monitor known as nmon.
- Appendix A, "Creating a YUM EPEL Repository"— Instructions how to setup a local yum repository or point your hosts to an EPEL repository.
- Appendix B, "Best Practises -Troubleshooting".

## **Document Conventions**

The following conventions are used throughout the Pivotal Command Center documentation to help you identify certain types of information.

- Text Conventions
- Command Syntax Conventions

About Pivotal, Inc.

### **Text Conventions**

Table 0.1 Text Conventions

<b>Text Convention</b>	Usage	Examples
italics	New terms where they are defined Database objects, such as schema, table, or columns names	The <i>master instance</i> is the postgres process that accepts client connections.  Catalog information for Pivotal Command Center resides in the <i>pg_catalog</i> schema.
monospace	File names and path names Programs and executables Command names and syntax Parameter names	Edit the postgresql.conf file.  Use gpstart to start Pivotal Command Center.
<monospace italics&gt;</monospace 	Variable information within file paths and file names  Variable information within command syntax	/home/gpadmin/ <config_file> COPY tablename FROM '<filename>'</filename></config_file>
monospace bold	Used to call attention to a particular part of a command, parameter, or code snippet.	Change the host name, port, and database name in the JDBC connection URL:  jdbc:postgresql://host:5432/m ydb
UPPERCASE	Environment variables SQL commands Keyboard keys	Make sure that the Java /bin directory is in your \$PATH.  SELECT * FROM my_table;  Press CTRL+C to escape.

## **Command Syntax Conventions**

**Table 0.2** Command Syntax Conventions

Text Convention	Usage	Examples
{ }	Within command syntax, curly braces group related command options. Do not type the curly braces.	FROM { 'filename'   STDIN }
[ ]	Within command syntax, square brackets denote optional arguments. Do not type the brackets.	TRUNCATE [ TABLE ] name

**Document Conventions** 

 Table 0.2
 Command Syntax Conventions

Text Convention	Usage	Examples
	Within command syntax, an ellipsis denotes repetition of a command, variable, or option. Do not type the ellipsis.	DROP TABLE name [,]
	Within command syntax, the pipe symbol denotes an "OR" relationship. Do not type the pipe symbol.	VACUUM [ FULL   FREEZE ]

## 1. Overview

This document provides an overview of Pivotal Command Center 2.0.x and instructions for installing the product. This chapter provides an overview of Pivotal Command Center, then briefly describes each component.

- Pivotal Command Center Overview
  - Pivotal Command Center UI
  - Pivotal HD Manager
  - Performance Monitor (nmon)
  - PostgreSQL Database
- Architectural Overview

#### **Pivotal Command Center Overview**

The Pivotal Command Center allows an administrative user to administer and monitor one or more Pivotal HD clusters. The Command Center has command-line tools to deploy and configure Pivotal HD clusters, as well as, an intuitive graphical user interface that is designed to help the user view the status of the clusters and take appropriate action. This release of Command Center allows only administering and monitoring of Pivotal HD Enterprise 1.0.x clusters.

Pivotal Command Center 2.0.x is comprised of the following:

#### **Pivotal Command Center UI**

The Pivotal Command Center UI provides the user with a single web-based graphical interface to monitor and manage one or more Pivotal HD clusters. This web application is hosted on a Ruby-on-Rails application which presents the status and metrics of the clusters. This data comes from multiple sources. All of the Hadoop specific data comes from the Pivotal HD Manager component. The system metrics data is gathered by our Performance Monitor (nmon) component.

See Chapter 3, "Using Pivotal Command Center UI" for more information.

#### **Pivotal HD Manager**

Pivotal HD Manager provides complete life cycle management for Pivotal HD Clusters. It performs the following two main groups of functions:

- Cluster installation, configuration and uninstalls
- Cluster monitoring and management

These functions are served through a set of RESTful web services that run as a web application on an Apache-Tomcat server on the Command Center admin host. This is called <code>gphdmgr-webservices</code>. This web application stores its metadata and cluster configuration for Pivotal HD cluster nodes and services in the Pivotal Command

Center PostgreSQL database. It makes use of a Puppet Server to perform most of its HD cluster installation and configuration. It also has a polling service that retrieves Hadoop metrics from the cluster and stores them in the Command Center PostgreSQL Database at periodic intervals.

Pivotal HD Manager provides a command-line interface (CLI) for installation, configuration and uninstalls. This tool invokes the <code>gphdmgr-webservice</code> APIs to install and configure the various Pivotal HD services. The CLI also provides a way to start and stop the clusters. For how to use this CLI, please refer to the *Pivotal HD Enterprise 1.0 Installation and Administrator Guide*.

The Command Center UI also invokes the <code>gphdmgr-webservice</code> APIs to retrieve all Hadoop-specific cluster metrics and status information. This includes the Hadoop metrics that was previously retrieved by the polling service.

#### **Performance Monitor (nmon)**

Pivotal Command Center comes with a Performance Monitor called nmon (for node monitor). This makes use of a highly scalable message passing architecture to gather performance metrics from each node that Command Center monitors. This consists of a nmon master daemon that runs on the Command Center admin host and an nmon daemon that runs on all the cluster nodes that report system metric information to the nmon master. This includes metrics such as CPU, memory, disk I/O and network usage information

The nmon master on the admin host dumps the system metrics it receives from the nmon agents on the cluster nodes into a PostgreSQL DB. This is then queried by the Command Center UI application to display its cluster analysis graphs.

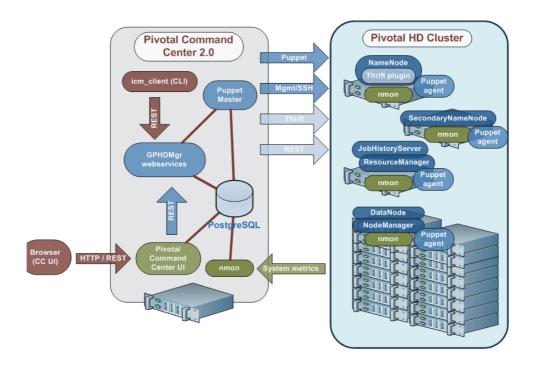
#### **PostgreSQL Database**

Pivotal Command Center makes use of a PostgreSQL Database to store the following:

- Cluster configurations
- Hadoop cluster metrics
- System metrics of the cluster
- Pivotal Command Center Metadata

## **Architectural Overview**

For more details about Pivotal HD Enterprise 1.0.x, refer to the *Pivotal HD 1.0 Installation and Administrator Guide*.



Architectural Overview 6

## 2. Installing Pivotal Command Center 2.0.x

This section describes how to install and configure Pivotal HD 2.0.x using the Pivotal Command Center Unified Installer.

This chapter includes the following sections:

- Supported Platforms
- Product Downloads
- Prerequisites
- Package Accessibility
- System Checks
- Installation Instructions
- Uninstalling Pivotal Command Center
- Upgrading Pivotal Command Center

## **Supported Platforms**

- RHEL 6.1 64-bit, 6.2 64-bit
- CentOS 6.1 64-bit, 6.2 64-bit

### **Product Downloads**

The following packages are required:

• PCC-2.0.x.\*.version build OS.x86 64.tar.gz

## **Prerequisites**

- Oracle JDK 1.6 installed on the Admin host.
- See Package Accessibility, below, for prerequisite packages.

Installation of Pivotal HD Manager assumes the user has a working knowledge of the following:

- Yum. Yum enables you to install or update software from the command line. See <a href="http://yum.baseurl.org/">http://yum.baseurl.org/</a>
- RPM (Redhat Package Manager).
- SSH (Secure Shell protocol).

Supported Platforms 7

## **Package Accessibility**

Pivotal Command Center and Pivotal HD Enterprise expect some prerequisite packages to be pre-installed on each host, depending on the software that gets deployed on a particular host. In order to have a smoother installation it is recommended that each host would have yum access to an EPEL yum repository. If you have access to the Internet, then you can configure your hosts to have access to the external EPEL repositories. However, if your hosts do not have Internet access (or you are deploying onto a large cluster), then having a local yum EPEL repo is highly recommended. This will also give you some control on the package versions you want deployed on your cluster. See Appendix A, "Creating a YUM EPEL Repository" for instructions on how to setup a local yum repository or point your hosts to an EPEL repository.

For Pivotal Command Center 2.0.x, here is a list of pre-requisites that need to either already be installed on the Command Center admin host or on an accessible yum repository:

- httpd
- mod ssl
- postgresql
- postgresql-devel
- postgresql-server
- compat-readline5
- createrepo
- sigar

You can run the following command on the admin node to make sure that you are able to install the prerequisite packages during installation.

```
$ sudo yum list httpd mod_ssl postgresql postgresql-devel
postgresql-server compat-readline5 createrepo sigar
```

If any of them are not available or not already installed, then you may have not added the repository correctly to your admin host.

For the cluster hosts (where you plan to install the cluster), the prerequisite packages depend on the software you will eventually install there, but you may want to verify that the following two packages are installed or accessible by yum on all hosts:

- nc
- postgresql-devel

## **System Checks**

**Important**: Avoid using hostnames with capital letters in them because Puppet has an issue generating certificates for domains with capital letters.

• Check that SELinux is disabled by running the following command:

```
# sestatus
```

Accepted return values are:

```
SELinuxstatus: disabled or:
SELinux status: permissive
```

If SELinux is enabled, you can temporarily disable it or make it permissive (this meets requirements for installation) by running the following command:

```
# echo 0 >/selinux/enforce
```

This only temporarily disables SELinux; once the host is rebooted, SE Linux will be re-enabled therefore we recommend disabling SELinux, described below, while running Pivotal HD/HAWQ.

**Note:** You can permanently disable SE Linux by editing the /etc/selinux/config file as follows, however disabling SELinux like this requires a system reboot for the change to take effect.

Change the value for the SELINUX parameter to:

```
SELINUX=disabled
```

Reboot the system.

- Every cluster node must be able to perform a forward and reverse DNS look-up for every other node.
- Verify that iptables is turned off, for example:

```
# chkconfig iptables off
# service iptables status
iptables: Firewall is not running.
```

#### **Installation Instructions**

Once you have met the prerequisites, you are ready to begin the installation. Perform the following installation steps as a root user.

**Upgrade Note:** If you are upgrading from Pivotal Command Center 2.0 to version 2.0.1, you must first stop the earlier version by running the following command:

```
$ service commander stop
```

**1.** Copy the Command Center tar file to your host. For example:

```
# scp ./PCC-2.0.x.version.build.os.x86_64.tar.gz
host:/root/phd/
```

System Checks 9

**2.** Log into the Command Center admin host as root user. cd to the directory where the Command Center tar files are located and untar. For example:

```
# cd /root/phd
# tar --no-same-owner -zxvf
PCC-2.0.x.version.build.os.x86 64.tar.gz
```

**3.** Still as root user, run the installation script. This installs the required packages and configures both Pivotal Command Center and Pivotal HD Manager, and starts services.

**Important**: You must run the installation script from the directory where it is installed, for example: PCC-2.0.x.version

For example:

```
# 1s
PCC-2.0.x.version
PCC-2.0.x.version.build.os.x86_64.tar.gz
# cd PCC-version
# ./install
```

You will see installation progress information on the screen. Once the installation successfully completes, you will see the following:

```
You have successfully installed PCC 2.0.x

You now need to install a GPHD cluster to monitor or sync

PCC to monitor an existing GPHD cluster. You can view your
cluster statuses here:

http://node0781.ic.analyticsworkbench.com:5000/status
```

**4.** Verify that your PCC instance is running by executing the following command:

```
$ service commander status
```

#### Starting, Stopping, and Restarting Command Center Services

To stop or restart Command Center services, run the following commands on the Pivotal Command Center admin host:

```
$ service commander stop
$ service commander start
$ service commander restart
```

#### **Launching Command Center**

Launch a browser and navigate to the host on which you installed Command Center. For example:

```
http://CommandCenterHost:5000
```

The Command Center login page is launched in your browser. The default username/password is <code>gpadmin/gpadmin</code>.

See Chapter 3, "Using Pivotal Command Center UI" for more details about the application, including how to change the default password.

Installation Instructions 10

#### **Next Steps**

See the *Pivotal HD 1.0 Enterprise Installation and AdministratorGuide* for instructions for using the command-line interface of Pivotal Command Center to deploy and configure a HD cluster.

## **Uninstalling Pivotal Command Center**

Follow the steps below to uninstall Pivotal Command Center and the Pivotal HD cluster:

- **1.** Stop services on all your clusters (See the *Pivotal HD Enterprise 1.0 Installation and Administrator Guide* for detailed steps).
- **2.** Uninstall all your clusters (See the *Pivotal HD Enterprise 1.0 Installation and Administrator Guide* for detailed steps).
- **3.** From the directory where you untarred the Pivotal Command Center, run the uninstall script:

```
# cd /root/phd/PCC-2.0.x.version/
# ./uninstall
```

## **Upgrading Pivotal Command Center**

Follow the steps below to upgrade your Pivotal CC to a newer version:

Note: We recommend backing up your data before performing any upgrades.

**1.** As gpadmin, stop the cluster:

```
# icm client -1 CLUSTERNAME stop
```

**2.** As root, stop the Command Center services:

```
# service commander stop
```

- **3.** Make sure that the tomcat server and puppet are not running any more (check for processes as well).
- **4.** Download the new PCC tarball and untar
- **5.** Run the installer from the new PCC location:

```
# ./install
```

Note: In the nmon configuration file (/etc/nmon/conf/nmon-site.xml), make sure the port number shown in <nmon-database-config> -> <uri> is 10432. If not update it to 10432 and restart the nmon service (service nmon restart).

#### **Important:**

- The upgrade replaces the Tomcat web server with Jetty.
- After this upgrade, the PostgreSQL database will run on port 10432 (pre-upgrade the default port was 5432)

## 3. Using Pivotal Command Center UI

This section provides an overview of the Pivotal Command Center 2.0 user interface.

### **Overview**

Pivotal Command Center UI is a browser-based application for viewing the status and performance of Pivotal HD clusters. At a high level, the screens consist of:

- Dashboard—Provides an overview of your Pivotal HD cluster. This screen shows at one glance the most important states and metrics that an administrator needs to know about the Pivotal HD cluster.
- Cluster Analysis—Provides detailed information about various metrics of your Pivotal HD cluster. This provides cluster-wide metrics all the way down to host-level metrics. This has hadoop-specific metrics such as MapReduce slot utilization and NameNode performance, as well as system metrics such as CPU, memory, disk and network statistics.
- MapReduce Job Monitor—Provides details about all, or a filtered set of MapReduce jobs.
- YARN App Monitor—Provides details about all, or a filtered set of YARN applications.
- HAWQ Query Monitor—When HAWQ (a revolutionary MPP database on Hadoop solution) is deployed on the cluster, Command Center can show the progress of all actively running queries on HAWQ.

#### **Status indicators**

Note that throughout the user interface the following indicators are used to indicate the status of nodes:

Green: Succeeded Blue: Running

Grey: Stopped/Pending

• Red: Killed/failed

## **Logging In**

The URL to access Pivotal Command Center UI from a browser is

http://CommandCenterHost:5000/login

To change the default port (5000), update the port settings in the following file:

/usr/local/greenplum-cc/config/app.yml

Overview 12

## **Browser Support**

The following browsers are supported by Pivotal Command Center 2.0:

- Firefox 16, 19
- IE 8, IE 9, both with Google Chrome Frame
- Chrome 25.0.1364.172

#### **Login Screen**

The first time you launch the Command Center UI, a login screen appears showing the hostname of the host for the Command Center.



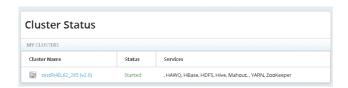
The default admin user/password is <code>gpadmin/gpadmin</code>. You can change this password via the Settings Menu.

Click the Login button to launch the Command Center UI.

#### Selecting a Cluster

Once you have launched Command Center, the Cluster Status screen appears, displaying a list of available clusters to monitor, the status of each cluster (**started**, **stopped**), and a list of services running on that cluster (Hive, mahout, and so on).

- Click the cluster name in the table to select a cluster.
- From any point within Command Center UI, you can always select a different cluster by using the **Select Cluster** drop-down menu in the upper right corner of the screen.



Browser Support 13

#### **Settings Menu**

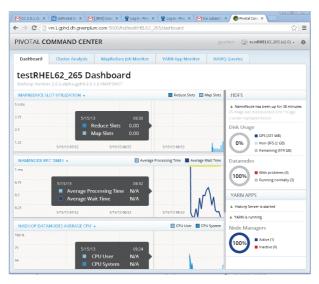
Click the **gear** icon in the upper right corner of the screen at any time to display the **Settings** menu. From the settings menu you can:

- Cluster Status. Click this to go back to a list of available clusters.
- Change Password. Click this to change your password.
- Logout.

### **Dashboard**

The dashboard gives you a high level view of a cluster at a glance. You are able to view the status of the most important cluster services, such as HDFS and YARN. It also shows you how the most important cluster metrics are trending in a visual way.

The graphs provide a unified view of the state of your system. They are also useful in detecting outliers and pinpointing specific problems that may be present in your system.



The right side of the Dashboard displays the state of both HDFS and YARN services. It answers the following questions:

- Is HDFS up?
- When did the last NameNode checkpoint occur?
- What percentage of cluster storage is used by HDFS and how much is free?
- How many DataNodes are up and are they running normally or with problems?
- Is YARN up?
- Is the History Serverup?

  Note: The History Server stores a history of the mapreduce jobs run on the cluster.
- How many NodeManagers are up?

The Dashboard provides metrics about:

Dashboard 14

- Mapreduce Slot Utilization
- Namenode RPC Times
- Hadoop Datanodes Average CPU
- Hadoop Datanodes Average Bandwidth
- Namenode Operations Per Second
- Hadoop Datanodes Average Disk Bandwidth
- Hadoop Datanodes Average Memory
- Mapreduce Jobs By Status

## **Cluster Analysis**

The Cluster Analysis screen provides detailed metrics on your Pivotal HD cluster.

It provides cluster-wide metrics all the way down to host-level metrics. It provides Hadoop-specific metrics, as well as system metrics that you can drill down to if needed.

The Cluster Analysis screen displays the same data that is shown in the dashboard but in greater detail.



By default the Cluster Analysis screen displays the metrics for all services, all categories, and all nodes. You can filter the information displayed by combinations of the following filters:

• By Service

Metrics can be filtered by services such as HDFS, YARN, or HAWQ.

• By Category

Metrics can be filtered by categories such as:

- namenode
- secondarynamenode
- datanode

Cluster Analysis 15

- yarn-resourcemanager
- yarn-nodemanager
- mapreduce-historyserver
- hawq-master
- hawq-segment

#### • Alphabetically

Metrics can be filtered alphabetically.

Based on the filters you select, the lower part of the Cluster Analysis screen provides detailed graphs that display data related to:

- Mapreduce Slot Utilization
- Namenode RPC Times
- Avg Namenode File Operations Per Second
- Mapreduce Jobs by Status
- Segment CPU
- Disk Bandwidth
- Network Bandwidth
- Segment Memory
- Load
- Swap Usage
- Swap I/O
- Network Operations
- Disk Operations

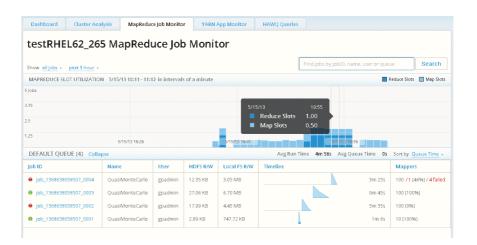
You can view either the **Performance Metrics**, which show the cluster/node utilization over-time, the **Real-time Metrics** which show the current metrics in real-time, or **Storage Metrics**, which show metrics about cluster storage.

If you select Cluster Analysis for **All Nodes** (the default), the Trending Metrics graph for the cluster is displayed:

Cluster Analysis 16

## **MapReduce Job Monitor**

The Job Monitor screen tracks the MapReduce jobs that are executed in the Pivotal HD cluster when the YARN MapReduce service is running. It provides details about all, or a filtered set of MapReduce jobs.



The MapReduce jobs displayed can be filtered by state and/or time range.

- By state:
  - all jobs (set by default)
  - currently pending jobs
  - currently running jobs
  - succeeded jobs
  - failed jobs
  - killed jobs
- By time range:

By selecting a preset time range in hours, weeks, months, year, or by specifying a custom time range.

The MapReduce jobs can also be filtered by searching for values for the following:

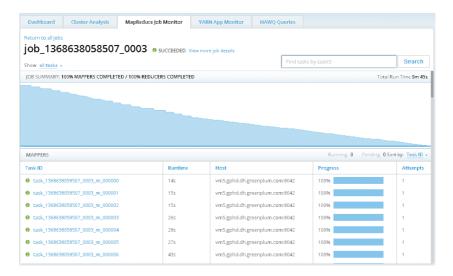
- jobID
- name
- user
- queue

Enter your search value in the search bar in the following format: searchKey=searchValue, where searchKey is one of **jobID**, **name**, **user**, or **queue**.

These are substring searches. For example: **jobID=1363920466130** will locate a job with **jobID=job\_1363920466130\_0002** 

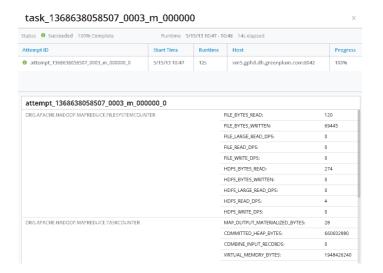
#### **Job Details**

When you click on any of the jobs in the Job Monitor more details of the job are shown.

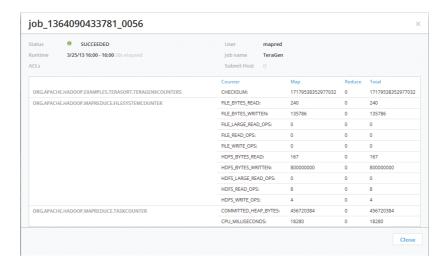


This screen displays all the tasks that are have been allocated for the selected job and their progress. You can see the mapper and the reducer tasks separately. In the above screen capture, the bars in the JOB SUMMARY section represent the two Mapper tasks that have run, one took 19 seconds, the other, 20 seconds.

Clicking on each task ID will show even more details about that particular task. You can also filter on a particular task ID in the search bar.

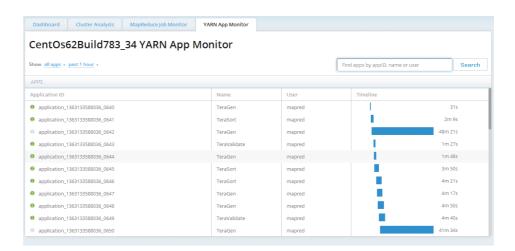


To see job related counters click on **View more job details** next to the job ID:



## **YARN App Monitor**

The YARN App Monitor screen tracks YARN applications that are executed in the Pivotal HD Cluster.



The YARN applications displayed can be filtered by category and/or time range:

- By Category:
  - all apps (set by default)
  - currently pending apps
  - currently running apps
  - succeeded apps
  - failed apps
  - killed apps
- By Time Range:

By selecting a preset time range in hours, weeks, months, year, or by specifying a custom time range.

The YARN applications can also be filtered by the following fields by entering it in the search bar in the following format: searchKey=searchValue:

- appID
- name
- user

These are substring searches. For example: appID=1363920466130 will locate the application with appID=application 1363920466130 0002

YARN App Monitor 20

## **HAWQ Query Monitor**

The HAWQ Query monitor is only displayed when HAWQ is installed on the cluster. This screen displays all **active** queries running on the HAWQ cluster:



In this release, this screen only displays **active** queries as can be seen when you run: SELECT \* FROM pg\_stat\_activity;

on the HAWQ cluster.

Click on a Query ID to get the syntax of that query:



HAWQ Query Monitor 21

## 4.

## **Pivotal Command Center Performance**

## **Monitor**

This section provides an overview of the Pivotal Command Center performance monitor known as nmon.

#### **Overview**

Pivotal Command Center comes with a Performance monitor called nmon (for node monitor). This makes use of a highly scalable message passing architecture to gather performance metrics from each node that Command Center monitors. This consists of a nmon master daemon that runs on the Command Center admin host and an nmon daemon that runs on all the cluster nodes that report system metric information to the nmon master. This includes metrics such as CPU, memory, disk I/O and network usage information.

The nmon master on the admin host dumps the system metrics it receives from the nmon agents on the cluster nodes into a PostgreSQL DB. This is then queried by the Command Center UI application to display its cluster analysis graphs.

The nmon agents hosts are deployed throughout the cluster during Pivotal HD cluster deployment itself (see *Pivotal HD Enterprise 1.0 Installation and Administrator Guide* for details).

The agents are deployed as services on each host, including on the Pivotal Command Center admin host. To stop or start the nmon service run the following as root:

```
# service nmon stop
```

# service nmon start

Overview 22

# A. Creating a YUM EPEL Repository

Pivotal Command Center and Pivotal HD Enterprise expect some prerequisite packages to be pre-installed on each host, depending on the software that gets deployed on a particular host. In order to have a smoother installation it is recommended that each host would have yum access to an EPEL yum repository. If you have access to the Internet, then you can configure your hosts to have access to the external EPEL repositories. However, if your hosts do not have Internet access (or you are deploying onto a large cluster), then having a local yum EPEL repo would be highly recommended. This will also give you some control on the package versions you want deployed on your cluster.

Following are the steps to create a local yum repo:

- 1. Mount the RHEL/CentOS DVD on a machine that will act as the local yum repo
- **2.** Install a webserver on that machine (e.g. httpd), making sure that HTTP traffic can reach this machine
- **3.** Install the following packages on the machine:

```
yum-utils
createrep
```

**4.** Go to the directory where the DVD is mounted and run the following command: # createrepo .

**5.** Create a repo file on each host with a descriptive filename in the /etc/yum.repos.d/ directory of each host (for example, CentOS-6.1.repo) with the following contents:

```
[CentOS-6.1]
name=CentOS 6.1 local repo for OS RPMS
baseurl=http://172.254.51.221/centos/$releasever/os/
$basearch/
enabled=1
gpgcheck=1
gpgkey=http://172.254.51.221/centos/$releasever/os/$basearch/
RPM-GPG-KEY-CentOS-6
```

**6.** Validate that you can access the local yum repos by running the following command:

Yum list

## Best Practises -Troubleshooting

#### nmon Issues

- If you have to restart the Admin node, ensure that the nmon service is started.
- If you notice any of the clusters are not being fully monitored, perform the following on the Admin node:
  - Make sure the nmon configuration (/etc/nmon/nmon-site.xml) includes all the clusters and their hosts. If it doesn't, update it and distribute the updated configuration to all the cluster hosts, then restart nmon on the Admin node as well as on the cluster hosts:

```
sudo service nmon restart
massh clusterHosts verbose 'sudo service nmon restart'
Where clusterHosts contains all the cluster hosts.
```

#### **Job Monitor Page**

• If an application is completed, but on the job monitoring page in the Command Center User Interface, it shows app/job as still running, then check whether History Server is running or not. If it is not running, start it.

Check using: http://<HistoryServerHost>:19888