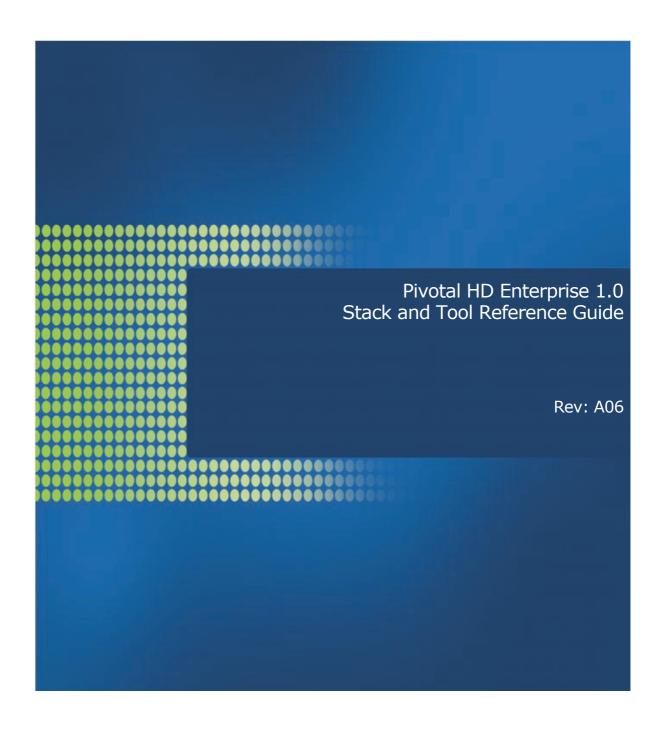
Pivotal



Use of Open Source

This product may be distributed with open source code, licensed to you in accordance with the applicable open source license. If you would like a copy of any such source code, EMC will provide a copy of the source code that is required to be made available in accordance with the applicable open source license. EMC may charge reasonable shipping and handling charges for such distribution. Please direct requests in writing to EMC Legal, 176 South St., Hopkinton, MA 01748, ATTN: Open Source Program Office.

Copyright © 2013 GoPivotal, Inc. All rights reserved.

GoPivotal, Inc. believes the information in this publication is accurate as of its publication date. The information is subject to change without notice. THE INFORMATION IN THIS PUBLICATION IS PROVIDED "AS IS." GOPIVOTAL, INC. ("Pivotal") MAKES NO REPRESENTATIONS OR WARRANTIES OF ANY KIND WITH RESPECT TO THE INFORMATION IN THIS PUBLICATION, AND SPECIFICALLY DISCLAIMS IMPLIED WARRANTIES OF MERCHANTABILITY OR FITNESS FOR A PARTICULAR PURPOSE. Use, copying, and distribution of any Pivotal software described in this publication requires an applicable software license.

All trademarks used herein are the property of Pivotal or their respective owners.

Table of Contents

About Pivotal, Inc	1
About This Guide	1
Getting Support	1
Product information and Technical Support	1
Chapter 1: MapReduce V1 (MR1) Based Clusters	
Prerequisities	2
Installation Notes	
Hadoop MR1	
Core Package Setup	
Hadoop-mr1 JobTracker Setup	
Hadoop-mr1 TaskTracker Setup	
Hadoop MR1 Configuration	
Usage	
Using Hadoop-mr1	
·	
Stop Hadoop-mr1	
Chapter 2: Yarn-Based Clusters	
Prerequisities	
Installation Notes	
Hadoop HDFS	7
Hadoop HDFS RPM Packages	8
Prerequisites: Core Package Setup	11
HDFS Namenode Setup	11
HDFS Datanode Setup	11
HDFS Secondary Namenode Setup	
Hadoop HDFS Configuration	12
Usage	12
Hadoop YARN	14
Hadoop YARN RPM Packages	14
Prerequisites: Core Package Setup	16
YARN ResourceManager Setup	17
YARN NodeManager Setup	17
Mapreduce HistoryServer Setup	17
YARN ProxyServer Setup	17
Hadoop HDFS Configuration	17
YARN Usage	18
Chapter 3: Additional Stack Components for Manual	
Installations	21
Prerequisities	
Installation Notes	
Zookeeper	
Zookeeper Server Setup	
Zookeeper Client Setup	
Zookeeper Configuration	
Usage	
HBase	
Prerequisites	
	· · · · · ·

HBase RPM Packag	es 2	24
HBase Master Setu	p	26
HBase RegionServe	er Setup	26
HBase Client Setup		26
	Setup [OPTIONAL]	
	[OPTIONAL]	
	n	
	tion Configuration2	
•	3	
	3	
•	3	
•	3	
•		
_		
•	Selver FORTIONAL I	
	Setup [OPTIONAL]	
	[OPTIONAL]	
	ver Setup [OPTIONAL]	
	on Configuration3	
	re Hive Cluster	
Hcatalog	3	36
Prerequisites		36
Hcatalog RPM Pack	ages 3	36
Hcatalog Client Set	up 3	37
Hcatalog Server Se	tup [OPTIONAL] 3	37
Webhcat Setup [OF	PTIONAL] 3	38
Webhcat Server Se	tup [OPTIONAL]	38
Hcatalog Configura	tion	38
_		
•		
•		
_		
_		
	Jes	
	p	
	•	
	on	
_		
•		
	es 4	
	[Optional]	
Flume Configuration	n	43

	Usage	
	Sqoop	43
	Prerequisites	43
	Sqoop RPM Packages	43
	Sqoop Client Setup	44
	Sqoop Metastore Setup [Optional]	44
	Sqoop Configuration	
	Usage	
	Spring Data	
	Installing Spring Data Hadoop	
	Installing Spring Data Hadoop through RPM	
	Spring Data Hadoop	
	HVE (Hadoop Virtualization Extensions)	
	Topology Awareness	
	Elasticity	
	HDFS Rack Awareness	
	Vaidya	
	Vaidya Overview	
	Installing Vaidya Files	
	Enabling and Disabling Vaidya	
	Using Vaidya to Analyze Jobs	
	Vaidya Configuration Rules	
	DataLoader	
Cŀ	napter 4: Unified Storage System (USS)	
	Overview	
	Supported Features	
	Architecture	
	Prerequisites	
	Getting Started	
	USS Namenode (Server)	
	USS Catalog	
	USS Agent	
	USS Client	70
	USS Directory Contents	
	USS RPM Packages	70
	Installing USS	73
	Configuring USS	
	Enabling USS on your Pivotal HD Cluster	73
	USS Namenode Configuration	74
	USS Client Configuration	75
	USS Catalog Configuration	76
	USS Catalog Administration - Using the USS CLI	76
	Basics	76
	General Operations	
	deficial operations	77
	USS Admin Operations	
	USS Admin Operations	77
		77 78
	USS Admin Operations	77 78 78
	USS Admin Operations Using USS Setting Up the Environment	77 78 78 79

Starting the service	83
Stopping the service	83
Restarting the service	83
Checking the service status	83
Uninstallation	84
Known limitations/Future Enhancements	
USS Namenode High Availability	
USS across multiple HDFS versions	
Security	
Catalog Service Integration	
Performance Testing/Benchmarking	
Limited Hive Support	
Troubleshooting/FAQs	
Terminology	
Chapter 5: Security	
Configuring Kerberos for HDFS and YARN (MapReduce)	
Kerberos Set-up	
Install Java Support Items	
Container and Script Modifications	
Site XML Changes	
Complete the HDFS/YARN Secure Configuration	
Turning Secure Mode Off	105
Building and Installing JSVC	
Installing the MIT Kerberos 5 KDC	
Zookeeper Secure Configuration	
Zookeeper Servers	109
HBase Secure Configuration	111
HBase Master and Regionservers	111
HBase Clients	113
Hive Secure Configuration	115
USS Secure Configuration	116
Securing USS	116
Accessing a secure HDFS cluster through USS	116
MapReduce Version 1 Configuration (MRv1)	117
MRv1 Kerberos Set-up Differences	117
Container and Script Modifications Differences	
MRv1 Site XML Differences	118
Auditing	119
Troubleshooting	
-	
Appendix A: Creating a YUM EPEL Repository	122
Appendix B: Installing Binaries	123
Prerequisites	
Hadoop	
HDFS setup	
YARN setup (Yarn distro only)	
MapReduce V1 (MR1) Setup and Run (MR1 Distro only).	
Zookeeper	

Table of Contents vi

Pivotal HD Enterprise 1.0 Stack and Tool Reference Guide

HBase	
Hive	130
HCatalog	
Pig	
Mahout	
Sqoop	134
Flume	
Oozie	

Preface

This guide provides reference information for system administrators and database superusers responsible for installing a Pivotal HD Enterprise system.

- About This Guide
- Getting Support

About Pivotal, Inc.

Greenplum is currently transitioning to a new corporate identity (Pivotal, Inc.). We estimate that this transition will be completed in 2013. During this transition, there will be some legacy instances of our former corporate identity (Greenplum) appearing in our products and documentation. If you have any questions or concerns, please do not hesitate to contact us through our web site:

http://gopivotal.com/about-pivotal/support.

About This Guide

This guide provides information and instructions for manually installing, configuring, and using a Pivotal HD Enterprise system. This guide is intended for system and database administrators responsible for managing a Pivotal HD system.

This guide assumes knowledge of Linux/UNIX system administration, database management systems, database administration, and structured query language (SQL).

Because HAWQ, ships with Pivotal HD is based on PostgreSQL, this guide assumes some familiarity with PostgreSQL.

Getting Support

Pivotal support, product, and licensing information can be obtained as follows.

Product information and Technical Support

For technical support, documentation, release notes, software updates, or for information about Pivotal products, licensing, and services, go to www.gopivotal.com.

Additionally you can still obtain product and support information from the EMC Support Site at: http://support.emc.com.

About Pivotal, Inc.

1. MapReduce V1 (MR1) Based Clusters

This chapter is applicable for those customers who want to deploy a MR1-based cluster instead of a Yarn-based cluster. If you are deploying a Yarn-based cluster, refer to Chapter 2, "Yarn-Based Clusters".

Pivotal HD 1.0.3 supports YARN (MR2) resource manager by default. If you don't want to deploy a YARN based cluster, we provide MR1 as optional manually-installable software.

Note that since MR1 needs to be installed manually, you won't be able to use Pivotal Command Center for monitoring and management of the cluster.

Note that Vaidya is not available if you are deploying a MR1-based cluster.

The base version of Hadoop MR1 is Apache Hadoop 1.0.3.

This chapter includes information and manual installation instructions (where applicable) about the following individual components:

Hadoop MR1

Once you have completed installation and configuration of Hadoop MR1; you can complete your manual stack installation using the instructions provided in Chapter 3, "Additional Stack Components for Manual Installations".

Prerequisities

- Oracle Java Development Kit (JDK) 1.6. Oracle JDK must be installed on every machine before getting started on each Hadoop component.
- You must ensure that time synchronization and DNS are functioning correctly on all client and server machines. For example, you can run the following command to sync the time with NTP server:
 - # service ntpd stop; ntpdate 10.32.97.146; service ntpd start

Installation Notes

In this chapter we install packages by running rpm -ivh <package_name.-<version>.rpm. You can also install them by using yum install <package_name> if you already have yum repository setup (see Appendix A, "Creating a YUM EPEL Repository").

In this document nn within rpm file names represents the rpm build number. It is different for different components.

Hadoop MR1

This section provides instructions for installing each of the following core Hadoop MR1 RPMs:

Prerequisities 2

Hadoop MR1 RPM Packages

hadoop-mr1-1.0.3_gphd_2_0_3_0-nn.x86_64.rpm	
Туре	Core
Requires	bigtop-utils, hdfs core
Description	Hadoop core packages provides the common core packages for running Hadoop
Install on nodes	Every node in the Hadoop cluster and the client workstation that will access the Hadoop service.

hadoop-mr1-jobtracker-1.0.3_gphd_2_0_3_0-nn.x86_64.rpm	
Туре	Daemon
Requires	hadoop-mr1
Description	Hadoop YARN core packages provides common files for running YARN.
Install on nodes	Install on the JobTracker node.

hadoop-mr1-tasktracker-1.0.3_gphd_2_0_3_0-nn.x86_64.rpm	
Туре	Daemon
Requires	hadoop-mr1
Description	Daemon scripts package for Hadoop YARN ResourceManager, which provides a convenient method to manage ResourceManager start/stop as a Linux service.
Install on nodes	Install on the TaskTracker node.

hadoop-mr1-conf-pseudo-1.0.3_gphd_2_0_3_0-nn.x86_64.rpm	
Туре	Configuration
Requires	hadoop-mr1, hadoop-mr1-jobtracker, hadoop-mr1-tasktracker, hadoop-hdfs-datanode, hadoop-hdfs-secondarynamenode, hadoop-yarn-resourcemanager, hadoop-hdfs-namenode
Description	A set of configuration files for running Hadoop in pseudo-distributed mode on one single server.
Install on nodes	Install on the pseudo-distributed host.

Core Package Setup

You must perform the following steps on all the nodes in the Hadoop cluster and its client nodes:

```
$ sudo rpm -ivh
working_dir/utility/rpm/bigtop-utils-0.4.0_gphd_2_0_3_0-nn.n
oarch.rpm
$ sudo rpm -ivh
working_dir/zookeeper/rpm/zookeeper-3.4.5_gphd_2_0_3_0-nn.no
arch.rpm
$ sudo rpm -ivh
working_dir/hadoop/rpm/hadoop-2.0.5_alpha_gphd_2_0_3_0-nn.x8
6_64.rpm
$ sudo rpm -ivh
working_dir/hadoop/rpm/hadoop-hdfs-2.0.5_alpha_gphd_2_0_3_0-nn.x86
64.rpm
```

Hadoop-mr1 JobTracker Setup

Install the Hadoop-mr1 JobTracker package on the workstation that will serve as the Hadoop-mr1 JobTracker:

```
$ sudo rpm -ivh
working_dir/hadoop-mr1/rpm/hadoop-mr1-jobtracker-1.0.3_gphd_
2  0  3  0-nn.x86  64.rpm
```

Hadoop-mr1 TaskTracker Setup

Install the Hadoop-mr1 TaskTracker package on the workstation that will serve as the Hadoop-mr1 TaskTracker:

```
$ sudo rpm -ivh
working_dir/hadoop-mr1/rpm/hadoop-mr1-tasktracker-1.0.3_gphd
2 0 3 0-nn.x86 64.rpm
```

Hadoop MR1 Configuration

The configuration files for Hadoop MR1 are located here:

/etc/gphd/hadoop/conf/

Out of the box by default it is a symbolic link to

/etc/gphd/hadoop-version/conf.empty template directory.

You can make modifications to these configuration templates or create your own configuration set. If you want to use a different configuration folders, adjust the /etc/gphd/hadoop/conf symbolic link to point to the folder you want to utilize at runtime.

If you want to run Hadoop MR1 in one host in pseudo-distributed mode on one single host, you can go through all the above setup steps on your host and then install the hadoop-mr1-conf-pseudo package:

```
$ sudo rpm -ivh
working_dir/hadoop-mr1/rpm/hadoop-mr1-conf-pseudo-1.0.3_gphd
_2_0_3_0-nn.x86_64.rpm
```

Usage

Starting Hadoop-MR1

Before you start Hadoop, you need to create some working directories on HDFS, as follows:

```
// create mapred.system.dir
# sudo -u hdfs hdfs dfs -mkdir -p /mapred/system
# sudo -u hdfs hdfs dfs -chown -R mapred:hadoop /mapred
// create mapreduce.jobtracker.staging.root.dir staging
directory
# sudo -u hdfs hdfs dfs -mkdir -p /user
```

Starting Hadoop-mr1 JobTracker

To start the hadoop-mr1 jobtracker daemon:

Ether:

```
$ sudo service hadoop-mr1-jobtracker start
or:
$ sudo /etc/init.d/hadoop-mr1-jobtracker start
```

Starting Hadoop-mr1 TaskTracker

To start the hadoop-mr1 tasktracker daemon:

Ether:

```
$ sudo service hadoop-mr1-tasktracker start
or:
$ sudo /etc/init.d/hadoop-mr1-tasktracker start
```

Using Hadoop-mr1

After JobTracker and TaskTracker are started, you can now submit MapReduce Jobs.

Note: Make sure HDFS daemons are running, and create the home directory /user/\${user.name} for each MapReduce user on hdfs. In these examples we use the user hadoop.

```
# sudo -u hdfs hdfs dfs -mkdir -p /user/hadoop
# sudo -u hdfs hdfs dfs -chown hadoop:hadoop /user/hadoop
```

Here is an example MapReduce job:

```
# su - hadoop
$ hadoop-mr1 jar
/usr/lib/gphd/hadoop-mr1-1.0.3_gphd_2_0_3_0/hadoop-examples-
*.jar pi 2 10000
```

This will run the PI generation example. You can track the progress of this job at the JobTracker dashboard: http://jobtracker-host:50030/.

Stop Hadoop-mr1

Stop Hadoop-mr1 JobTracker

To stop the hadoop-mr1 jobtracker daemon:

```
Ether:
```

```
$ sudo service hadoop-mr1-jobtracker stop
or:
$ sudo /etc/init.d/hadoop-mr1-jobtracker stop
```

Stop Hadoop-mr1 TaskTracker

To stop the hadoop-mrl tasktracker daemon:

Ether:

```
$ sudo service hadoop-mr1-tasktracker stop
or:
$ sudo /etc/init.d/hadoop-mr1-tasktracker stop
```

Now you have completed installation of MR1 components; you can complete your manual stack installation using the instructions provided in Chapter 3, "Additional Stack Components for Manual Installations".

2. Yarn-Based Clusters

This chapter is applicable for those customers who are manually deploying a Yarn-based cluster. If you are deploying a MapReduce V1 (MR1)-based cluster, refer to Chapter 1, "MapReduce V1 (MR1) Based Clusters".

Pivotal HD 1.0.3 supports YARN (MR2) resource manager by default. If you don't want to deploy a YARN based cluster, we provide MR1 as optional manually-installable software.

The base version of Hadoop YARN/MapReduce is Apache Hadoop 2.0.5-alpha.

This chapter describes how to manually install, configure, and use the Pivotal HD 1.0.3 distribution of the following Apache Hadoop components::

- HDFS
- YARN
- MapReduce

Once you have completed installation and configuration of these YARN components; you can complete your manual stack installation using the instructions provided in Chapter 3, "Additional Stack Components for Manual Installations".

Prerequisities

- Oracle Java Development Kit (JDK) 1.6. Oracle JDK must be installed on every machine before getting started on each Hadoop component.
- You must ensure that time synchronization and DNS are functioning correctly on all client and server machines. For example, you can run the following command to sync the time with NTP server:
 - # service ntpd stop; ntpdate 10.32.97.146; service ntpd start

Installation Notes

In this chapter we install packages by running rpm -ivh <package_name.-<version>.rpm. You can also install them by using yum install <package_name> if you already have yum repository setup (see Appendix A, "Creating a YUM EPEL Repository").

In this document nn within rpm file names represents the rpm build number. It is different for different components.

Hadoop HDFS

This section provides instructions for installing each of the following core Hadoop RPMs:

HDFS Namenode Setup

Prerequisities 7

- HDFS Datanode Setup
- HDFS Secondary Namenode Setup

Hadoop HDFS RPM Packages

Pivotal provides the following RPMs as part of this release. The core packages provide all executables, libraries, configurations, and documentation for Hadoop and is required on every node in the Hadoop cluster as well as the client workstation that will access the Hadoop service. The daemon packages provide a convenient way to manage Hadoop HDFS daemons as Linux services, which rely on the core package.

hadoop-2.0.5_alpha_gphd_2_0_3_0-nn.x86_64.rpm	
Туре	Core
Requires	bigtop-utils, zookeeper-core
Description	Hadoop core packages provides the common core packages for running Hadoop
Install on Nodes	Every node in the Hadoop cluster and the client workstation that will access the Hadoop service.

hadoop-hdfs-2.0.5_alpha_gphd_2_0_3_0-nn.x86_64.rpm	
Туре	Core
Requires	hadoop, bigtop-jsvc
Description	Hadoop HDFS core packages provides the common files for running HFS.
Install on Nodes	Every node in the HDFS cluster and the client workstation that will access the HDFS.

hadoop-hdfs-namenode-2.0.5_alpha_gphd_2_0_3_0-nn.x86_64.rpm	
Туре	Daemon
Requires	hadoop-hdfs
Description	Daemon scripts package for Hadoop Namenode, which provides a convenient method to manage Namenode start/stop as a Linux service.
Install on Nodes	Only on HDFS Namenode server.

hadoop-hdfs-datanode-2.0.5_alpha_gphd_2_0_3_0-nn.x86_64.rpm	
Туре	Daemon
Requires	hadoop-hdfs

hadoop-hdfs-datanode-2.0.5_alpha_gphd_2_0_3_0-nn.x86_64.rpm	
Description	Daemon scripts package for Hadoop Datanode, which provides a convenient method to manage datanode start/stop as a Linux service.
Install on Nodes	Install on all HDFS Datanodes.

hadoop-hdfs-secondarynamenode-2.0.5_alpha_gphd_2_0_3_0-nn.x86_64.rpm	
Туре	Daemon
Requires	hadoop-hdfs
Description	Daemon scripts package for Hadoop SecondaryNamenode, which provides a convenient method to manage SecondaryNamenode start/stop as a Linux service.
Install on Nodes	Install on one server that will be acting as the Secondary Namenode.

hadoop-hdfs-journalnode-2.0.5_alpha_gphd_2_0_3_0-nn.x86_64.rpm	
Туре	Daemon
Requires	hadoop-hdfs
Description	Daemon scripts package for Hadoop JournalNode, which provides a convenient method to manage journalnode start/stop as a Linux service.
Install on Nodes	Install on all HDFS JournalNodes.

hadoop-hdfs-zkfc-2.0.5_alpha_gphd_2_0_3_0-nn.x86_64.rpm	
Туре	Daemon
Requires	hadoop-hdfs
Description	Daemon scripts package for Hadoop zkfc, which provides a convenient method to manage zkfc start/stop as a Linux service.
Install on Nodes	Install on all HDFS zkfc nodes.

hadoop-hdfs-fuse-2.0.5_alpha_gphd_2_0_3_0-nn.x86_64.rpm	
Туре	Core
Requires	hadoop-libhdfs, hadoop-client
Description	Binaries that can be used to mount hdfs as a local directory.
Install on Nodes	Install on the servers that want to mount the HDFS.

hadoop-libhdfs-2.0.5_alpha_gphd_2_0_3_0-nn.x86_64.rpm	
Туре	Core
Requires	hadoop-hdfs
Description	Native implementation of the HDFS.
Install on Nodes	Install on servers that you want to run native hdfs.

hadoop-httpfs-2.0.5_alpha_gphd_2_0_3_0-nn.x86_64.rpm	
Туре	Core
Requires	bigtop-tomcat, hadoop, hadoop-hdfs
Description	HttpFS is a server that provides a REST HTTP gateway supporting all HDFS File System operations (read and write).
Install on Nodes	Install on servers that will be serving REST HDFS service

hadoop-doc-2.0.5_alpha_gphd_2_0_3_0-nn.x86_64.rpm	
Туре	Doc
Requires	N/A
Description	Document package provides the Hadoop document.
Install on Nodes	Install on whichever host that user want to read hadoop documentation.

hadoop-conf-pseudo-2.0.5_alpha_gphd_2_0_3_0-nn.x86_64.rpm	
Туре	Configuration
Requires	hadoop-hdfs-datanode, hadoop-hdfs-secondarynamenode, hadoop-yarn-resourcemanager, hadoop-hdfs-namenode, hadoop-yarn-nodemanager, hadoop-mapreduce-historyserver
Description	A set of configuration files for running Hadoop in pseudo-distributed mode on one single server.
Install on Nodes	Install on the pseudodistributed host.

hadoop-client-2.0.5_alpha_gphd_2_0_3_0-nn.x86_64.rpm	
Туре	Library
Requires	N/A

hadoop-client-2.0.5_alpha_gphd_2_0_3_0-nn.x86_64.rpm	
Description	A set of symbolic link which gathers the libraries for programming Hadoop and submit Hadoop jobs.
Install on Nodes	Clients nodes that will be used to submit hadoop jobs.

Prerequisites: Core Package Setup

You must perform the following steps on all the nodes in the Hadoop cluster and its client nodes:

```
$ sudo rpm -ivh
working_dir/utility/rpm/bigtop-utils-0.4.0_gphd_2_0_3_0-nn.noarch.rpm
$ sudo rpm -ivh
working_dir/zookeeper/rpm/zookeeper-3.4.5_gphd_2_0_3_0-nn.noarch.rpm
$ sudo rpm -ivh
working_dir/hadoop/rpm/hadoop-2.0.5 alpha gphd 2 0 3 0-nn.x86 64.rpm
```

Where working dir is the directory where you want the rpms expanded.

HDFS Namenode Setup

Install the Hadoop Namenode package on the workstation that will serve as HDFS Namenode:

```
$ sudo rpm -ivh
working_dir/utility/rpm/bigtop-jsvc-1.0.15_gphd_2_0_3_0-nn.x86_64.rpm
$ sudo rpm -ivh
working_dir/hadoop/rpm/hadoop-hdfs-2.0.5_alpha_gphd_2_0_3_0-nn.x86_64.rpm
$ sudo rpm -ivh
working_dir/hadooop/rpm/hadoop-hdfs-namenode-2.0.5_alpha_gphd_2_0_3_0-nn.x8
6 64.rpm
```

HDFS Datanode Setup

Install the Hadoop Datanode package on the workstation that will serve as HDFS Datanode:

```
$ sudo rpm -ivh
working_dir/utility/rpm/bigtop-jsvc-1.0.15_gphd_2_0_3_0-nn.x86_64.rpm
$ sudo rpm -ivh
working_dir/hadoop/rpm/hadoop-hdfs-2.0.5_alpha_gphd_2_0_3_0-nn.x86_64.rpm
$ sudo rpm -ivh
working_dir/hadoop/rpm/hadoop-hdfs-datanode-2.0.5_alpha_gphd_2_0_3_0-nn.x86
64.rpm
```

HDFS Secondary Namenode Setup

Install the Hadoop Secondary Namenode package on the workstation that will serve as HDFS Secondary Namenode:

```
$ sudo rpm -ivh
working_dir/utility/rpm/bigtop-jsvc-1.0.10_gphd_2_0_3_0-nn.x86_64.rpm
$ sudo rpm -ivh
working_dir/hadoop/rpm/hadoop-hdfs-2.0.5_alpha_gphd_2_0_3_0-nn.x86_64.rpm
$ sudo rpm -ivh
working_dir/hadoop/rpm/hadoop-secondarynamenode-2.0.5_gphd_2_0_3_0-nn.x86_6
4.rpm
```

Hadoop HDFS Configuration

The configuration files for Hadoop are located here: /etc/gphd/hadoop/conf/

Out of the box by default it is a symbolic link to

/etc/gphd/hadoop-version/conf.empty template directory.

You can make modifications to these configuration templates or create your own configuration set. If you want to use a different configuration folder, edit the /etc/gphd/hadoop/conf symbolic link to point to the folder you want to utilize at runtime.

If you want to run Hadoop 2.0 in one host in pseudo-distributed mode on one single host, you can make sure all the dependent packages of hadoop-conf-pseudo have been installed on your host and then install the hadoop-conf-pseudo package:

```
$ sudo rpm -ivh
working_dir/hadoop/rpm/hadoop-conf-pseudo-2.0.5_alpha_gphd_2
_0_3_0-nn.x86_64.rpm
```

Refer to Apache Hadoop 2.0.5-alpha documentation for how to configure Hadoop in distributed mode. This document describes how to use Hadoop in a pseudo-distributed mode.

Usage

After installing the daemon package for Hadoop, you can start the daemons.

Starting HDFS

HDFS includes three main components: Namenode, Datanode, Secondary Namenode.

To start the Namenode daemon:

You need to format the Namenode before starting it, as follows:

```
$ sudo -u hdfs hdfs namenode -format
```

Note: You only have to do this once. But if you have changed the hadoop namenode configuration, you may need to run this again.

Then start the Namenode by running either:

```
$ sudo service hadoop-hdfs-namenode start
or:
$ sudo hdfs hadoop-hdfs-namenode start
```

When Namenode is started, you can visit its dashboard at:

```
http://localhost:50070/
```

To start the Datanode daemon:

Run either:

```
$ sudo service hadoop-hdfs-datanode start
or:
$ sudo hdfs hadoop-hdfs-datanode start
```

To start the Secondary Namenode daemon:

Run either:

```
$ sudo service hadoop-hdfs-secondarynamenode start
or:
$ sudo hdfs hadoop-hdfs-secondarynamenode start
```

Using HDFS

When the HDFS components are started, you can try some HDFS usage, for example:

```
$ sudo -u hdfs hdfs dfs -ls /
$ sudo -u hdfs hdfs dfs -mkdir -p /user/hadoop
$ sudo -u hdfs hdfs dfs -chown -R hadoop:hadoop /user/hadoop
#you can see a full list of hdfs dfs command options
$ hdfs dfs
$ bin/hdfs dfs -copyFromLocal /etc/passwd /user/hadoop/
```

Note: by default, the root folder is owned by user hdfs, so you have to use sudo -u hdfs *** to execute the first few commands.

Shutting down HDFS

Stop the Namenode Daemon:

Run either:

```
$ sudo service hadoop-hdfs-namenode stop
or:
$ sudo hdfs hadoop-hdfs-namenode stop
```

Stop the Datanode Daemon:

Run either:

```
$ sudo service hadoop-hdfs-datanode stop
or:
$ sudo hdfs hadoop-hdfs-datanode stop
```

Stop the Secondary Namenode Daemon:

Run either:

```
$ sudo service hadoop-hdfs-secondarynamenode stop
```

or:

\$ sudo hdfs hadoop-hdfs-secondarynamenode stop

Hadoop YARN

This section provides instructions for installing each of the following core Hadoop YARN RPMs:

- YARN ResourceManager Setup
- YARN NodeManager Setup
- Mapreduce HistoryServer Setup
- YARN ProxyServer Setup

Hadoop YARN RPM Packages

Pivotal provides the following RPMs as part of this release. The core packages provide all executables, libraries, configurations, and documentation for Hadoop and is required on every node in the Hadoop cluster as well as the client workstation that will access the Hadoop service. The daemon packages provide a convenient way to manage Hadoop YARN daemons as Linux services, which rely on the core package.

hadoop-2.0.5_alpha_gphd_2_0_3_0-nn.x86_64.rpm	
Туре	Core
Requires	bigtop-utils, zookeeper-core
Description	Hadoop core packages provides the common core packages for running Hadoop.
Install on Nodes	Every node in the Hadoop cluster and the client workstation that will access the Hadoop service.

hadoop-yarn-2.0.5_alpha_gphd_2_0_3_0-nn.x86_64.rpm	
Туре	Core
Requires	hadoop
Description	Hadoop YARN core packages provides common files for running YARN.
Install on Nodes	Install on all YARN nodes.

hadoop-yarn-resourcemanager-2.0.5_alpha_gphd_2_0_3_0-nn.x86_64.rpm	
Туре	Daemon
Requires	hadoop-yarn

hadoop-yarn-resourcemanager-2.0.5_alpha_gphd_2_0_3_0-nn.x86_64.rpm	
Description	Daemon scripts package for Hadoop YARN ResourceManager, which provides a convenient method to manage ResourceManager start/stop as a Linux service.
Install on Nodes	Install on the Resource Manager node.

hadoop-yarn-nodemanager-2.0.5_alpha_gphd_2_0_3_0-nn.x86_64.rpm	
Туре	Daemon
Requires	hadoop-yarn
Description	Daemon scripts package for Hadoop YARN NodeManager, which provides a convenient method to manage NodeManager start/stop as a Linux service.
Install on Nodes	Install on all the Node Manager nodes.

hadoop-yarn-proxyserver-2.0.5_alpha_gphd_2_0_3_0-nn.x86_64.rpm	
Туре	Daemon
Requires	hadoop-yarn
Description	Daemon scripts package for Hadoop YARN ProxyServer, which provides a convenient method to manage ProxyServer start/stop as a Linux service.
Install on Nodes	Install on the node that will act as a proxy server from the user to applicationmaster

hadoop-mapreduce-2.0.5_alpha_gphd_2_0_3_0-nn.x86_64.rpm	
Туре	Core
Requires	hadoop-yarn
Description	Hadoop Mapreduce core libraries.
Install on Nodes	Install on all ResourceManager and NodeManager nodes.

hadoop-mapreduce-historyserver-2.0.5_alpha_gphd_2_0_3_0-nn.x86_64.rpm	
Туре	Daemon
Requires	hadoop, hadoop-mapreduce
Description	Daemon scripts package for Hadoop MapReduce HistoryServer, which provides a convenient method to manage MapReduce HistoryServer start/stop as a Linux service.
Install on Nodes	Install on the host that will be acting as the MapReduce History Server.

hadoop-doc-2.0.5_alpha_gphd_2_0_3_0-nn.x86_64.rpm	
Туре	Doc
Requires	N/A
Description	Document package provides the Hadoop documentation.
Install on Nodes	Install on whichever host that user want to read hadoop doc.

hadoop-conf-pseudo-2.0.5_alpha_gphd_2_0_3_0-nn.x86_64.rpm	
Туре	Configuration
Requires	hadoop-hdfs-datanode, hadoop-hdfs-secondarynamenode, hadoop-yarn-resourcemanager, hadoop-hdfs-namenode hadoop-yarn-nodemanager, hadoop-mapreduce-historyserver
Description	A set of configuration files for running Hadoop in pseudo-distributed mode on one single server.
Install on Nodes	Install on the pseudu-distributed host.

hadoop-client-2.0.5_alpha_gphd_2_0_3_0-nn.x86_64.rpm	
Туре	Library
Requires	hadoop, hadoop-hdfs, hadoop-yarn, hadoop-mapreduce
Description	A set of symbolic link which gathers the libraries for programming Hadoop and submit Hadoop jobs.
Install on Nodes	Clients nodes that will be used to submit hadoop jobs.

Prerequisites: Core Package Setup

You must perform the following steps on all the nodes in the Hadoop cluster and its client nodes:

```
$ sudo rpm -ivh
working_dir/utility/rpm/bigtop-utils-0.4.0_gphd_2_0_3_0-nn.noarch.rpm
$ sudo rpm -ivh
working_dir/zookeeper/rpm/zookeeper-3.4.5_gphd_2_0_3_0-nn.noarch.rpm
$ sudo rpm -ivh
working_dir/hadoop/rpm/hadoop-2.0.5_alpha_gphd_2_0_3_0-nn.x86_64.rpm
```

Where working_dir is the directory where you want the rpms expanded.

YARN ResourceManager Setup

Install the YARN ResourceManager package on the workstation that will serve as YARN ResourceManager:

```
$ sudo rpm -ivh
working_dir/hadoop/rpm/hadoop-yarn-2.0.5_alpha_gphd_2_0_3_0-nn.x86_64.rpm
$ sudo rpm -ivh
working_dir/hadoop/rpm/hadoop-yarn-resourcemanager-2.0.5_alpha_gphd_2_0_3_0
-nn.x86_64.rpm
```

YARN NodeManager Setup

Install the YARN NodeManager package on the workstation that will serve as YARN nodes:

```
$ sudo rpm -ivh
working_dir/hadoop/rpm/hadoop-yarn-2.0.5_alpha_gphd_2_0_3_0-nn.x86_64.rpm
$ sudo rpm -ivh
working_dir/hadoop/rpm/hadoop-yarn-nodemanager-2.0.5_alpha_gphd_2_0_3_0-nn.
x86_64.rpm
```

Mapreduce HistoryServer Setup

Install the YARN Mapreduce History Manager package and its dependency packages on the workstation that will serve as the MapReduce History Server:

```
$ sudo rpm -ivh
working_dir/hadoop/rpm/hadoop-yarn-2.0.5_alpha_gphd_2_0_3_0-nn.x86_64.rpm
$ sudo rpm -ivh
working_dir/hadoop/rpm/hadoop-mapreduce-2.0.5_alpha_gphd_2_0_3_0-nn.x86_64.rpm
$ sudo rpm -ivh
working_dir/hadoop/rpm/hadoop-mapreduce-historyserver-2.0.5_alpha_gphd_2_0_
3 0-nn.x86 64.rpm
```

YARN ProxyServer Setup

Install the YARN Proxy Server package and its dependency packages on the workstation that will serve as the YARN Proxy Server.

```
$ sudo rpm -ivh
working_dir/hadoop/rpm/hadoop-yarn-2.0.5_alpha_gphd_2_0_3_0-nn.x86_64.rpm
$ sudo rpm -ivh
working_dir/hadoop/rpm/hadoop-yarn-proxyserver-2.0.5_alpha_gphd_2_0_3_0-nn.
x86_64.rpm
```

Hadoop HDFS Configuration

The configuration files for Hadoop are located here: /etc/gphd/hadoop/conf/

Out of the box by default it is a symbolic link to

/etc/gphd/hadoop-version/conf.empty template.

You can make modifications to these configuration templates or create your own configuration set. If you want to use a different configuration folder, adjust the /etc/gphd/hadoop/conf symbolic link to point to the folder you want to utilize at runtime.

If you want to run Hadoop 2.0 in one host in pseudo-distributed mode on one single host, you can go through all the above setup steps on your host and then install the hadoop-conf-pseudo package, as follows:

```
$ sudo rpm -ivh
working_dir/hadoop/rpm/hadoop-conf-pseudo-2.0.5_alpha_gphd_2_0_3_0
-nn.x86 64.rpm
```

Refer to Apache Hadoop 2.0.5-alpha documentation for how to configure Hadoop in distributed mode. This document describes how to use Hadoop in a pseudo-distributed mode.

YARN Usage

Starting YARN

YARN includes three services: ResourceManager (RM), NodeManager (NM), MapReduce HistoryManager (MRHM). RM and NM are required, MRHM is optional.

Before you start these services, you need to create some working directories on HDFS, as follows:

Create working directories on HDFS

```
$ sudo -u hdfs hdfs dfs -mkdir /tmp
$ sudo -u hdfs hdfs dfs -chmod 777 /tmp
$ sudo -u hdfs hdfs dfs -mkdir -p /var/log/gphd/hadoop-yarn
$ sudo -u hdfs hdfs dfs -chown yarn:hadoop
/var/log/gphd/hadoop-yarn
$ sudo -u hdfs hdfs dfs -mkdir -p /user/history
$ sudo -u hdfs hdfs dfs -chown mapred:hadoop /user/history
$ sudo -u hdfs hdfs dfs -chmod -R 777 /user/history
$ sudo -u hdfs hdfs dfs -mkdir -p /user/hadoop
$ sudo -u hdfs hdfs dfs -chown hadoop:hadoop /user/hadoop
```

Starting ResourceManager

RM daemon need to be started only on the master node.

Run either:

```
$ sudo service hadoop-yarn-resourcemanager start
or:
$ sudo yarn hadoop-yarn-resourcemanager start
```

When RM is started, you can visit its dashboard at: http://localhost:8088/

Starting NodeManager

NM daemon needs to be started on all hosts that will be used as working nodes.

Run either

```
$ sudo service hadoop-yarn-nodemanager start
or:
$ sudo yarn hadoop-yarn-nodemanager start
```

Start MapReduce HistoryServer

MapReduce HistoryServer only needs to be run on the server that is meant to be the history server. It is an optional service and should only be enabled if you want to keep track of the MapReduce jobs that have been run.

Run:

```
$ sudo service hadoop-mapreduce-historyserver start
```

When the MR HistoryServer is started, you can visit its dashboard at: http://localhost:19888/

Using YARN

After RM and NM are started, you can now submit YARN applications.

For simplicity, we assume you are running Hadoop in pseudo-distributed mode using the default pseudo configuration.

Note: Make sure HDFS daemons are running.

Here is an example MapReduce job:

```
$ hadoop jar
/usr/lib/gphd/hadoop-mapreduce/hadoop-mapreduce-examples-2.0
.5-alpha-gphd-2.0.3.0.jar pi 2 200
```

This will run the PI generation example. You can track the progress of this job at the RM dashboard: http://localhost:8088/

You can also run other MapReduce examples, the following command will print a list of available examples:

```
$ hadoop jar
/usr/lib/gphd/hadoop-mapreduce/hadoop-mapreduce-examples-2.0
.5-alpha-gphd-2.0.3.0.jar
```

Stopping YARN

You can stop the YARN daemons manually by using the following commands.

To stop the MapReduce HistoryServer Daemon:

Run:

```
$ sudo service hadoop-mapreduce-historyserver stop
```

To stop the NodeManager Daemon:

Run:

\$ sudo service hadoop-yarn-nodemanager stop

To stop the ResourceManager Daemon:

Run:

\$ sudo service hadoop-yarn-resourcemanager stop

Now you have completed installation of MR2 components; you can complete your manual stack installation using the instructions provided in Chapter 3, "Additional Stack Components for Manual Installations".

3. Additional Stack Components for Manual Installations

If you are deploying an MR1-based cluster, you should have already manually installed HDFS and MR1, as described in Chapter 1, "MapReduce V1 (MR1) Based Clusters".

If you are manually deploying a YARN-based cluster, you should have already have manually installed MR2, as described in Chapter 2, "Yarn-Based Clusters".

This chapter describes how to manually install, configure and use other components of the Pivotal HD 1.0.2 distribution of Apache Hadoop.

• Zookeeper: Zookeeper 3.4.5

HBase: HBase 0.94.8

• Hive: Hive 0.11.0

Pig: Pig 0.10.1

Mahout: Mahout 0.7

• Flume: Flume 1.3.1

Sqoop: Sqoop 1.4.2

Spring Data

• HVE (Hadoop Virtualization Extensions)

Vaidya

DataLoader

For information about the installation, configuration, and use of the USS component is provided in Chapter 4, "Unified Storage System (USS)"

Prerequisities

If not already met:

- Oracle Java Development Kit (JDK) 1.6. Oracle JDK must be installed on every machine before getting started on each Hadoop component.
- You must ensure that time synchronization and DNS are functioning correctly on all client and server machines. For example, you can run the following command to sync the time with NTP server:

```
# service ntpd stop; ntpdate 10.32.97.146; service ntpd start
```

In this document nn within rpm file names represents the rpm build number. It is different for different components.

Prerequisities 21

Installation Notes

In this chapter we install packages by running rpm -ivh package_name-<version>.rpm. You can also install them by using yum install <package_name> if you already have yum repository setup (see Appendix A, "Creating a YUM EPEL Repository").

In this document nn within rpm file names represents the rpm build number. It is different for different components.

Zookeeper

The base version of ZooKeeper is Apache ZooKeeper 3.4.5.

ZooKeeper is a high-performance coordination service for distributed applications.

This section describes how to install, configure, and use Zookeeper.

Zookeeper RPM Packages

Pivotal HD provides the following RPMs as part of this release. The core package provides all executable, libraries, configurations, and documentation for Zookeeper and is required on every node in the Zookeeper cluster as well as the client workstation that will access the Zookeeper service. The daemon packages provide a convenient way to manage Zookeeper daemons as Linux services, which rely on the core package.

Note: Zookeeper doesn't require Hadoop Core Packages.

zookeeper-3.4.5_gphd_2_0_3_0-nn.noarch.rpm	
Туре	Core
Requires	N/A
Description	Zookeeper core package which provides the executable, libraries, configuration files and documentations.
Install on Nodes	Every node in the ZooKeeper cluster, and the client workstations which will access the ZooKeeper service.

zookeeper-server-3.4.5_gphd_2_0_3_0-nn.noarch.rpm	
Туре	Deamon
Requires	ZooKeeper Core Package
Description	Daemon scripts package for Zookeeper server, which provides a convenient method to manage Zookeeper server start/stop as a Linux service.
Install on Nodes	N/A

Installation Notes 22

zookeeper-doc-3.4.5_gphd_2_0_3_0-nn.noarch.rpm	
Туре	Documentation
Requires	N/A
Description	Zookeeper documentation.
Install on Nodes	N/A

Zookeeper Server Setup

Install the Zookeeper core package and the Zookeeper server daemon package on the workstation that will serve as the zookeeper server, as follows:

```
$ sudo rpm -ivh
working_dir/zookeeper/rpm/zookeeper-3.4.5_gphd_2_0_3_0-nn.noarch.r
pm
$ sudo rpm -ivh
working_dir/zookeeper/rpm/zookeeper-server-3.4.5_gphd_2_0_3_0-nn.n
oarch.rpm
```

Where working dir is the directory where you want the rpms expanded.

Zookeeper Client Setup

Install the Zookeeper core package on the client workstation to access the Zookeeper service, as follows:

```
$ sudo rpm -ivh
working_dir/zookeeper/rpm/zookeeper-3.4.5_gphd_2_0_3_0-nn.no
arch.rpm
```

Zookeeper Configuration

Zookeeper configuration files are in the following location:

```
/etc/gphd/zookeeper/conf
```

This is the default configuration for quick reference and modification. It is a symbolic link to /etc/gphd/zookeeper-version/conf.dist template set.

You can make modifications to these configuration templates or create your own configuration set. If you want to use a different configuration folders, adjust the symbolic link /etc/gphd/zookeeper to point to the folder you want to utilize at runtime.

Usage

Starting the Zookeeper Daemon

After installing the daemon package for Zookeeper, the Zookeeper server daemon will start automatically at system startup by default.

Zookeeper 23

You can start the daemons manually by using the following commands.

Run:

```
$ sudo service zookeeper-server start
```

Accessing the Zookeeper service

To access the Zookeeper service on a client machine, use the command zookeeper-client directly in shell:

```
$ zookeeper-client
```

In the ZK shell:

```
> ls
> create /zk_test my_data
> get /zk_test
> quit
```

You can get a list of available commands by inputting "?" in the zookeeper shell.

Stopping the Zookeeper Daemon

You can stop the Zookeeper server daemon manually using the following commands:

Run:

```
$ sudo service zookeeper-server stop
```

HBase

The base version of HBase changed to Apache HBase 0.94.8.

HBase is a scalable, distributed database that supports structured data storage for large tables.

This section specifies how to install, configure, and use HBase.

Prerequisites

As HBase is built on top of Hadoop and Zookeeper, the Hadoop and Zookeeper core packages must be installed for HBase to operate correctly.

HBase RPM Packages

Pivotal HD provides the following RPMs as part of this release. The core package provides all executables, libraries, configurations and documentation for HBase and is required on every node in HBase cluster as well as the client workstation that wants to access the HBase service. The daemon packages provide a convenient way to manage HBase daemons as Linux services, which rely on the core package.

hbase-0.94.8_gphd_2_0_3_0-nn.noarch.rpm	
Туре	Core
Requires	Hadoop HDFS Packages and ZooKeeper Core Package
Description	HBase core package provides all executables, libraries, configuration files and documentations.

hbase-master-0.94.8_gphd_2_0_3_0-nn.noarch.rpm	
Туре	Daemon
Requires	HBase Core Package
Description	Daemon scripts package for HMaster, which provides a convenient method to manage HBase HMaster server start/stop as a Linux service.

hbase-regionserver-0.94.8_gphd_2_0_3_0-nn.noarch.rpm	
Туре	Daemon
Requires	HBase Core Package
Description	Daemon scripts package for HRegionServer, which provides a convenient method to manage HBase HRegionServer start/stop as a Linux service.

hbase-thrift-0.94.8_gphd_2_0_3_0-nn.noarch.rpm	
Туре	Daemon (thrift service)
Requires	HBase Core Package
Description	Daemon scripts package to provide HBase service through thrift.

hbase-rest-0.94.8_gphd_2_0_3_0-nn.noarch.rpm	
Туре	Daemon (Restful service)
Requires	HBase Core Package
Description	Daemon scripts package to provide HBase service through REST.

hbase-doc-0.94.8_gphd_2_0_3_0-nn.noarch.rpm	
Туре	Documentation
Requires	HBase Core Package
Description	HBase documentation.

HBase Master Setup

Install the HBase core package and the HBase master daemon package on the workstation that will serve as the HMaster:

```
$ sudo rpm -ivh
working_dir/hbase/rpm/hbase-0.94.8_gphd_2_0_3_0-nn.noarch.rp
m
$ sudo rpm -ivh
working_dir/hbase/rpm/hbase-master-0.94.8_gphd_2_0_3_0-nn.no
arch.rpm
```

HBase RegionServer Setup

Install the HBase core package and the HBase regionserver daemon package on the workstation that will serve as the HRegionServer:

```
$ sudo rpm -ivh
working_dir/hbase/rpm/hbase-0.94.8_gphd_2_0_3_0-nn.noarch.rp
m
$ sudo rpm -ivh
working_dir/hbase/rpm/hbase-regionserver-0.94.8_gphd_2_0_3_0
-nn.noarch.rpm
```

HBase Client Setup

Install the HBase core package on the client workstation that will access the HBase service:

```
$ sudo rpm -ivh
working_dir/hbase/rpm/hbase-0.94.8_gphd_2_0_3_0-nn.noarch.rp
m
```

HBase Thrift Server Setup [OPTIONAL]

Install the HBase core package and the HBase thrift daemon package to provide HBase service through Apache Thrift:

```
$ sudo rpm -ivh
working_dir/hbase/rpm/hbase-0.94.8_gphd_2_0_3_0-nn.noarch.rp
m
$ sudo rpm -ivh
working_dir/hbase/rpm/hbase-thrift-0.94.8_gphd_2_0_3_0-nn.no
```

```
arch.rpm
```

REST Server Setup [OPTIONAL]

Install the HBase core package and the HBase rest daemon package to provide HBase service through Restful interface:

```
$ sudo rpm -ivh
working_dir/hbase/rpm/hbase-0.94.8_gphd_2_0_3_0-nn.noarch.rp
m
$ sudo rpm -ivh
working_dir/hbase/rpm/hbase-rest-0.94.8_gphd_2_0_3_0-nn.noar
ch.rpm
```

HBase Configuration

The configuration files for HBase are located here: /etc/gphd/hbase/conf/

This is the default configuration for quick reference and modification.

/etc/gphd/hbase is a symbolic link to /etc/gphd/hbase-version/; and the conf folder is a symbolic link to the exact configuration directory.

You can make modifications to these configuration templates or create your own configuration set. If you want to use a different configuration folders, adjust the symbolic link /etc/gphd/hbase/conf to point to the folder you want to utilize at runtime

HBase Post-installation Configuration

- **1.** Login to one of the cluster nodes.
- 2. Create the hbase.rootdir
 \$ sudo -u hdfs hdfs dfs -mkdir -p /hbase
- **3.** Set permissions for the hbase. rootdir

```
# sudo -u hdfs hdfs dfs -chown hbase:hadoop /hbase
```

4. Set the ownership for the hbase. rootdir

```
$ sudo -u hdfs hadoop fs -chown hbase:hadoop /hbase
```

5. Add hbase user to the hadoop group if not already present using

```
$ usermod -G hadoop hbase
```

Usage

Starting the HBase Daemon

After installing the daemon package for HBase, the HBase server daemons will start automatically at system startup by default.

You can start the daemons manually by using the following commands:

Run:

```
$ sudo service hbase-master start
```

Starting the HRegionServer daemon

You can start the daemons manually by using the following commands:

Run:

```
$ sudo service hbase-regionserver start
```

Starting the Hbase Thrift server daemon [OPTIONAL]

You can start the daemons manually by using the following commands:

Run:

```
$ sudo service hbase-thrift start
```

Starting the Hbase Rest server daemon [OPTIONAL]

You can start the daemons manually by using the following commands:

Run:

```
$ sudo service hbase-rest start
```

Accessing the HBase service

To access the HBase service on a client machine, use the command hbase directly in shell:

```
$ hbase
```

Or you can use this command to enter the hbase console:

```
$ hbase shell
```

In the HBase shell, you can run some test commands, for example:

```
hbase(main):003:0> create 'test', 'cf'
hbase(main):003:0> list 'test'
hbase(main):004:0> put 'test', 'row1', 'cf:a', 'value1'
hbase(main):005:0> put 'test', 'row2', 'cf:b', 'value2'
hbase(main):006:0> put 'test', 'row3', 'cf:c', 'value3'
hbase(main):007:0> scan 'test'
hbase(main):008:0> get 'test',□ 'row1'
hbase(main):012:0> disable 'test'
hbase(main):013:0> drop 'test'
hbase(main):014:0> quit
```

Type help to get help for the HBase shell.

Stopping the HBase daemon

You can stop the daemons manually by using the following commands:

Run:

```
$ sudo service hbase-master stop
```

Stopping the HRegionServer daemon

You can stop the daemons manually by using the following commands:

Run:

```
$ sudo service hbase-regionserver stop
```

Stopping the Hbase Thrift server daemon [OPTIONAL]

You can stop the daemons manually by using the following commands:

Run:

\$ sudo service hbase-thrift stop

Stopping the Hbase Rest server daemon [OPTIONAL]

You can stop the daemons manually by using the following commands:

Run:

\$ sudo service hbase-rest stop

Hive

The base version of Hive is Apache Hive 0.11.0.

Hive is a data warehouse infrastructure that provides data summarization and ad hoc querying.

This section specifies how to install, configure, and use Hive.

Hive Components

A Hive installation consists of the following components:

- hive-server
- hive-metastore
- hive-dbserver

Prerequisites

As Hive is built on top of Hadoop, HBase and Zookeeper, the Hadoop, HBase and Zookeeper core packages must be installed for Hive to operate correctly.

The following prerequisites must be also met before installing Hive:

- PostgresSQL Server
- Hive Metastore backed by a DB Server.

ve/gphd/warehouse
hive.metastore.local = false

Hive RPM Packages

Hive consists of one core package and a thrift sever daemon package that provides Hive service through thrift.

hive-0.11.0_gphd_2_0_3_0-nn.noarch.rpm	
Туре	Core
Requires	Hadoop, HBase Core Packages
Description	Hive core package provides the executables, libraries, configuration files and documentations.
Install on Nodes	Hive Client workstation

hive-server-0.11.0_gphd_2_0_3_0-nn.noarch.rpm	
Туре	Daemon (thrift server)
Requires	Hive Core Package

hive-server-0.11.0_gphd_2_0_3_0-nn.noarch.rpm	
Description	Daemon scripts package to provide Hive service through thrift
Install on Nodes	Hive Thrift server node

hive-metastore-0.11.0_gphd_2_0_3_0-nn.noarch.rpm	
Туре	Deamon (Metastore server)
Requires	Hive Core Package
Description	Daemon scripts package to provide Hive metadata information through metastore server.
Install on Nodes	Hive Metastore server node

hive-server2-0.11.0_gphd_2_0_3_0-nn.noarch.rpm	
Туре	Daemon (hive server2)
Requires	Hive Core Package
Description	Daemon scripts package to provide Hive Server2.
Install on Nodes	Hive Thrift server node

Installing Hive

Set up PostgreSQL on the HIVE_METASTORE Node

- **1.** Choose one of the cluster nodes to be the HIVE METASTORE.
- **2.** Login to the nominated ${\tt HIVE_METASTORE}$ node as root.
- **3.** Execute the following commands yum install postgresql-server
- **4.** Initialize the database:

service postgresql initdb

5. Open the /var/lib/pgsql/data/postgresql.conf file and set the following values:

```
listen_addresses = '*'
standard conforming strings = off
```

6. Open the /var/lib/pgsql/data/pg_hba.conf file and comment out all the lines starting with host and local by adding # to start of the line.

Add following lines:

```
local all all trust host all all 0.0.0.0 0.0.0.0 trust
```

7. Create /etc/sysconfig/pgsql/postgresql file and add

```
PGPORT=10432
```

8. Start the database:

```
service postgresql start
```

9. Create the user, database.

```
sudo -u postgres createuser -U postgres -p 10432 -a hive sudo -u postgres createdb -U postgres -p 10432 metastore
```

Set up the HIVE_METASTORE

1. Install Hive metastore using:

```
yum install hive-metastore postgresql-jdbc
```

2. Open the /etc/gphd/hive/conf/hive-site.xml and change it to following:

```
<configuration>
cproperty>
   <name>javax.jdo.option.ConnectionPassword
   <value>hive</value>
</property>
cproperty>
   <name>hive.metastore.uris
   <value>thrift://<CHANGE TO</pre>
HIVE METASTORE ADDRESS>:9083</value>
</property>
cproperty>
   <name>javax.jdo.option.ConnectionURL</name>
<value>jdbc:postgresql://<CHANGE TO HIVE METASTORE ADDRESS>:
10432/metastore</value>
</property>
cproperty>
   <name>hive.metastore.warehouse.dir
   <value>/hive/gphd/warehouse</value>
</property>
cproperty>
   <name>hive.hwi.war.file
<value>/usr/lib/gphd/hive/lib/hive-hwi-0.9.1-gphd-2.0.1.0.wa
r</value>
</property>
cproperty>
    <name>javax.jdo.option.ConnectionDriverName
```

```
<value>org.postgresgl.Driver</value>
</property>
cproperty>
    <name>datanucleus.autoCreateSchema</name>
    <value>false</value>
</property>
cproperty>
    <name>hive.metastore.local
    <value>false</value>
</property>
cproperty>
    <name>javax.jdo.option.ConnectionUserName
    <value>hive</value>
</property>
cproperty>
    <name>hive.metastore.execute.setugi</name>
    <value>true</value>
</property>
</configuration>
```

Note: Replace < CHANGE TO HIVE METASTORE ADDRESS> in above file.

3. Create file /etc/gphd/hive/conf/hive-env.sh and add the following:

```
export HADOOP_HOME="/usr/lib/gphd/hadoop"
export HADOOP_CONF_DIR="/etc/gphd/hadoop/conf"
export HADOOP_MAPRED_HOME="/usr/lib/gphd/hadoop-mapreduce"
export HIVE_CONF_DIR="/etc/gphd/hive/conf"
Make it executable using:
chmod +x /etc/gphd/hive/conf/hive-env.sh
```

4. Edit file /etc/gphd/hadoop/conf/hadoop-env.sh and add the following before export HADOOP CLASSPATH:

```
export HIVELIB_HOME=$GPHD_HOME/hive/lib
export HADOOP_CLASSPATH=$HADOOP_CLASSPATH:\
$HIVELIB_HOME/hive-service-0.9.1-gphd-2.0.1.0.jar:\
$HIVELIB_HOME/libthrift-0.7.0.jar:\
$HIVELIB_HOME/hive-metastore-0.9.1-gphd-2.0.1.0.jar:\
$HIVELIB_HOME/libfb303-0.7.0.jar:\
$HIVELIB_HOME/hive-common-0.9.1-gphd-2.0.1.0.jar:\
$HIVELIB_HOME/hive-exec-0.9.1-gphd-2.0.1.0.jar:\
$HIVELIB_HOME/hive-exec-0.9.1-gphd-2.0.1.0.jar:\
$HIVELIB_HOME/hive-exec-0.9.1-gphd-2.0.1.0.jar:\
```

5. Link postgresql jar:

```
ln -s /usr/share/java/postgresql-jdbc.jar
/usr/lib/gphd/hive/lib/postgresql-jdbc.jar
```

6. Create the schema:

```
sudo -u postgres psql -U hive -d metastore -p 10432 -f /usr/lib/gphd/hive-0.9.1_gphd_2_0_2_0/scripts/metastore/upgr ade/postgres/hive-schema-0.9.0.postgres.sql
```

7. Start the hive-metastore:

```
service hive-metastore start
```

Note: MySQL is no longer supported. Please migrate from MySQL to PostgreSQL.

Hive Client Setup

Hive is a Hadoop client-side library. Install the Hive core package on the client workstation:

```
$ sudo rpm -ivh
working_dir/hive/rpm/hive-0.11.0_gphd_2_0_3_0-nn.noarch.rpm
```

Hive Thrift Server Setup [OPTIONAL]

Install the Hive core package and Hive thrift daemon package to provide Hive service through thrift.

```
$ sudo rpm -ivh
working_dir/hive/rpm/hive-0.11.0_gphd_2_0_3_0-nn.noarch.rpm
$ sudo rpm -ivh
working_dir/hive/rpm/hive-server-0.11.0_gphd_2_0_3_0-nn.noar
ch.rpm
```

Hive Server2 Setup [OPTIONAL]

Install the Hive core package and Hive thrift daemon package to provide Hive service through thrift.

```
$ sudo rpm -ivh
working_dir/hive/rpm/hive-0.11.0_gphd_2_0_3_0-nn.noarch.rpm
$ sudo rpm -ivh
working_dir/hive/rpm/hive-server2-0.11.0_gphd_2_0_3_0-nn.noa
rch.rpm
```

Hive MetaStore Server Setup [OPTIONAL]

Install the Hive core package and Hive Metastore daemon package to provide Hive metadata information through centralized Metastore service:

```
$ sudo rpm -ivh
working_dir/hive/rpm/hive-0.11.0_gphd_2_0_3_0-nn.noarch.rpm
$ sudo rpm -ivh
working_dir/hive/rpm/hive-metastore-0.11.0_gphd_2_0_3_0-nn.n
oarch.rpm
```

Hive Configuration

The configuration files for Hive are located here: /etc/gphd/hive/conf/

This is the default configuration for quick reference and modification. It is a symbolic link to /etc/gphd/hive-version/conf

You can make modifications to this configuration template or create your own. If you want to use a different configuration folder, adjust the symbolic link /etc/gphd/hive/conf to point to the folder you want to utilize at runtime.

Hive Post-installation Configuration

- **1.** Login to one of the cluster nodes as root.
- 2. Create the hive.warehouse.dir

```
$ sudo -u hdfs hadoop fs -mkdir -p /hive/gphd/warehouse
```

3. Set permissions for the hive.warehouse.dir

```
$ sudo -u hdfs hadoop fs -chmod 775 /hive/gphd/warehouse
```

4. Set the ownership for the hive.warehouse.dir

```
$ sudo -u hdfs hadoop fs -chown hadoop:hadoop
/hive/gphd/warehouse
```

5. Add hive user to hadoop group if not already present using

```
$ usermod -G hadoop hive
```

Hive Usage

Start Hive Client

To run Hive on a client machine, use the hive command directly in shell:

```
$ hive
```

You can check the Hive command usage by:

```
$ hive -help
```

Start Beeline Client

HiveServer2 supports a new command shell Beeline that works with HiveServer2:

```
$ beeline
```

Start/Stop Hive Thrift Server [Optional]

You can start/stop Hive thrift server daemon as follows:

Run:

```
$ sudo service hive-server start
$ sudo service hive-server stop
```

Start/Stop Hive Server2 [Optional]

You can start/stop Hive server2 daemon as follows:

Run:

```
$ sudo service hive-server2 start
$ sudo service hive-server2 stop
```

Start/Stop Hive Metastore Server [Optional]

You can start/stop Hive Metastore server daemon as follows:

Run:

```
$ sudo service hive-metastore start
$ sudo service hive-metastore stop
```

Configuring a Secure Hive Cluster

If you are running Hive in a standalone mode using a local or embedded MetaStore you do not need to make any modifications.

The Hive MetaStore supports Kerberos authentication for Thrift clients. Follow the instructions provided in the section about "Security" on page 91 to configure Hive for a security-enabled HD cluster.

Hcatalog

The base version of Heatalog is Apache Heatalog 0.11.0.

HCatalog is a metadata and table management system.

This section specifies how to install, configure, and use Heatalog.

Prerequisites

Heatalog is built on top of Hadoop, HBase, Hive and Zookeeper, so the Hadoop, HBase, Hive and Zookeeper core packages must be installed for Heatalog to operate correctly.

Hcatalog RPM Packages

Heatalog consists of one core package and a thrift sever daemon package that provides Hive service through thrift.

hcatalog-0.11.0_gphd_2_0_3_0-nn.noarch.rpm	
Туре	Core
Requires	Hadoop, HBase and Hive Core Packages.

Hcatalog 36

hcatalog-0.11.0_gphd_2_0_3_0-nn.noarch.rpm	
Description	Hcatalog core package provides the executables, libraries, configuration files and documentations.
Install on Nodes	Hcatalog Client workstation.

hcatalog-server-0.11.0_gphd_2_0_3_0-nn.noarch.rpm	
Туре	Daemon (hcatalog server).
Requires	Hcatalog Core Package.
Description	Daemon scripts package to provide Hive service through thrift.
Install on Nodes	Hcatalog server node.

webhcat-0.11.0_gphd_2_0_3_0-nn.noarch.rpm	
Туре	Libraries.
Requires	Hcatalog Core Package.
Description	Daemon scripts package to provide Hive metadata information through metastore server.
Install on Nodes	Webhcat server node.

webhcat-server-0.11.0_gphd_2_0_3_0-nn.noarch.rpm	
Туре	Daemon(webhcata server).
Requires	Hcatalog and Webhcat Core Package.
Description	Daemon scripts package to provide Webhcat Server.
Install on Nodes	Webhcat server node.

Hcatalog Client Setup

Heatalog is a Hadoop client-side library. Install the Heatalog core package on the client workstation.

```
$ sudo rpm -ivh
working_dir/hive/rpm/hcatalog-0.11.0_gphd_2_0_3_0-nn.noarch.rpm
```

Hcatalog Server Setup [OPTIONAL]

Install the Heatalog core package and Heatalog thrift daemon package to provide Heatalog service.

\$ sudo rpm -ivh

Hcatalog 37

```
working_dir/hcatalog/rpm/hcatalog-0.11.0_gphd_2_0_3_0-nn.noa
rch.rpm
$ sudo rpm -ivh
working_dir/hcatalog/rpm/hcatalog-server-0.11.0_gphd_2_0_3_0
-nn.noarch.rpm
```

Webhcat Setup [OPTIONAL]

Install the Hcatalog core package and Webhcat package to provide Webhcat libraries.

```
$ sudo rpm -ivh
working_dir/hcatalog/rpm/hcatalog-0.11.0_gphd_2_0_3_0-nn.noa
rch.rpm
$ sudo rpm -ivh
working_dir/hcatalog/rpm/webhcat-0.11.0_gphd_2_0_3_0-nn.noar
ch.rpm
```

Webhcat Server Setup [OPTIONAL]

Install the Heatalog core package and Hive Metastore daemon package to provide Hive metadata information through centralized Metastore service.

```
$ sudo rpm -ivh
working_dir/hcatalog/rpm/hcatalog-0.11.0_gphd_2_0_3_0-nn.noa
rch.rpm
$ sudo rpm -ivh
working_dir/hcatalog/rpm/webhcat-0.11.0_gphd_2_0_3_0-nn.noar
ch.rpm
$ sudo rpm -ivh
working_dir/hcatalog/rpm/webhcat-server-0.11.0_gphd_2_0_3_0-
nn.noarch.rpm
```

Hcatalog Configuration

The configuration files for Heatalog are located here: /etc/qphd/hive/conf/

This is the default configuration for quick reference and modification. It is a symbolic link to /etc/gphd/hive-version/conf

You can make modifications to this configuration template or create your own. If you want to use a different configuration folder, adjust the symbolic link /etc/gphd/hive/conf to point to the folder you want to utilize at runtime.

Usage

Start Heatalog Client

To run Heatalog on a client machine, use the hive command directly in shell:

\$ hcat

You can check the hive command usage by running:

```
$ hcat -help
```

Hcatalog 38

Start/Stop Hcatalog Server [Optional]

You can start/stop Heatalog server daemon as follows:

Either:

```
$ sudo service hcatalog-server start
$ sudo service hcatalog-server stop

or:
$ sudo /etc/init.d/hcatalog-server start
$ sudo /etc/init.d/hcatalog-server stop
```

Start/Stop Webhcat Server [Optional]

You can start/stop Webhcat server daemon as follows:

Either:

```
$ sudo service webhcat-server start
$ sudo service webhcat-server stop

or:
$ sudo /etc/init.d/webhcat-server start
$ sudo /etc/init.d/webhcat-server stop
```

Pig

The base version of Pig is Apache Pig 0.10.1.

Pig is a high-level data-flow language and execution framework for parallel computation.

This section specifies how to install, configure, and use Pig.

Prerequisites

As Pig is built on top of Hadoop the Hadoop package must be installed to run Pig correctly.

Pig RPM Packages

Pig has only one core package.

pig-0.10.1_gphd_2_0_3_0-nn.noarch.rpm	
Туре	Core
Requires	Hadoop Core Packages
Description	Pig core package provides executable, libraries, configuration files and documentation.
Install on Nodes	Pig client workstation

Pig 39

pig-doc-0.10.1_gphd_2_0_3_0-nn.noarch.rpm	
Туре	Documentation
Requires	N/A
Description	Pig documentation.
Install on Nodes	N/A

Pig Client Setup

Pig is a Hadoop client-side library. Install the Pig package on the client workstation:

```
$ sudo rpm -ivh
working dir/pig/rpm/pig-0.10.1 gphd 2 0 3 0-nn.noarch.rpm
```

Pig Configuration

The configuration files for Pig are located here: /etc/gphd/pig/conf/

This is the default configuration templates for quick reference and modification.

You can modify these configuration templates or create your own configuration set. If you want to use a different configuration folder, adjust the symbolic link conf under /etc/gphd/pig/ to point to the folder you want to utilize at runtime.

Usage

To run Pig scripts on a client machine, use the command pig directly in shell:

```
$ pig COMMAND
```

You can check the pig command usage by:

```
$ pig -help
```

Mahout

The base version of Mahout is Apache Mahout 0.7.

Mahout is a scalable machine learning and data mining library.

This section specifies how to install, configure, and use Mahout.

Prerequisites

Mahout is built on top of Hadoop, so the Hadoop package must be installed to get Mahout running.

Mahout RPM Packages

Mahout has only one core package.

Mahout 40

mahout-0.7_gphd_2_0_3_0-nn.noarch.rpm	
Туре	Core
Requires	Hadoop Core Packages
Description	Mahout core package provides executable, libraries, configuration files and documentations.
Install on Nodes	Mahout clie.nt workstation

Mahout Client Setup

Mahout is a Hadoop client-side library. Install the Mahout package on the client workstation:

```
$ sudo rpm -ivh
working dir/mahout/rpm/mahout-0.7 gphd 2 0 3 0-nn.noarch.rpm
```

Mahout Configuration

You can find the configuration files for Mahout in the following location:

```
/etc/gphd/mahout/conf/
```

This is the default configuration templates for quick reference and modification.

You can modify these configuration templates or create your own configuration set. If you want to use a different configuration folder, adjust the symbolic link conf under /etc/gphd/mahout/ to point to the folder you want to utilize at runtime.

Usage

To run Mahout scripts on a client machine, use the command mahout directly in shell:

```
$ mahout PROGRAM
```

You can check the full list of mahout programs by running:

\$ mahout

Flume

The base version of Flume is Apache Flume 1.3.1.

Flume is a distributed, reliable, and available service for efficiently collecting, aggregating, and moving large amounts of log data. It has a simple and flexible architecture based on streaming data flows. It is robust and fault tolerant with tunable reliability mechanisms and many failover and recovery mechanisms. It uses a simple extensible data model that allows for online analytic application. For more info, please refer to the Apache Flume page: http://flume.apache.org/

This section specifies how to install, configure, and use Flume.

Flume 41

Prerequisites

As Flume is built on top of Hadoop, the Hadoop package must be installed to get Flume running correctly.

(Hadoop core and hadoop hdfs should be installed)

Flume RPM Packages

Flume consists of one core package and a flume-agent sever daemon package.

flume-1.3.1_gphd_2_0_3_0-nn.noarch.rpm	
Туре	Core
Requires	Hadoop Core Packages
Description	Flume core package provides executable, libraries, configuration files and documentations.
Install on Nodes	Flume client workstation.

flume-agent-1.3.1_gphd_2_0_3_0-nn.noarch.rpm	
Туре	Daemon (Flume Agent server)
Requires	Flume core Package
Description	Daemon scripts package to provide Flume service for generating, processing, and delivering data.
Install on Nodes	Flume agent server node.

Flume Client Setup

Flume is a Hadoop client-side library. Install the Flume package on the client workstation: Installation can be done using yum or rpm -ivh options.

```
$ sudo rpm -ivh
working dir/flume/rpm/flume-1.3.1 gphd 2 0 3 0-nn.noarch.rpm
```

Note: User flume and group flume should be created with correct configuration, including uid, gid, home_dir and shell. Check in following paths: /etc/passwd, /etc/group

Flume Agent Setup [Optional]

Install the Flume core package and Flume agent daemon package to provide Flume service for generating, processing, and delivering data:

```
$ sudo rpm -ivh
working_dir/flume/rpm/flume-1.3.1_gphd_2_0_3_0-nn.noarch.rpm
$ sudo rpm -ivh
```

Flume 42

```
working_dir/flume/rpm/flume-agent-1.3.1_gphd_2_0_3_0-nn.noar
ch.rpm
```

Flume Configuration

The configuration files for Flume are located here: /etc/gphd/flume/conf/

This is the default configuration templates for quick reference and modification.

You can modify these configuration templates or create your own configuration set. If you want to use a different configuration folder, adjust the symbolic link conf under /etc/qphd/flume/ to point to the folder you want to utilize at runtime.

Usage

Starting Flume Client

To run Flume scripts on a client machine, use the command flume-ng directly in shell:

```
$ flume-ng
```

You can check the flume-ng command usage by running:

```
$ flume-ng --help
```

Starting/Stopping Flume Agent Server [Optional]

You can start/stop Flume agent server daemon as follows:

Run:

```
$ sudo service flume-agent start
$ sudo service flume-agent stop
$ sudo service flume-agent status
```

Sqoop

The base version of Sqoop is Apache Sqoop 1.4.2.

Sqoop is a tool designed for efficiently transferring bulk data between Apache Hadoop and structured datastores such as relational databases. For more details, please refer to the Apache Sqoop page: http://sqoop.apache.org/

This section specifies how to install, configure, and use Sqoop.

Prerequisites

As Sqoop is built on top of Hadoop and HBase, the Hadoop and HBase package must be installed to get Flume running correctly.

Sqoop RPM Packages

Flume consists of one core package and a sqoop-metastore sever daemon package.

Sqoop 43

sqoop-1.4.2_gphd_2_0_3_0-nn.noarch.rpm	
Туре	Core
Requires	Hadoop, HBase Core Packages
Description	Sqoop core package provides executable, libraries, configuration files and documentations.
Install on Nodes	Sqoop. client workstation

sqoop-metastore-1.4.2_gphd_2_0_3_0-nn.noarch.rpm		
Туре	pe Daemon (Sqoop Metastore server)	
Requires	Sqoop core Package	
Description	Daemon scripts package to provide shared metadata repository for Sqoop.	
Install on Nodes	Sqoop metastore server node	

Sqoop Client Setup

Sqoop is a Hadoop client-side library. Install the Sqoop package on the client workstation:

```
$ sudo rpm -ivh
working_dir/sqoop/rpm/sqoop-1.4.2_gphd_2_0_3_0-nn.noarch.rpm
```

Note: User sqoop and group sqoop should be created with correct configuration: uid sqoop, gid sqoop, homedir /home/sqoop, Shell /sbin/nologin. Check in following path: /etc/passwd and /etc/group.

Sqoop Metastore Setup [Optional]

Install the Sqoop core package and Sqoop agent daemon package to provide shared metadata repository for Sqoop. sqoop-metastore has the dependency with sqoop-core package:

```
$ sudo rpm -ivh
working_dir/sqoop/rpm/sqoop-1.4.2_gphd_2_0_3_0-nn.noarch.rpm
$ sudo rpm -ivh
working_dir/sqoop/rpm/sqoop-metastore-1.4.2_gphd_2_0_3_0-nn.
noarch.rpm
```

Sqoop Configuration

The configuration files for Flume are located here: /etc/gphd/sgoop/conf/

This is the default configuration templates for quick reference and modification.

Sqoop 44

You can modify these configuration templates or create your own configuration set. If you want to use a different configuration folder, adjust the symbolic link conf under /etc/gphd/sqoop/ to point to the folder you want to utilize at runtime.

Usage

Starting Sqoop Client

To run Sqoop scripts on a client machine, use the command sqoop directly in shell:

```
$ sqoop
```

You can check the sqoop command usage by running:

```
$ sqoop help
```

Starting/Stopping Sqoop Metastore Server [Optional]

You can start/stop Sqoop metastore server daemon as follows:

Run:

```
$ sudo service sqoop-metastore start
$ sudo service sqoop-metastore stop
$ sudo service sqoop-metastore status
```

Spring Data

Spring for Apache Hadoop provides support for writing Apache Hadoop applications that benefit from the features of Spring, Spring Batch, and Spring Integration. For more information, please refer to the Spring Data official page:

```
http://www.springsource.org/spring-data/hadoop
```

Installing Spring Data Hadoop

1. Download and copy Pivotal HD Tools Tarball to /home/gpadmin/. Make sure the tarball has read permission for user gpadmin. To extract the PHDTools tarball execute the following command:

```
[root@hdp2-w17 gpadmin]# chown gpadmin:gpadmin
PHDTools-version.tar.gz
[root@hdp2-w17 gpadmin]# ls -lrt PHDTools-version.tar.gz
-rw-r--r-- 1 gpadmin gpadmin 499930679 Mar 20 20:12
PHDTools-version.tar.gz

[root@hdp2-w17 gpadmin]# sudo su - gpadmin

[gpadmin@hdp2-w17 ~]$ tar xzvf PHDTools-version.tar.gz
[gpadmin@hdp2-w17 ~]$ ls -lrt GPHD*
drwxrwxr-x 5 gpadmin gpadmin 4096 Mar 20 00:35
PHDTools-version
-rw-r--r-- 1 gpadmin gpadmin 499930679 Mar 20 20:12
```

Spring Data 45

```
PHDTools-version.tar.gz

[gpadmin@hdp2-w17 ~]$ cd

PHDTools-version/spring-data-hadoop/rpm/
[gpadmin@hdp2-w17 rpm]$ ls -lrt

total 1580

-rw-rw-r-- 1 gpadmin gpadmin 1610604 Mar 20 00:04

spring-data-hadoop-1.0.1.RC1-3.noarch.rpm

-rw-rw-r-- 1 gpadmin gpadmin 76 Mar 20 00:44

spring-data-hadoop-1.0.1.RC1-3.noarch.rpm.md5
```

Installing Spring Data Hadoop through RPM

To install Spring Data Hadoop through RPM execute the following command:

Spring Data Hadoop

By default, Spring Data Hadoop is installed to /usr/local/gphd/ directory.

The following documentation is installed:

```
[gpadmin@hdp2-w17 ~]$ cd
/usr/local/gphd/spring-data-hadoop-1.0.1.RC1
[gpadmin@hdp2-w17 spring-data-hadoop-1.0.1.RC1]$ ls -lrt
total 36
-rw-r--r- 1 root root 861 Oct 11 01:32 readme.txt
-rw-r--r- 1 root root 11357 Oct 11 01:32 license.txt
-rw-r--r- 1 root root 1151 Mar 4 06:19 notice.txt
drwxr-xr-x 2 root root 4096 Mar 20 20:49 dist
drwxr-xr-x 4 root root 4096 Mar 20 20:49 docs
drwxr-xr-x 3 root root 4096 Mar 20 20:49 schema
drwxr-xr-x 2 root root 4096 Mar 20 20:49 samples
```

Please refer to the readme.txt and files within docs/ directory to start using Spring Data Hadoop.

Spring Data 46

HVE (Hadoop Virtualization Extensions)

Hadoop was first designed for typically running on commodity hardware and native OS. Hadoop Virtualization Extensions were developed to enhance Hadoop running in cloud and virtualization environments

Topology Awareness

Hadoop Virtualization Extensions (HVE) allow Hadoop clusters implemented on virtualized infrastructure full awareness of the topology on which they are running, thus enhancing the reliability and performance of these clusters.

HVE should be enabled in the following situations:

- When there is more than one Hadoop VM per physical host in virtualized environments
- When Datanodes and TaskTrackers exist in separate virtual machines in virtualized environments, so as to achieve graceful scaling of the compute component of the Hadoop cluster.
- When there is a topology layer between host and rack (e.g. chassis), which can affect the failure/locality group between hosts, in non-virtualized environments.

Topology Awareness Configuration and Verification Sample Setup

This setup has 2 logical racks, 2 physical hosts (install ESXi and managed by vCenter) per rack, and 2 DN/CN (VM in ESXi) nodes per host. There is also one namenode/jobtracker and a client node that can be used to start jobs.

In this setup, each DN/CN node has 4 vCPUs, 16G memory and 200G (Non-SSD) disks.

The namenode and jobtracker are installed on another dedicated VM with 4vCPU, 4G Memory and 100G disks.

Node Distribution on Hosts

Rack 1	Host 1	Namenode and Jobtracker	DN1
	Host 2	DN2	NN3
Rack 2	Host 3	DN4	DN5
	Host 4	DN6	DN7

Enable topology awareness (Hadoop V2)

1. Add the following line to core-site.xml:

```
<name>topology.script.file.name</name>
     <value>/hadoop/hadoop-smoke/etc/hadoop/topology.sh</value
     <!-- point to topology.sh location.-->
```

```
</property>
   cproperty>
      <name>net.topology.impl</name>
      <value>org.apache.hadoop.net.NetworkTopologyWithNodeGroup
      </value>
      <description> The default implementation of
      NetworkTopology which is classic three layer one.
      </description>
   </property>
   cproperty>
      <name>net.topology.nodegroup.aware
      <value>true</value>
      <description> By default, network topology is not aware
      of nodegroup layer.
      </description>
   </property>
   cproperty>
      <name>dfs.block.replicator.classname</name>
      <value>org.apache.hadoop.hdfs.server.blockmanagement.Bloc
      kPlacementPolicyWithNodeGroup</value>
      <description> The default implementation of
      ReplicationTargetChooser.
      </description>
   </property>
2. Add the following line to yarn-site.xml:
   cproperty>
      <description>The class to use as scheduled
      requests.</description>
      <name>yarn.resourcemanager.scheduled.requests.class/name>
      <value>org.apache.hadoop.mapreduce.v2.app.rm.ScheduledReques
      tsWithNodeGroup</value>
   </property>
   cproperty>
      <description> The boolean value to identify if the cluster
      is deployed on an environment which needs an additional
      layer (node group) between node and rack for network
      topology.
      </description>
      <name>net.topology.with.nodegroup</name>
```

```
<value>true</value>
   </property>
   cproperty>
      <description>The class to use as
      AbstractSchedulerElementsFactory in RM
      scheduler.</description>
      <name>yarn.resourcemanager.scheduler.elements.factory.impl/
      name>
      <value>org.apache.hadoop.yarn.server.resourcemanager.schedul
      er.SchedulerElementsFactoryWithNodeGroup</value>
   </property>
Topology.data sample
   [root@namenode enable]# cat topology.data
   10.111.57.223 (VM IP) /Rack1/NodeGroup1
   10.111.57.224 /Rack1/NodeGroup1
   10.111.57.225 /Rack1/NodeGroup2
   10.111.57.226 /Rack2/NodeGroup1
   10.111.57.227 /Rack2/NodeGroup1
   10.111.57.228 /Rack2/NodeGroup2
   10.111.57.229 /Rack2/NodeGroup2
Topology.sh sample:
   [root@namenode enable]# cat topology.sh
   #! /bin/bash
   HADOOP CONF=/hadoop/hadoop-smoke/etc/hadoop
   # this is the location of topology.data
   while [ $# -gt 0 ] ; do
     nodeArg=$1
     exec< ${HADOOP CONF}/topology.data
     result=""
     while read line ; do
       ar=( $line )
       if [ "\${ar[0]}" = "\$nodeArg" ] ; then
         result="${ar[1]}"
       fi
     done
     shift
     if [ -z "$result" ] ; then
       echo -n "/default/rack "
     else
       echo -n "$result "
     fi
```

done

3. Verify HVE is enabled:

Run the TestDFSIO script.

The output is as follows:

```
1) HVE enabled:
Job Counters
Launched map tasks=100
Launched reduce tasks=1
Data-local map tasks=26
NODEGROUP_LOCAL_MAPS=49
Rack-local map tasks=25
2) HVE disabled:
Job Counters
Launched map tasks=100
Launched reduce tasks=1
Data-local map tasks=20
Rack-local map tasks=80
```

Elasticity

HVE Elastic Resource Extension enables the adaption of MapReduce tasks to changing resources on nodes/clusters where Hadoop clusters are deployed on virtualization by sharing resource with VMs from other clusters or applications.

Overview

Currently, the Hadoop resource model is static at the node level, assuming the node resources are not changed while the cluster is running. This design and implementation are based on an assumption that all cluster resources are dedicated for Hadoop MapReduce jobs, so they are fully available at all times. This assumption does not hold when users want to deploy multiple applications on the same cluster, e.g. deploying HBase and MapReduce on the same HDFS cluster. In particular, in an era of cloud computing, it is common for Hadoop clusters to be deployed on virtualization by sharing resource with VMs from other clusters or applications.

The HVE elastic resource feature addresses scenarios in which nodes' resources are possibly changed, so that scheduling of MapReduce tasks on these nodes can adapted to changing resources, as represented in the figure below.

With this feature, APIs (CLI and JMX interface) and script tools are provided to get/set map and reduce task slots on Hadoop cluster nodes for MR jobs.

Function List

Below are functionalities included in this elastic feature:

Table 3.1 Function List

Function	Description
Configuration	Enable/disable elastic resource feature on Hadoop cluster by specifying a configuration property when starting MR cluster.
List nodes' status	List the status of all the nodes or nodes specified by user, including its tracker_name, hostname, health status, slot number, etc.
Set map/reduce slot on specific node	Set map/reduce slot number to a node specified by user via CLI or JMX interface.
Set map/reduce slot on nodes in batch mode	Set map/reduce slot number on nodes specified by a node-slot mapping file.

Configuration

Elastic resource feature is disabled by default. To enable elastic resource, make the following changes to the Hadoop configuration.

In mapred-site.xml, add the following property to enable the elastic resource feature:

Also, specifying the node's map and reduce slot number in mapred-site.xml will also be the upper limit of node's slot resource. Any slot number update, if greater than this value, will be replaced with the value specified here. In following example, the maximum map slot number is 5 and maximum reduce slot number is 2.

Map:

Command Line Interface

List all CLIs of MRAdmin

```
hadoop mradmin
    [-refreshServiceAcl]
    [-refreshQueues]
    [-refreshUserToGroupsMappings]
    [-refreshSuperUserGroupsConfiguration]
    [-refreshNodes]
    [-listNodes]
    [-nodeStatus <tracker name>]
    [-setSlot <tracker name> <map | reduce> <slot number>]
    [-setSlotByHost <hosname> <map|reduce> <slot number>]
    [-help [cmd]]
```

List Nodes

hadoop mradmin -listNodes

Get Node Status

hadoop mradmin -nodeStatus <TRACKER-NAME>

Set Slot Number with Tracker_Name

Map:

hadoop mradmin -setSlot <TRACKER_NAME> map <NUM>

Reduce:

hadoop mradmin -setSlot <TRACKER NAME> reduce <NUM>

Set Slot Number with Hostname

Map:

hadoop mradmin -setSlotByHost <HOSTNAME> map <NUM>

Reduce:

hadoop mradmin -setSlotByHost <HOSTNAME> reduce <NUM>

Updating Slots in Batch mode

Batch Updating Slots Script

Location: \${HADOOP HOME}/bin/update-slots.sh

The script will apply the slot updating operation in batch mode on nodes by resolving a slot description file.

Slot Description File

Location of this file is by default as below:

```
${HADOOP HOME}/conf/slots
```

Also, you can also specify the location of this file in running this script tool:

```
Usage: update-slots.sh PATH TO SLOT FILE
```

The slot description file format is as following:

```
hostname1 map_slots reduce_slots
hostname2 map_slots reduce_slots
hostname3 map_slots reduce_slots
...
```

Below is an example:

```
hadoop@Hadoop:~ $ cat ${HADOOP_HOME}/conf/slots
Hadoop-01 2 2
Hadoop-02 1 1
```

Example

The setup of cluster with 3 nodes is as following:

Table 3.2 Cluster Setup

HostName	Role
Namenode	JobTracker / Namenode
Hadoop-01	TaskTracker / Datanode
Hadoop-02	TaskTracker / Datanode

1. CLI

a. listNodes:

```
hadoop@NameNode:~$ hadoop mradmin -listNodes

TaskTracker
Health UsedMap MaxMap UsedReduce MaxReduce

tracker_Hadoop-02:localhost.localdomain/127.0.0.1:35722

OK 0 2 0 2

tracker_Hadoop-01:localhost.localdomain/127.0.0.1:48516

OK 0 1 0 1
```

b. nodeStatus:

```
hadoop@NameNode:~$ hadoop mradmin -nodeStatus
tracker_Hadoop-01:localhost.localdomain/127.0.0.1:48516

{
    "failures" : 0,
    "slots" : {
        "map_slots" : 1,
        "reduce_slots_used" : 0,
        "map_slots_used" : 0
    },
    "tracker_name" :
"tracker_Hadoop-01:localhost.localdomain/127.0.0.1:48516",
    "hostname" : "Hadoop-01",
    "health" : "OK",
```

```
"last seen" : 1359570978699
hadoop@NameNode:~$ hadoop mradmin -nodeStatus
tracker Hadoop-02:localhost.localdomain/127.0.0.1:35722
  "failures" : 0,
  "slots" : {
    "map slots" : 2,
    "reduce slots" : 2,
    "reduce slots used" : 0,
    "map slots used" : 0
  },
  "tracker name" :
"tracker Hadoop-02:localhost.localdomain/127.0.0.1:35722",
  "hostname": "Hadoop-02",
  "health" : "OK",
  "last seen" : 1359570988442
}
c. setSlot (map):
hadoop@NameNode:~$ hadoop mradmin -setSlot
tracker Hadoop-02:localhost.localdomain/127.0.0.1:35722 map
Set map slot of
tracker Hadoop-02:localhost.localdomain/127.0.0.1:35722
successfully.
hadoop@NameNode:~$ hadoop mradmin -listNodes
TaskTracker
Health UsedMap MaxMap UsedReduce MaxReduce
tracker Hadoop-02:localhost.localdomain/127.0.0.1:35722
tracker Hadoop-01:localhost.localdomain/127.0.0.1:48516
d. setSlot (reduce):
hadoop@NameNode:~$ hadoop mradmin -setSlot
tracker Hadoop-02:localhost.localdomain/127.0.0.1:35722
reduce 1
Set reduce slot of
tracker Hadoop-02:localhost.localdomain/127.0.0.1:35722
successfully.
hadoop@NameNode:~$ hadoop mradmin -listNodes
TaskTracker
Health UsedMap MaxMap UsedReduce MaxReduce
tracker Hadoop-02:localhost.localdomain/127.0.0.1:35722
                                                            1
```

```
tracker Hadoop-01:localhost.localdomain/127.0.0.1:48516
  OK 0
                            1
                                           0
                                                               1
  e. setSlotByHost (map):
  hadoop@NameNode:~$ hadoop mradmin -setSlotByHost Hadoop-01
  map 1
   Set map slot of
   tracker Hadoop-01:localhost.localdomain/127.0.0.1:48516
  successfully.
  hadoop@NameNode:~$ hadoop mradmin -listNodes
  TaskTracker
  Health UsedMap MaxMap UsedReduce MaxReduce
  tracker Hadoop-02:localhost.localdomain/127.0.0.1:35722
                                                                2
                             2
   tracker Hadoop-01:localhost.localdomain/127.0.0.1:48516
                                                                2
              \cap
                             1
                                            \cap
  f. setSlotByHost (reduce):
  hadoop@NameNode:~$ hadoop mradmin -setSlotByHost Hadoop-01
  reduce 1
  Set reduce slot of
  tracker Hadoop-01:localhost.localdomain/127.0.0.1:48516
  successfully.
  hadoop@NameNode:~$ hadoop mradmin -listNodes
  TaskTracker
  Health UsedMap MaxMap UsedReduce MaxReduce
  tracker Hadoop-02:localhost.localdomain/127.0.0.1:35722
                              2
   tracker Hadoop-01:localhost.localdomain/127.0.0.1:48516
  OK
               \cap
                              1
                                                                1
2. Batch Updating Slots
   a. Use the slot description file by default:
  hadoop@NameNode:~/hadoop-1.2.0$ bin/update-slots.sh
  Set map slot of
  tracker Hadoop-01:localhost.localdomain/127.0.0.1:48516
  successfully.
  Set reduce slot of
  tracker Hadoop-01:localhost.localdomain/127.0.0.1:48516
  successfully.
  Set map slot of
  tracker Hadoop-02:localhost.localdomain/127.0.0.1:35722
  successfully.
  Set reduce slot of
  tracker Hadoop-02:localhost.localdomain/127.0.0.1:35722
  successfully.
  hadoop@NameNode:~/hadoop-1.2.0$ hadoop mradmin -listNodes
```

```
TaskTracker

Health UsedMap MaxMap UsedReduce MaxReduce

tracker_Hadoop-02:localhost.localdomain/127.0.0.1:35722

OK 0 1 0 1

tracker_Hadoop-01:localhost.localdomain/127.0.0.1:48516
```

b. Specify the slot description file:

```
hadoop@NameNode:~/hadoop-1.2.0$ bin/update-slots.sh
slot desc file
Set map slot of
tracker Hadoop-01:localhost.localdomain/127.0.0.1:48516
successfully.
Set reduce slot of
tracker Hadoop-01:localhost.localdomain/127.0.0.1:48516
successfully.
Set map slot of
tracker Hadoop-02:localhost.localdomain/127.0.0.1:35722
successfully.
Set reduce slot of
tracker Hadoop-02:localhost.localdomain/127.0.0.1:35722
successfully.
hadoop@NameNode:~/hadoop-1.2.0$ hadoop mradmin -listNodes
TaskTracker
Health UsedMap MaxMap UsedReduce MaxReduce
tracker Hadoop-02:localhost.localdomain/127.0.0.1:35722
                   1
tracker Hadoop-01:localhost.localdomain/127.0.0.1:48516
                  1
                                                          1
```

HDFS Rack Awareness

HDFS rack awareness is a key feature to achieve localized I/O (locality).

With respect to read and write separately, HDFS has:

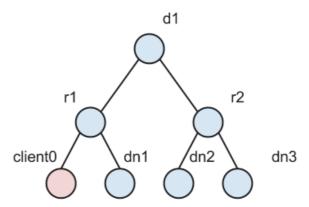
- BlockPlacementPolicy for write locality: namenode will look up network topology and construct a list of chosen nodes (pipeline) for a requesting a block to locate, based on algorithms provided by a BlockPlacementPolicy.
- Block pseudo distance sort for read locality: when reading a block, after obtaining all the located blocks, namenode sorts these located blocks based on their topological distance with client. The closer nodes get higher priority for read.

Both operations need to reference network topology, which is managed by the rack awareness feature. The rack awareness feature includes:

- A topology resolving framework: when datanodes register themselves on a namenode, that namenode will resolve their network location using their host name or ip, using DNSToSwitchMapping. This is a pluggable component that allows users to define their own topology based on their network layout. The most commonly used DNSToSwitchMapping is ScriptBasedMapping, which calls a shell script.
- An in-memory topology tree: all registered datanodes' network locations are kept in a topology tree.

Problem: Ignored off-cluster clients

The problem of the current implementation is that it do not support off-cluster clients. The figure below is an example of off-cluster clients:



In that figure, node **dn1** is a datanode and its network location is /d1/r1, and so on for **dn2** and **dn3**. Node **client0** is an off-cluster node, which means there is no datanode deployed on **client0**. In this case, **client0** has no chance to register itself in the topology tree of the namenode. Therefore both read and write operations select random nodes even though **dn1** is closer (more preferable) than either **dn2** or **dn3**.

This problem will cause performance issues in the following cases:

- When a mapreduce cluster is not exactly co-located: some mapreduce clusters share the same hdfs cluster with other mapreduce clusters, or in some cases a mapreduce cluster will cover several hdfs clusters. In those cases, a big portion of I/O will be off-cluster client operations which can not benefit from localized I/O.
- When a physical cluster is not dedicated to hadoop: a physical cluster may not be
 dedicated to hadoop and other supporting systems such as data loading tools may
 share the same cluster. In that case, the data loading tool can not benefit from
 localized I/O, even if the tool and hdfs shares the same rack/data center. The
 problem may even more common in virtualized environment.

Solution: Design

To tackle this problem, we changed the logic in block placement policy and block pseudo distance sort. We also resolved the network location of the client.

Resolving client location

Resolving the client location: we reused the framework that resolves datanodes. However, since we did not add client network locations into topology tree (as explained below), we have to cache client locations to avoid unnecessary resolve operations.

As a result, we introduced two LRU caches:

- A black list for those clients who have no valid location or whose locations do not share the same rack with any datanode.
- A white list opposite to the black list, for those clients who are not datanodes but share the same rack with at least one datanode.

Referring to the diagram of ignored off-cluster clients, the table below lists some examples of location cache.

Table 3.3 Location Cache Examples

HostName	Location	Cache
client1	d1/r1	white list
client2	d2/r1	black list
client4	null	black lis

The size of LRU cache is configurable so you can limit the memory usage of namenode.

Block placement policy

The tables below demonstrate how the BlockPlacementPolicy has been changed to support non-datanode clients.

Table 3.4 Former block placement algorithm

Replica	Rule
1	Client Local
2	Random node whose rack is different from replica 1
3	Random node who share the same rack with replica 2
>=4	Random node

New Block Placement Algorithm:

Table 3.5 Changed block placement algorithm

Replica	Rule
1	Client Local if client is datanode, or random node who shares the same rack with client if client is not a datanode
2	Random node whose rack is different from replica 1

Table 3.5 Changed block placement algorithm

Replica	Rule
3	Random node who shares the same rack with replica 2
>=4	Random node

Usage

The client rack aware feature is disabled by default. To enable, add the following to the hdfs-site.xml file:

```
properties>
   cproperty>
      <name>dfs.rackawareness.with.client
      <value>true</value>
   </property>
</properties>
cproperties>
   cproperty>
      <name>dfs.rackawareness.with.client.blacklist.size/name>
      <description>Black list size of client cache, 5000 by
      default.</description>
      <value>5000</value>
   </property>
</properties>
cproperties>
   cproperty>
      <name>dfs.rackawareness.with.client.cache.size
      <description>White list size of client cache, best set it
      equals the size of cluster. 2000 by default.</description>
      <value>2000</value>
   </property>
</properties>
```

Note that you need to restart DFS after changing the configuration.

Vaidya

- Vaidya Overview
- Installing Vaidya Files
- Enabling and Disabling Vaidya
- Using Vaidya to Analyze Jobs
- Vaidya Configuration Rules

Vaidya Overview

Vaidya is a diagnostic tool installed with PHD for Map/Reduce jobs. After a job is executed successfully, it uses a job history log and job configuration information to identify any performance or scalability problems with the job. Upon execution, it provides a job analysis report indicating specific problems with the job along with the remedy to correct them. The report element includes, "rule title", "rule description", "rule importance", "rule severity", "reference details" and "remedy/prescription" to rectify the problem. The "rule severity", is a product of rule impact and the rule importance.

Note: The Vaidya tool does *not* analyze failed jobs either for performance or scalability problems nor for the reasons of failure.

The Vaidya tool includes diagnostic rules (also referred to as "tests") where each rule analyzes a specific problem with the M/R job. Diagnostic rule is written as a Java class and captures the logic of how to detect a specific problem condition with the M/R job. Each diagnostic rule takes job history log and job configuration information provided to it using a standard structured interface. The standard interface allows administrators and developers to independently add more diagnostic rules in the Vaidya tool.

Note that Vaidya is installed together with PHD and is by default enabled. No additional installation and configuration needed.

Note that Vaidya is not available if you are deploying a MR1-based cluster.

Installing Vaidya Files

By default, Vaidya files are installed at:

- The Vaidya JAR library is installed into /usr/lib/qphd/hadoop-mapreduce/
- The Vaidya default test configuration file is installed into /etc/qphd/hadoop/conf/

Enabling and Disabling Vaidya

By default, Vaidya is enabled after installation, there is normally no need to enable it manually.

Enabling Vaidya

In cases where Vaidya is not enabled and you want to enable it explicitly:

On the job tracker node, go to the PHD configuration folder (by default, /etc/gphd/hadoop/conf), and add the following lines into the file mapred-site.xml.

mapred-site.xml

Disabling Vaidya

To disable Vaidya:

Set the property mapreduce.vaidya.enabled value to be false, or remove these lines from mapred-site.xml.

Note:

- The value of property mapreduce.vaidya.enabled should be changed to point to the correct jar file you installed. By default, this is /usr/lib/gphd/hadoop/contrib/vaidya/hadoop-vaidya-<hadroop_PHD_VERSION>.jar or /usr/lib/gphd/hadoop-mapreduce/hadoop-vaidya-<hadroop_PHD_VERSION>.jar.
- Once you edit the xml file, restart the job history server service to ensure the change takes effect.

Using Vaidya to Analyze Jobs

To use Vaidya with PHD 1.0.x and to ensure your job history server service is running:

- **1.** Successfully run a map-reduce job for Vaidya to analyze. Refer to *Pivotal HD Enterprise 1.0 Installation and Administrator Guide* for instructions about how to run map-reduce job with PHD.
- **2.** Ensure your job history server service is running.

3. Open the following URL in a web browser:

http://<historyserver_host>:<historyserver_port>/jobhistory
Where:

- < historyserver_host > refers to the host name or IP address of the machine where you run job history server service.
- < historyserver_port > refers to the HTTP port job history server web where the UI listens. By default, this value is 19888. Your browser should show you the job history server UI page.
- **4.** You will see a list of jobs that have run, including the most recent job. Click the job id of any job in this list, and you should see the detailed information of the job.
- **5.** On the left side of the navigation area, there should be a link called **Vaidya report** under the navigation item **Job**. Click the **Vaidya report** link and Vaidya will analyze the job for you and show a report.

Vaidya Configuration Rules

After you installed Vaidya with PHD, rules configuration is installed as a postex diagnosis tests.xml XML file, here: /etc/gphd/hadoop/conf/

You can find all rules to be run on a selected job in that XML file, where each rule is defined as an XML PostExPerformanceDiagnosisTests/DiagnosticTest element, for example:

</DiagnosticTest>

The Title and Description elements provide a brief summary about what this rule is doing.

By editing postex diagnosis tests.xml, you can configure the rules.

Notes:

- Remember to backup original configuration file before editing the configuration file, invalid xml config file may cause Vaidya behavior incorrectly.
- Before you start editing rules, you should have background knowledge about XML syntax and how XML represents data (for example, what the CDATA element represents).

Disabling a Rule

Comment out or remove the whole DiagnosticTest element.

Changing the Importance of a Rule

Importance indicates how relatively important a rule is, relative to other rules in the same set. You can change the importance value by editing Importance element in the XML file. A level served as a factor which is multiplied to impact value returned by each rule.

There are three values valid for this attribute: Low, Medium and High; their corresponding values are: 0.33, 0.66 and 0.99.

In the displayed Vaidya report, there is a value named Severity for each rule. A severity level is the result of multiplying the impact value (returned by rule) and the importance factor (defined in XML file).

For example, a rule returns impact of 0.5, its importance is marked as Medium, then its severity is 0.5 * 0.66 = 0.33.

Changing Success Threshold

Each rule calculates a value between 0 and 1 (inclusively) to indicate how healthy a job is according to the given rule, this value is called impact. The smaller the impact is (that is, closer to 0), the healthier the job is.

To give a more straight forward result, you can set a threshold for each rule, therefore a rule whose impact value is larger than the threshold will be marked as "failed", otherwise, it is marked as "passed".

Note that threshold is compared with impact value, rather than severity (which means make a rule less important will not make a failed rule succeed).

You can change the threshold value by editing the SuccessThreshold element in the XML file.

Changing Input Parameters

Some rules may need additional input parameters to complete their logic. You can specify additional parameters by editing/adding elements under the InputElement element of each rule.

Other

For a full explanation and instruction about the meaning of each XML element and how to change them, refer the Apache's Official Vaidya Guide for more information.

Adding a New Rule

A Vaidya rule consists of the following two parts:

- A java class that consists of the logic of the rule
- A paragraph of XML in the configuration file

Creating a Java Binary for a New Rule

Important: This section assumes a working knowledge of how to write, compile, and package Java code.

- 1. From where you installed PHD, download the correct hadoop-vaidya-<HADOOP_PHD_VERSION>.jar file (which you specified in mapred-site.xml) to your development machine, if you plan on writing Java code on another machine than the one where you installed PHD (This is a typical case).
- **2.** Create a java file with an IDE or editor, which defines a class that extends the org.apache.hadoop.vaidya.DiagnosticTest class:

myrule.java

```
package com.greenplum.vaidya.rules;
import org.apache.hadoop.vaidya.DiagnosticTest;
import org.apache.hadoop.vaidya.statistics.job.JobStatistics;
public class MyRule extends DiagnosticTest {
    @Override
    public String getReferenceDetails() {
        return "";
    }
    @Override
    public String getPrescription() {
        return "";
    }
    @Override
    public double evaluate(JobStatistics jobStatistics) {
        return 0.5;
    }
}
```

- **3.** Edit the three methods getReferenceDetails, getPrescription and evaluate to construct the logic.
 - evaluate method should return a **double** value between 0.0 and 1.0 and represents the impact as the analysis result.

- getPrescription method should return some text providing user suggestions/remedies about how to optimize your Map/Reduce configuration accordingly.
- getReferenceDetails method should return some text indicating the meaningful counters and their values which can help you to diagnose your Map/Reduce configuration accordingly.
- **4.** Compile the java class and package compiled class to a jar file, e.g., myrule.jar. Note that you need to put the Vaidya jar file you just downloaded into your class path to make your code compile.

Creating XML Configuration For a New Rule

Add a DiagnosticTest element into the postex_diagnosis_tests.xml file (the file you set in mapred-site.xml file), according to the sample given in the configuration part. Ensure the value of ClassName element is set to be the full class name of the java rule class you just created.

Deploying files

- 1. Upload the packaged jar file (myrule.jar for example) to the node where you installed PHD job tracker, and store it in a folder where hadoop service has the permission to read and load it. We recommend you place it under /usr/lib/gphd/hadoop-mapreduce/lib/
- 2. Edit mapred-site.xml, append the jar file you just uploaded to the mapred.vaidya.jar.file or mapreduce.vaidya.jarfiles property value, for example:

Important:

- Do not remove the default Vaidya jar file from this property, Vaidya needs this property to load basic Vaidya classes to make it run.
- Multiple jar files are separated by different separator characters on different platforms. On the Linux/Unix platform, the ":" character should be used. You can look at theFile.pathSeparator attribute of your java platform to ensure it.
- **3.** To make your settings take affect, restart job history server service.
- **4.** Run your Vaidya analysis again, you should see your rules run and results appear in the report.

DataLoader

See the *Greenplum DataLoader Installation and User Guide* for detailed information.

DataLoader 66

4. Unified Storage System (USS)

This chapter installation, configuration, usage, and troubleshooting about the USS service.

Overview

USS is a service on Pivotal HD that provides a unified namespace view of data across multiple filesystems.

USS enables users to access data across multiple filesystems, without copying the data to and from HDFS. It is the underlying technology for enabling Data Tiering/Retention and Migration from one Hadoop distribution to another.

USS is implemented as a 'pseudo Hadoop File System' that delegates File System operations directed at it to other filesystems in a HDFS-like way. It mounts multiple filesystems and maintains a centralized view of the mount points, which are accessible through the URI scheme of uss://. It relies on a catalog service for managing metdata about mount points and delegated filesystems.

Versions

This section contains information about USS version 0.4.0, which is compatible with Pivotal HD 1.0.1 and above (Hadoop version 2.0.x).

Supported Features

USSv 0.4.0 supports the following features:

Table 4.1 Supported Features

Feature	Description	
Compatibility	Pivotal HD 1.0.1and above (Hadoop 2.0.x)	
Supported Delegate File Systems	HDFS 2.x Isilon (OneFS) LocalFS (Mount point target paths assumed to be present on all datanodes/nodemanagers) NFS (Mounted on all datanodes/nodemanagers) FTP (with known issues)	
Modes of access/interfaces	 HDFS CLI Native Java Map-Reduce Jobs Hadoop Streaming Pig Hive (External tables) 	
Installation/Configuration	Fully managed through Pivotal HD Enterprise 1.0.1 above.	
Catalog	Postgres database	
Security	USS 0.4.0 adds support for secure hadoop (using kerberos). See USS Secure Configuration.	

Overview 67

Architecture

Unified Storage System Client hdfs cli MR pig hive streaming uss://.. Hadoop FileSystem Client USS USS USS USS URI USS Client Protocol Resolution Namenode Catalog Unified Storage System Layer hdfs://Namenode1... hdfs://lsi NN.. hdfs://Namenode2.. file:///... NFS HDFS 1 HDFS 2 Isilon (Locally (Isi_NN) (Namenode1) (Namenode2) Mounted) **Delegate Filesystems**

USS is implemented as a 'pseudo Hadoop FileSystem' that:

- Accepts Hadoop File System requests containing USS URIs
- Extracts **USS Mount Points** from USS URIs
- Resolves USS Mount Points to their Actual Paths on supported Delegate File Systems
- Delegates the File System requests to the Delegate File Systems

Refer to "Terminology" on page 89 for a description of these terms.

Overview 68

Prerequisites

- Platform
 - RHEL 6.1/6.2 64Bit or CentOS 6.1/6.2 64Bit
- Pivotal HD Hadoop 1.0.1 or above.
 - Pivotal HD is available through via the EMC Download Center (paid version) or the Pivotal HD product page (Community Edition). Please contact your local Pivotal HD Support if you need help downloading Pivotal HD.
- For USS to support a storage system, it must either be a Hadoop File System, or it must have a Hadoop File System compatible implementation.
- Installation/Setup/Configuration/Maintenance of delegate file systems is beyond the purview of USS. USS assumes that these file systems are available for serving requests.
- Oracle Java 1.6 JDK needs to be installed in all USS nodes in the cluster. Use RPM install to setup necessary paths for JDK on the machine. Download and execute self-executing RPM binary (for example: jdk-6u43-linux-x64-rpm.bin) from the Oracle site http://www.oracle.com/technetwork/java/javase/downloads/jdk6downloads-1902814.html after accepting the license.

Getting Started

USS (v0.4.0) is part of Pivotal HD 1.0.1 and above. USS is distributed as a set of three RPM packages that are part of the PHDTools-version tarball. You can download the PHDTools tarball from the EMC Download Center (paid version) or the Pivotal HD product page (Community Edition).

Architecturally, USS has the following four components:

USS Namenode (Server)

The USS Namenode is a stateless server-side component. It is implemented as a Hadoop RPC server, that is capable of interacting with a catalog system to delegate file system calls to appropriate underlying file systems. It is expected to run on a dedicated server. In this alpha version, this runs as a single server, with no HA and failover. It is designed to be the central place where all file system calls converge, before they are delegated to a multitude of underlying file systems. It abstracts datasets on underlying file systems as mount-points. It stores the mapping between mount-points and the directories that they point to in the USS catalog. The USS Namenode does not store any state, its state is contained within the USS Catalog.

USS Catalog

The USS Catalog is a metadata store for the USS Namenode. In this version, the USS catalog is contained in a postgres database. USS ships scripts that can help users manage the USS Catalog. The USS RPMs do not install the postgres database to be

Getting Started 69

used as the USS catalog. It is expected that users have a postgres database ready. However, if users use Pivotal Command Center to deploy USS, it can deploy a postgres database as the USS Catalog.

USS Agent

The USS Agent is expected to be present on all nodes in a Pivotal HD cluster. This is the USS client-side library that accepts a FileSystem request (containing a mount-point), maps the request to the appropriate File System and path, by making an RPC request to the USS Namenode.

USS Client

The USS Client hosts allow administrators to update the USS Catalog using the USS CLI. This is an optional role. If not configured, users can use the USS Namenode or the USS Catalog host as the client node.

USS Directory Contents

The contents of USS 0.4.0 within the PHDTools-version tarball are located here:

```
PHDTools-version/uss/
|--rpm
|--uss-0.4.0-5.noarch.rpm
|--uss-0.4.0-5.noarch.rpm.md5
|--uss_catalog-0.4.0-5.noarch.rpm
|--uss_catalog-0.4.0-5.noarch.rpm.md5
|--uss_namenode-0.4.0-5.noarch.rpm
|--uss_namenode-0.4.0-5.noarch.rpm.md5
```

USS RPM Packages

USS is distributed as a set of three RPM packages:

Package Descriptions:

 Table 4.2
 rpm package descriptions

PRPM Package Name	Description	Dependencies	Hosts to install on
uss-0.4.0-5.noarch.rpm	The base USS package. It installs the USS library and required configuration files.	Pivotal HD hadoop package	All hosts using USS (agent, namenode, client and catalog hosts)
uss_catalog-0.4.0-5.no arch.rpm	This package installs the scripts and configuration files required to interact with and administer the USS Catalog.	uss (version >= 0.4.0)	USS Client hosts. Installed on the USS Namenode host by default as the uss_namenode RPM package has a dependency on it.
uss_namenode-0.4.0-5.n oarch.rpm	This is the USS Namenode package. It installs the USS namenode service daemon and its runtime configuration.	uss_catalog package (version >= 0.4.0).	USS Namenode host

USS RPM Packages 71

Package Contents uss-0.4.0-5.noarch.rpm

Table 4.3 uss-0.4.0-5.noarch.rpm

File	Description
/usr/lib/gphd/uss/uss-0.4.0.jar	The USS library
/usr/lib/gphd/uss/uss-javadoc-0.4.0.jar	The USS javadocs
/etc/gphd/uss/conf/commons-logging.proper ties and /etc/gphd/uss/log4j.properties	USS log configuration files
/etc/gphd/uss/conf/uss-client-site.xml	USS client side configuration file
/etc/gphd/uss/conf/uss-nn-site.xml	USS Namenode configuration file
/etc/gphd/uss/conf/uss-catalog.xml	USS Catalog configuration file
/etc/gphd/uss/conf/uss-env.sh	USS Client side environment setup script
/usr/lib/gphd/uss/docs/README.txt	The USS installation, configuration and usage document.
/usr/lib/gphd/uss/docs/CHANGES.txt	The USS changelog.
/usr/lib/gphd/uss/docs/NOTICE.txt	The USS notice document.
/usr/lib/gphd/uss/docs/README.txt	The USS license document.

uss_catalog-0.4.0-5.noarch.rpm

Table 4.4 uss_catalog-0.4.0-5.noarch.rpm

File	Description
/usr/bin/uss/	This is the USS CLI to help admins/users to interact with the USS catalog.
/usr/lib/gphd/uss/migration	Contains utilities for migrating the USS catalog database.
/usr/lib/gphd/uss/lib	Contains the java libraries required to access the USS catalog.
/usr/lib/gphd/uss/libexec	Contains USS utility scripts.
/etc/default/uss	USS Catalog and Namenode runtime configuration script.

uss_namenode-0.4.0-5.noarch.rpm

Table 4.5 uss_namenode-0.4.0-5.noarch.rpm

File Description	
/usr/sbin/uss-namenode-daemon	The USS Namenode daemon script.
/etc/rc.d/init.d/uss-namenode	USS Namenode service script.

USS RPM Packages 72

Table 4.5 uss namenode-0.4.0-5.noarch.rpm

File	Description
/var/log/gphd/uss	USS Namenode log directory.
/var/log//gphd/uss-namenode	Log directory for USS Namenode

Installing USS

With Pivotal HD 1.0.1 and above users can install USS on their clusters using Pivotal HD Manager. This is the recommended way to install USS. For more information on installing USS using Pivotal HD Manager, please refer to the *Pivotal HD Manager 1.0 Installation and Administrator Guide* for details.

Configuring USS

Where:

The following sections describe the USS configurations in Pivotal HD 1.0.1 and above. Pivotal HD Manager sets up all these configurations and no user action is necessary.

Enabling USS on your Pivotal HD Cluster

To enable USS, the following properties are set in /etc/gphd/hadoop/conf/core-site.xml.

Installing USS 73

- fs.uss.impl—USS is implemented as a pseudo hadoop file system. USS URI's use the scheme uss. Hadoop compatible file systems set the property fs.<scheme>.impl to the class that implements the file system. The class that implements the USS pseudo file system is com.greenplum.hd.uss.USSProtocolTranslator.
- mapred.outdir.resolverClass—If the output directory of a mapreduce job is set to a USS URI, hadoop uses this property to refer to a class that is capable of resolving the USS URI to a 'real' URI on an underlying file system. USS provides the class com.greenplum.hd.uss.USSPathResolver to resolve USS URIs

USS Namenode Configuration

These are the server-side configurations for USS.

```
uss-nn-site.xml
<?xml version="1.0"?>
<?xml-stylesheet type="text/xsl" href="configuration.xsl"?>
<configuration>
   property>
      <name>uss.namenode.address
      <value>uss://USS NAMENODE HOST:16040</value>
      <description>Uri for USS Namenode.</description>
   </property>
   cproperty>
      <name>uss.catalog.type</name>
      <value>db</value>
      <description>Type of catalog to use.</description>
   </property>
   property>
      <name>uss.db.name</name>
      <value>dbApplicationContext.xml</value>
      <description>Settings for postgres catalog.</description>
   </property>
</configuration>
```

Where:

- uss.namenode.address—The fully-qualified URI of the USS Namenode. Specifies the hostname and port on which the USS Namenode listens for requests.
- uss.catalog.type—Type of catalog to use. Defaults to db. We do not recommend changing this property.

Configuring USS 74

• uss.db.name—Settings for postgres catalog. Defaults to dbApplicationContext.xml. We do not recommend changing this property.

USS Client Configuration

These are the client-side configurations for USS.

</configuration>

Where:

- fs.uss.name—This property specifies the hostname and port of the USS Namenode on the client side. By default, it is set to uss://localhost:16040/. If the client is a non-USS Namenode host, then this property can be used to specify the hostname and port of the USS Namenode on the client.
- uss.tmp.fs—This property indicates the location that USS uses to store temporary data. All Hadoop-compatible file systems require a location to store temporary. Hadoop uses this location for creating the staging area for map-reduce jobs (.staging directory) as well as to store the .Trash directory. As USS is implemented as a pseudo Hadoop file system, it also requires a location for storing temporary data. However, as there is no storage/file system tied to USS, USS stores temporary data on the file system specified against this parameter. Users could choose to set this to a filesystem of their choice, as long as the file system is compatible with Hadoop (i.e. it has an HDFS implementation available). This file system should be available for USS to store temporary data. It could either be one of the underlying file systems, or a dedicated file system for storing temporary data.

Configuring USS 75

USS Catalog Configuration

The USS Namenode as well as the USS CLI uses the configuration file /etc/gphd/uss/conf/uss.properties to read connection properties for the USS Catalog. If you wish to modify the USS Catalog Database URI, please update the property uss.db.url from this file. The value for this property is of the form:

```
uss.properties - uss.db.url
uss.db.url=jdbc:postgresql://[USS_CATALOG_DB_HOST]:[USS_CATA
LOG DB PORT]/usscatalog
```

We recommend that users only make sure that the host and port of the postgres server designated to be the USS Catalog is correct. Do not modify any other property from this file.

USS Catalog Administration - Using the USS CLI

USS provides a Command Line Interface for administrators to manage the USS catalog from the USS Client hosts. The purpose of this tool is to provide a central/consistent way to manage the USS Catalog. This section describes the usage of the USS Admin CLI. Currently, USS Catalog resides in a postgres database and contains mainly USS Mount Point information. It is envisioned that the catalog will be enriched further.

Basics

- The *USS Catalog Management Tool* is accessed via the command: uss. This is a script located at /usr/bin/uss
- All nodes part of the USS Client role have this command available to them. It is a part of the uss catalog RPM package.
- These nodes need to have network reachability to the catalog (node on which the postgres database is installed)

Usage:

General Operations

Operations which do not inherently modify the catalog and can be used freely by all cluster users.

List existing mount points

Displays Mount Point Name, Mount Point ID and Mount Point Target URI for all mount points currently defined in the catalog.

```
$ uss list
click_data (1) > hdfs://NAMENODE_HOST1:8020/data/website/clicks
conv_data (2) > hdfs://NAMENODE_HOST2:8020/data/website/conversions
view data (3) > file:///data/website/views
```

USS Classpath

General utility to view the java classpath used by USS.

USS Admin Operations

This encompasses the admin level operations that actually modify the catalog. It is encouraged that these operations be restricter to just a few admin users.

- Initialize USS Catalog
- Add a New Mount Point
- Delete an Existing Mount Point

USS Admin Usage

```
$ uss admin --help
usage: uss admin
-a, --add
                                Add a new mount point to the USS
                                catalog
-d,--delete
                                Delete an existing mount point from
                                the USS catalog
-h,--help
                                Print help for this application
-i,--initialize
                                Setup the database for use with USS
-1,--list
                                List all mount points defined in
                                the USS catalog
-n,--mount-point-name <arg> Name of the mount point to be added
                                or deleted
-t,--mount-point-target <arg> Target (full URI) of the mount
                                point to be added
```

Initialize USS Catalog

• Create the usscatalog database on the host which has the postgres installation to use as the USS Catalog Database

- Create the usscatalog database owner
- Create the usscatalog database DDL
- Add default (as well as sample) data to the database
- Note: USS uses flyway to manage schema versioning and migration.
- Note: Pivotal HD Manager initializes the USS catalog database using this command automatically during cluster deployment.

```
$ uss admin --initialize
```

Add a New Mount Point

Accepts the mount point name and mount point target URI as user input and adds a new mountpoint to the catalog

```
$ uss admin --add --mount-point-name NewMountPoint
--mount-point-target hdfs://NAMENODE1:8020/newPath
$ uss list
click_data (1) > hdfs://NAMENODE1:8020/data/website/clicks
conv_data (2) > hdfs://NAMENODE2:8020/data/website/conversions
view_data (3) > file:///data/website/views
NewMountPoint (5) > hdfs://NAMENODE1:8020/newPath
```

Delete an Existing Mount Point

Accepts the mount point name as user input and deletes the mount point from the catalog (if it exists).

```
$ uss admin --delete --mount-point-name NewMountPoint
```

```
$ uss list
click_data (1) > hdfs://namenode1.mycompany.com:8020/data/website/clicks
conv_data (2) > hdfs://namenode2.mycompany.com:8020/data/website/conversions
view data (3) > file:///data/website/views
```

Using USS

Setting Up the Environment

Once USS is configured, users can access it by using the Hadoop Filesystem CLI or Map-reduce jobs. However, clients also need to set some environment variables prior to using USS. These variables make the USS library available in the classpath for hadoop commands. Pivotal HD Manager adds these variables to

/etc/gphd/hadoop/conf/hadoop-env.sh by default, so users don't have to add them manually. These environment settings are listed below:

USS Environment Settings

```
$ export
HADOOP_CLASSPATH=/usr/lib/gphd/uss/uss-0.3.0.jar:/etc/gphd/uss/conf/
$ export HADOOP USER CLASSPATH FIRST=true
```

The above environment settings have also been captured in the client side environment setup script /etc/gphd/uss/conf/uss-env.sh. Users may also edit this file and execute it before using USS.

Using USS Paths in map-reduce jobs

If users wish to use USS paths in map-reduce jobs, then they need one more extra configuration step. When the nodemanagers try to access USS paths, they need the USS library to resolve these paths. As a result, users need to add the same settings mentioned above in the /etc/gphd/hadoop/conf/hadoop-env.sh script and then restart the nodemanager on all nodes. To do this, add the following lines to /etc/gphd/hadoop/conf/hadoop-env.sh on all nodemanager nodes -

```
/etc/gphd/hadoop/conf/hadoop-env.sh
export HADOOP_CLASSPATH=/usr/lib/gphd/uss/uss-0.2.0.jar:/etc/gphd/uss/conf/
export HADOOP USER CLASSPATH FIRST=true
```

The yarn configuration file also needs to be updated, as follows:

In /etc/gphd/hadoop/conf/yarn-site.xml update the
yarn.application.classpath property to include the following classpath entries:
 \$USS_HOME
 \$USS_CONF

Testing USS

Execute hadoop filesystem commands (using the Hadoop CLI) or mapreduce jobs to test

To test USS, users can execute some hadoop filesystem CLI commands, mapreduce jobs, hadoop streaming jobs, pig scripts or hive scripts using mount-points defined in the catalog.

USS URIS

A USS URI is a URI that USS can understand and resolve. It does not point to a 'real' location on a FileSystem, but contains an index into a Catalog that can be used to lookup for the actual URI on an underlying 'Delegate FileSystem'. A USS URI has the form:

uss://<USSNameNode Host>:<USSNameNode Port>/<Mount-Point>/<Sub-path under
the mount-point>

Examples using USS URIs in Hadoop Filesystem CLI

```
# Prepare input directory
hadoop fs -mkdir /tmp/test uss input
```

```
# Copy data from /etc/passwd
hadoop fs -copyFromLocal /etc/passwd /tmp/test uss input
# Add mount-point for input. Do this on the
uss-admin/uss-namenode/uss-catalog host
$ uss admin --add --mount-point-name input mount point --mount-point-target
hdfs://NAMENODE HOST:8020/tmp/test uss input
# Add mount-point for output. Do this on the
uss-admin/uss-namenode/uss-catalog host
$ uss admin --add --mount-point-name output mount point
--mount-point-target hdfs://NAMENODE HOST:8020/tmp/test uss output
# List contents of the input directory using USS
$ hadoop fs -ls uss://USS NAMENODE:16040/input mount point
Found 1 items
-rw-r--r-- 1 user wheel
                               1366 2012-08-17 10:08
hdfs://NAMENODE HOST:8020/tmp/test_uss_input/passwd
```

Examples using USS URIs in Mapreduce jobs

```
$ hadoop jar
/usr/lib/gphd/hadoop-mapreduce/hadoop-mapreduce-examples-2.0.2-alpha-gphd-2
.0.1.0.jar wordcount uss://USS_NAMENODE_HOST:16040/input_mount_point
uss://USS_NAMENODE_HOST:16040/output_mount_point/uss_mr_output
```

Examples using USS URIs in Pig Scripts

```
\# pig script
$ cat uss.pig
A = load 'uss://USS_NAMENODE_HOST:16040/input_mount_point';
B = foreach A generate flatten(TOKENIZE((chararray)$0, ':')) as word;
C = filter B by word matches '
w+';
D = group C by word;
E = foreach D generate COUNT(C), group;
STORE E into
'uss://USS_NAMENODE_HOST:16040/output_mount_point/uss_pig_output';
$ pig -Dpig.additional.jars=/usr/lib/gphd/uss/uss-0.3.0.jar uss.pig //
Execute pig script, by adding the uss jar as an additional jar to include.
```

Examples using USS URIs in Hive Scripts

```
$ cat uss-hive.sql
```

```
-- creates an external table with location pointed to by a USS URI.
DROP TABLE test uss external;
CREATE EXTERNAL TABLE test uss external (testcol1 STRING, testcol2 STRING)
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ':'
LINES TERMINATED BY '\n'
STORED AS TEXTFILE
LOCATION 'uss://USS NAMENODE HOST:16040/input mount point';
SELECT * FROM test uss external;
$ hive -f uss-hive.sql
            Examples using USS URIs in Streaming jobs
\# streaming mapper
$ cat uss-streaming-mapper.py
#!/usr/bin/env python
import sys
# input comes from STDIN (standard input)
for line in sys.stdin:
    # remove leading and trailing whitespace
    line = line.strip()
    # split the line into words
    words = line.split()
    # increase counters
    for word in words:
        # write the results to STDOUT (standard output);
        # what we output here will be the input for the
        # Reduce step, i.e. the input for USSStreamingReducer.py
        #
```

```
print '%s\t%s' % (word, 1)
\# streaming reducer
$ cat uss-streaming-reducer.py
#!/usr/bin/env python
from operator import itemgetter
import sys
current word = None
current count = 0
word = None
# input comes from STDIN
for line in sys.stdin:
    # remove leading and trailing whitespace
    line = line.strip()
    # parse the input we got from mapper.py
    word, count = line.split('\t', 1)
    # convert count (currently a string) to int
    try:
        count = int(count)
    except ValueError:
        # count was not a number, so silently ignore/discard this line
        continue
    # this IF-switch only works because Hadoop sorts map output by key
    # (here: word) before it is passed to the reducer
    if current word == word:
        current count += count
    else:
        if current word:
        # write result to STDOUT
```

USS Namenode Service

USS exposes the USS Namenode service. This service listens for Hadoop RPC requests from clients at a specific port (configurable via uss.namenode.address in uss-nn-site.xml on the USS Namenode host). This service can be started using service script /etc/rc.d/init.d/uss-namenode.

Starting the service

\$ service uss-namenode start

Stopping the service

\$ service uss-namenode stop

Restarting the service

service uss-namenode restart

Checking the service status

service uss-namenode status

This service internally uses the USS daemon script installed at /usr/sbin/uss-namenode-daemon to perform the desired function.

USS Namenode Service 83

It uses the following runtime files:

- /var/run/uss-namenode.pid-USS Namenode service pid file.
- /var/lock/subsys/uss-namenode-USS Namenode service lock file.
- /var/log//gphd/uss-namenode/uss-namenode.log-USS Namenode log file

Uninstallation

USS can be uninstalled by erasing the USS RPM packages from the hosts on which it was installed.

```
$ yum erase uss_namenode -y
$ yum erase uss -y
```

Known limitations/Future Enhancements

USS Namenode High Availability

In this release, we have not implemented high availability for the USS Namenode. High Availability will be implemented in a future release.

USS across multiple HDFS versions

In this release, USS works with Pivotal HD 2.0.x HDFS. In a future release, we will implement support for other versions/implementations of HDFS like 1.x and MapR.

Security

The USS Namenode currently does not implement security. It re-uses file system permissioning/access control from the underlying file systems. In a future release, there may be a security layer in USS.

Catalog Service Integration

As described earlier in this document, currently, USS supports a postgres catalog. In a subsequent release, the catalog will reside in a central catalog service. The USS Namenode will communicate with the catalog service to resolve mount points.

Performance Testing/Benchmarking

For the alpha release, there has been no performance testing and benchmarking of USS. This will happen in a subsequent release and results will be published once available.

Limited Hive Support

In its alpha release, USS has limited support for hive as a higher level language. Hive is supported for the following scenario -

Uninstallation 84

- The Hive warehouse directory is located on HDFS.
- External tables can be created with locations specified by USS URIs.

Support for more scenarios may be added in a subsequent release.

Troubleshooting/FAQs

- I am trying to access a file system through USS. However, I keep getting errors like ipc.Client: Retrying connect to server: <USS Namenode Hostname>/<USS Namenode ip address>:<USS Namenode port>.
 Already tried 1 time(s).'
 - Please check if the USS Namenode is running on the host and port specified in configuration against the 'uss.namenode.address' property.
- I am trying to access a file system through USS. However, I keep getting errors like No FileSystem for scheme: uss'
 - Make sure fs.uss.impl is set to com.greenplum.hd.uss.USSProtocolTranslator in /etc/gphd/hadoop/conf/core-site.xml
- Where do I specify the address of the USS Namenode and why?
 - The address of the USS Namenode can be specified on client and USS namenode hosts.

Client Side:On the client host, the USS Namenode address can be specified by the property uss.fs.name in the file /etc/gphd/uss/conf/uss-client-site.xml. By default, on the client host, the USS Namenode address is set to localhost:16040. If users have set up the USS Namenode on a different host:port location, they can use this property to set the address of the USS Namenode

USS Namenode Side: The USS Namenode starts at localhost:16040 by default. If clients wish to change the port, they can specify the new port by setting the property uss.namenode.address on the USS Namenode host at the location /etc/gphd/uss/conf/uss-nn-site.xml.

- How do I configure USS as the default filesystem?
 - The cluster wide default file system can be over ridden to use USS by default. This implies that all files default to the file system URI scheme of uss:// as opposed to hdfs://. This can be done by setting fs.default.name in /etc/gphd/hadoop/conf/core-site.xml to uss://uss-namenode-host:uss-namenode-port on all nodes in the cluster (datanodes, nodemanager, client hosts). The advantage of setting this in configuration is that you can use the USS mount-point only, to access a USS Path, as opposed to using the fully qualified USS URI. An example of this is:

```
$ hadoop fs -ls /fs-test-mp
Found 1 items
-rw-r--r- 3 root supergroup 199 2012-10-15 21:20
/user/testuser/fstest/README.txt
```

When using USS as the default FileSystem fpr mapreduce jobs, however, you need to set <code>fs.defaultFS</code> to <code>uss://<uss_namenode>:<uss_port> in the /etc/gphd/hadoop/conf/core-site.xml on all nodemanager hosts. Also add /etc/gphd/uss/conf/uss-client-site.xml into

HADOOP_CLASSPATH in /etc/gphd/hadoop/confhadoop-env.sh (or /etc/default/hadoop)</code>

• Why do I see the exception below when I run a mapreduce job using USS Paths for input/output directories?

```
INFO mapred.JobClient: Task Id : attempt 201210151759 0001 m 000004 0,
Status : FAILED
java.lang.RuntimeException: java.lang.ClassNotFoundException:
com.greenplum.hd.uss.USSProtocolTranslator
org.apache.hadoop.conf.Configuration.getClass(Configuration.java:867)
org.apache.hadoop.fs.FileSystem.createFileSystem(FileSystem.java:1380)
    at org.apache.hadoop.fs.FileSystem.access$200(FileSystem.java:66)
    at org.apache.hadoop.fs.FileSystem$Cache.get(FileSystem.java:1404)
    at org.apache.hadoop.fs.FileSystem.get(FileSystem.java:254)
    at org.apache.hadoop.fs.FileSystem.get(FileSystem.java:123)
    at org.apache.hadoop.mapred.Child$4.run(Child.java:254)
    at java.security.AccessController.doPrivileged(Native Method)
    at javax.security.auth.Subject.doAs(Subject.java:416)
org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.j
ava:1121)
    at org.apache.hadoop.mapred.Child.main(Child.java:249)
Caused by: java.lang.ClassNotFoundException:
com.greenplum.hd.uss.USSProtocolTranslator
    at java.net.URLClassLoader$1.run(URLClassLoader.java:217)
    at java.security.AccessController.doPrivileged(Native Method)
    at java.net.URLClassLoader.findClass(URLClassLoader.java:205)
    at java.lang.ClassLoader.loadClass(ClassLoader.java:321)
    at sun.misc.Launcher$AppClassLoader.loadClass(Launcher.java:294)
    at java.lang.ClassLoader.loadClass(ClassLoader.java:266)
```

```
at java.lang.Class.forName0(Native Method)
at java.lang.Class.forName(Class.java:264)
at
org.apache.hadoop.conf.Configuration.getClassByName(Configuration.java:820)
at
org.apache.hadoop.conf.Configuration.getClass(Configuration.java:865)
... 10 more
```

• This is because the nodemanagers cannot find the USS library to resolve USS Paths at runtime. To fix this, please add the following lines to /etc/gphd/hadoop/conf/hadoop-env.sh and restart nodemanager on all nodes using service hadoop-yarn-nodemanager restart.

```
/etc/gphd/hadoop/conf/hadoop-env.sh
export HADOOP_CLASSPATH=/usr/lib/gphd/uss/uss-0.2.0.jar:/etc/gphd/uss/conf/
export HADOOP USER CLASSPATH FIRST=true
```

- How do I list all the mount-points defined in the catalog?
 - USS provides an API to list all the mount-points that they point to. Users can list the mount-points by running a hadoop fs -ls query on the path '/' on any datanode or hadoop client node.

Additionally, on the USS client host, administrators can also use- uss list to list all the mount points and their targets.

```
$ uss list
click_data (1) >
hdfs://namenode1.mycompany.com:8020/data/website/clicks
conv_data (2) >
hdfs://namenode2.mycompany.com:8020/data/website/conversions
view_data (3) > file:///data/website/views
```

• What hadoop services are dependent on USS? In other words, which services need the USS library in their classpath at runtime?

• In order to use USS for map-reduce, NodeManagers need to have the USS library in their classpath.

To add the USS library to the classpath, which can be done by <code>source uss-env.sh(or set HADOOP_CLASSPATH</code> to the path of the USS jar, and set <code>HADOOP_USER_CLASSPATH_FIRST=true</code>). However, when you export these settings in the environment and run <code>service hadoop-yarn-nodemanager restart</code>, the <code>/etc/init.d/hadoop-yarn-nodemanager script does a sudo <code>-u mapred hadoop-daemon.sh</code>. Because of the <code>sudo -u mapred</code>, a new shell gets spawned, and the environment variables do not propagate. As a result, even if you set these variables in your shell before starting the nodemanager, the nodemanager still does not have the USS jar in the classpath.</code>

Instead, list the 2 settings here in

/etc/gphd/hadoop/conf/hadoop-env.sh on all nodemanager nodes. This script is sourced in the bin/hadoop script, which is the last script executed while starting services (as a result, these environment settings will be present in the same shell).

- If a map-reduce job does file-system CRUD operations for exmple, it does a FileSystem.open() on a USS Path, then does any other service require USS?
 - In this case, the datanodes may need to have the uss.jar library available at runtime as well. Use the same technique described in the above answer to add the USS jar to the classpath of the datanode processes.
- If a USS configuration parameter value changes, what services need to be restarted for the new setting to take effect?
 - As of now, no USS setting is used by any of the hadoop services.

 /etc/gphd/uss/conf/uss-client-site.xml contains settings that are entirely used by the hadoop filesystem client.

 /etc/gphd/uss/conf/uss-nn-site.xml contains settings that are entirely used by the USS Namenode. Currently, we do not have any settings that may require us to restart hadoop services like namenode, secondary namenode, datanode, resourcemanager or nodemanager.
- I have configured USS to access HDFS running on a independent set of hosts, whose namenode is identified by the URI hdfs://hdfs-nn.mydomain.com:8020/. However, when I execute some Hadoop CLI commands using USS paths, I keep seeing Connection Refused errors. The errors seem to indicate that a connection attempt is being made to an HDFS Namenode on a different host/ip from my HDFS namenode host.
 - Please check the setting uss.tmp.fs in /etc/gphd/uss/conf/uss-client-site.xml. This setting is described in '"Configuring USS" on page 73 section in this document. Ensure that it is set to a location that is available to serve Hadoop Filesystem requests. USS uses it to store temporary data.

- When I run a pig script using USS, I see errors like ClassNotFoundException: com.greenplum.hd.uss.USSProtocolTranslator
 - Pig needs to include the USS library. You can achieve this by running your pig script as

```
$. pig -Dpig.additional.jars=/usr/lib/gphd/uss/uss-0.2.0.jar
<pig script using uss paths>
```

- Why do I get an error relation "mount_point" does not exist when I try to add a mount-point using uss admin -a?
 - This error indicates that the table mount_point does not exist in the uss-catalog database. Please check that you have initialized the uss-catalog database using uss admin --initialize.

Terminology

• USS URI—A USS URI is a URI that USS can understand and resolve. It does not point to a 'real' location on a FileSystem, but contains an entry into a Catalog that can be used to lookup for the actual URI on an underlying 'Delegate FileSystem'. A USS URI has the form:

```
uss://<USSNameNode Host>:<USSNameNode Port>/<Mount-Point>/<Sub-path under the mount-point>
```

- Mount Point—USS URI's are of the form: uss://[Authority]/[Mount Point]/[Subdirectory]. The mount point is an index in a 'mount-table' like Catalog system that USSNameNode can access. This Catalog could reside in Configuration or it could be part of the Centralized Catalog Service at a later stage. Given a USS URI, it is the responsibility of the USSNameNode to lookup the Catalog Service for the mount-point and return the 'Resolved URI' on a 'Delegate FileSystem'.
- Catalog—The Catalog is the metadata store that the USS Namenode uses to look up mount points and resolve USS URIs to their actual URIs on Delegate File Systems.
- Resolved URI—USS URI's are of the form 'uss://[Authority]/[Mount Point]/[Subdirectory]'. This URI does not point to a 'real' FileSystem, but is just an index into a Catalog. It is the responsibility of the USSNameNode to lookup the Catalog for the USS URI and resolve it to the correct location on an underlying 'Delegate FileSystem'. The 'real' URI returned after looking up the Catalog System is referred to in this document as a 'Resolved URI'. The 'Resolved URI' points to a real location on the 'Delegate FileSystem'
- **Delegate FileSystem**—USSProtocolTranslator provides an abstraction over a range of FileSystems in USS. All URIs in USS look like 'uss://[Authority]/[Mount Point]/[Subdirectory]'. Every time a URI lookup is desired, the call goes to USSProtocolTranslator which calls the USSNameNode to figure out the actual location of the URI, which resides on some underlying FileSystem (S3, HDFS, Isilon, Atmos, FTP, etc). The FileSystem call is then 'delegated' to this underlying FileSystem. This underlying FileSystem is referred to in this document as 'Delegate FileSystem'.

Terminology 89

• Client—A client is the entity that sends requests to USS. A client could be a user using FsShell (Hadoop command line) to request access to a USS URI, a Java program that uses the Hadoop FileSystem API or a Map-Reduce program that uses the JobClient to setup a map-reduce job.

Terminology 90

5. Security

Kerberos is a network authentication protocol that provides strong authentication for client/server applications using secret-key cryptography.

You must install and configure Kerberos to enable security in Pivotal HD 1.0.x.

HDFS, Mapreduce, Yarn, and Hive, and Pivotal HAWQ can be enabled for Kerberos.

Note: HAWQ does not currently work if HDFS is configured to use Kerberos.

This chapter contains the following:

- Configuring Kerberos for HDFS and YARN (MapReduce)
- Zookeeper Secure Configuration
- HBase Secure Configuration
- Hive Secure Configuration
- USS Secure Configuration
- MapReduce Version 1 Configuration (MRv1)
- Troubleshooting

Configuring Kerberos for HDFS and YARN (MapReduce)

At a minimum Kerberos provides protection against user and service spoofing attacks, and allows for enforcement of user HDFS access permissions. The installation is not difficult, but requires very specific instructions with many steps, and suffers from the same difficulties as any system requiring distributed configuration. Pivotal is working to automate the process to make it simple for users to enable/disable secure PHD clusters. Until then these instructions are intended to provide a step by step process for getting a cluster up and running in secure mode.

Note that after the initial HDFS/YARN configuration other services that need to be set-up to run on secure HDFS (for example, HBase) or that you want to also secure (for example, Zookeeper) need to configured.

Important: Save your command history, it will help in checking for errors when troubleshooting

Kerberos Set-up

Install the KDC

If you do not have a pre-existing KDC see "Installing the MIT Kerberos 5 KDC" on page 106.

Note: CentOS and RedHat use AES-256 as the default encryption strength. If you want to use AES-256 you need to install the JCE security policy file (described below) on all cluster hosts. If not, disable this encryption type in the KDC configuration. To disable AES-256 on an MIT kerberos 5 KDC remove aes256-cts:normal from the supported encrypes parameter in kdc.conf.

Integrating Cluster Security with an Organizational KDC

If your organization runs Active Directory or other Kerberos KDC it is not recommended this be used for cluster security. Instead install an MIT Kerberos KDC and realm for the cluster(s) and create all the service principals in this realm as per the instructions below. This KDC will be minimally used for service principals whilst Active Directory (or your organizations's MIT KDC) will be used for cluster users. Next configure one-way cross-realm trust from this realm to the Active Directory or corporate KDC realm.

Important: This is strongly recommended as a large PHD cluster requires large numbers of service principals be created by the IT manager for your organizations' Active Directory or organizational MIT KDC. For example a 100 node PHD cluster requires 200+ service principals. In addition when a large cluster starts up it may impact the performance of your organizations' IT systems as all the service principals make requests of the AD or MIT Kerberos KDC at once.

Install Kerberos Workstation and Libraries on Cluster Hosts

If you are using MIT krb5 run the following command:

yum install krb5-libs krb5-workstation

Distribute the Kerberos Client Configuration File to all Cluster Hosts

If you are using Kerberos 5 MIT this is /etc/krb5.conf. This file must exist on all cluster hosts. For PHD you can use massh to push the files, and then to copy them to the proper place.

Create the Principals

These instructions are for MIT Kerberos 5, command syntax for other Kerberos versions may be different.

Principals (Kerberos users) are of the form: name/role@REALM. For our purposes the name will be a PHD service name (for example: hdfs) and the role will be a DNS resolvable hostname (one you can use to connect to the host in question).

Important:

- Replace REALM with the KDC realm you are using for your PHD cluster where it appears.
- The host names used MUST be resolvable to an address on all the cluster hosts and MUST be of the form host.domain as some Hadoop components require at least one "." part in the host names used for principals.
- The names of the principals matter as some processes may throw exceptions if you change them. Hence it is safest to use the specified Hadoop principal names.

- Hadoop does support an _HOST tag in the site XML that is interpreted as the hostname but this can be problematic to use as it must resolve to the exact same name used in the principals.
- The mapping is to hostname not hostname -f and hence typically does not include the domain part. Using _HOST.domain does not work as Hadoop does not substitute properly in this case; whereas _HOST@REALM is interpreted properly. Given some of the issues surrounding _HOST these instructions do not assume its usage.

For the HDFS services you will need to create an hdfs/hostname principal for each host running an HDFS service (namenode, secondary namenode, datanode). For YARN services you will need to create a yarn/hostname principal for each host running a YARN service (resource manager, node manager). For MapReduce services you need to create a principal mapred/hostname for the Job History Server.

To create the required secure HD principals (using krb5 command syntax):

- For each cluster host (except client-only hosts) run:
 addprinc -randkey HTTP/hostname@REALM
- HDFS (namenode, datanodes), for each HDFS service host, run:
 addprinc -randkey hdfs/hostname@REALM
- YARN (resource manager, node managers), for each YARN service host, run: addprinc -randkey yarn/hostname@REALM
- MAPRED (job history server): for each JHS service host, run: addprinc -randkey mapred/hostname@REALM

Important: If you have 1000 cluster hosts running HDFS and YARN you will need 2000 HDFS and YARN principals, and need to distribute their keytab files. It is recommended you use a cluster-local KDC for this purpose and configure cross-realm trust to your organizational Active Directory or other Kerberos KDC.

Create the Keytab Files

Important: You MUST use kadmin.local (or the equivalent in your KDC) for this step on the KDC as kadmin does not support -norandkey.

Important: You can put the keytab files anywhere during this step, in this document we are creating a directory /etc/security/keytab/ and using that on cluster hosts, and so for consistency are placing them in a similarly named directory on the KDC. If the node you are on already has files in /etc/security/keytab/ it may be best to create a separate, empty, directory for this step.

Each service's keytab file for a given host will contain the service principal for that host and the HTTP principal for that host.

HDFS key tabs

For each host having an HDFS process (resource manager or node manager) run:

```
kadmin.local: ktadd -norandkey -k
/etc/security/keytab/hdfs-hostid.service.keytab
hdfs/hostname@REALM HTTP/hostname@REALM
```

Where hostid is just a short name for the host, for example, vm1, vm2, etc. This is to differentiate the files by host. You can use the hostname if desired.

For example for a 3 node cluster (one namenode, 2 datanodes):

```
kadmin.local: ktadd -norandkey -k
/etc/security/keytab/hdfs-vm2.service.keytab
hdfs/centos62-2.localdomain@BIGDATA.COM
HTTP/centos62-2.localdomain@BIGDATA.COM
kadmin.local: ktadd -norandkey -k
/etc/security/keytab/hdfs-vm3.service.keytab
hdfs/centos62-3.localdomain@BIGDATA.COM
HTTP/centos62-3.localdomain@BIGDATA.COM
kadmin.local: ktadd -norandkey -k
/etc/security/keytab/hdfs-vm4.service.keytab
hdfs/centos62-4.localdomain@BIGDATA.COM
HTTP/centos62-4.localdomain@BIGDATA.COM
```

YARN keytabs

For each host having a YARN process (resource manager or node manager) run:

```
kadmin.local: ktadd -norandkey -k
/etc/security/keytab/yarn-hostid.service.keytab
yarn/hostname@REALM HTTP/hostname@REALM
```

For example, for a 3 node cluster (one node resource manager; 2 node managers):

```
kadmin.local: ktadd -norandkey -k
/etc/security/keytab/yarn-vm2.service.keytab
yarn/centos62-2.localdomain@BIGDATA.COM
HTTP/centos62-2.localdomain@BIGDATA.COM
kadmin.local: ktadd -norandkey -k
/etc/security/keytab/yarn-vm3.service.keytab
yarn/centos62-3.localdomain@BIGDATA.COM
HTTP/centos62-3.localdomain@BIGDATA.COM
kadmin.local: ktadd -norandkey -k
/etc/security/keytab/yarn-vm4.service.keytab
yarn/centos62-4.localdomain@BIGDATA.COM
HTTP/centos62-4.localdomain@BIGDATA.COM
```

MAPRED keytabs

For each host having a MapReduce job history server run:

```
kadmin.local: ktadd -norandkey -k
/etc/security/keytab/mapred-hostid.service.keytab
mapred/hostname@REALM HTTP/hostname@REALM
```

For example:

```
kadmin.local: ktadd -norandkey -k
/etc/security/keytab/mapred-vm2.service.keytab
mapred/centos62-2.localdomain@BIGDATA.COM
HTTP/centos62-2.localdomain@BIGDATA.COM
```

Distribute the Keytab Files

- **1.** On each cluster node create the directory for the keytab files; in this document we are using /etc/security/keytab.
- 2. Move all the keytab files for a given host to the keytab directory on that host. For example: hdfs-vm2.service.keytab, yarn-vm2.service.keytab, and mapred-vm2.service.keytab go to host vm2)
- **3.** On each host:
 - **a.** Change the permissions on all keytabs to read-write by owner only: chmod 400 *.keytab
 - **b.** Change the group on all keytab files to hadoop: chgrp hadoop *
 - **c.** Change the owner of each keytab to the relevant principal name. For example, for yarn-vm2.service.keytab run: chown yarn yarn-vm2.service.keytab
 - d. Create links to the files of the form principalname.service.keytab. For example, for yarn-vm2.service.keytab run: ln -s yarn-vm2.service.keytab yarn.service.keytab

Important: The last step above allows you to maintain clear identification of each keytab file while also allowing you to have common site xml files across cluster hosts.

Below is an example keytab directory for a cluster control node (namenode, resource manager, JHS):

```
lrwxrwxrwx 1 root
                     root
                             23 Jun 10 23:50
hdfs.service.keytab -> hdfs-vm2.service.keytab
-rw----- 1 hdfs hadoop 954 Jun 10 23:44
hdfs-vm2.service.keytab
lrwxrwxrwx 1 root root
                             25 Jun 10 23:51
mapred.service.keytab -> mapred-vm2.service.keytab
-rw----- 1 mapred hadoop 966 Jun 10 23:44
mapred-vm2.service.keytab
lrwxrwxrwx 1 root
                             23 Jun 10 23:51
                     root
yarn.service.keytab -> yarn-vm2.service.keytab
-rw----- 1 yarn
                      hadoop 954 Jun 10 23:44
yarn-vm2.service.keytab
```

Below is an example keytab directory for a cluster node (datanode, node manager):

```
lrwxrwxrwx 1 root root 23 Jun 11 01:58 hdfs.service.keytab
-> hdfs-vm3.service.keytab
-rw----- 1 hdfs hadoop 954 Jun 10 23:45
hdfs-vm3.service.keytab
lrwxrwxrwx 1 root root 23 Jun 11 01:58 yarn.service.keytab
-> yarn-vm3.service.keytab
```

```
-rw----- 1 yarn hadoop 954 Jun 10 23:45 yarn-vm3.service.keytab
```

Install Java Support Items

Install JCE on all Cluster Hosts

Important: This step is only needed if you are using AES-256

Note: These files will already exist in your environment and appear to be the same, but are the limited strength encryption files, you must replace them with the unlimited strength files to use AES-256

- **1.** Download and unzip the JCE file for your JDK version (Java Cryptography Extension (JCE) Unlimited Strength Jurisdiction Policy Files 7 for JDK 7 and Java Cryptography Extension (JCE) Unlimited Strength Jurisdiction Policy Files 6 for JDK6).
- **2.** Place the local_policy.jar and US_export_policy.jar files in the /usr/java/default/jre/lib/security/ directory on all cluster hosts.

Check JSVC on all Datanodes

JSVC allows a Java process to start as root and then switch to a less privileged user, and is required for the datanode process to start in secure mode. Your distribution comes with a pre-built JSVC, you need to verify it can find a JVM as follows:

1. Run:

```
/usr/libexec/bigtop-utils/jsvc -help
```

2. Look under the printed -jvm item in the output and you should see something like:

```
use a specific Java Virtual Machine. Available JVMs: 'server'
```

If you do not see the server line this jsvc will not work for your platform so try:

- **3.** Install JSVC using yum and run the check again; if it fails try **4.**) below:
- **4.** Build from source and install manually (see "Building and Installing JSVC" on page 105).

If you have datanode start-up problems and no other errors are obvious it might be a JSVC problem and you may need to do item **4.)** above. JSVC is very precise about platform and JDK matching, so use the instructions in "Building and Installing JSVC" on page 105 for your system OS and JDK.

Container and Script Modifications

Configure the Linux Container

1. Edit the

```
/usr/lib/gphd/hadoop-yarn/etc/hadoop/container-executor.cfg to
be:

# NOTE: these next two should be set to the same values they
have in yarn-site.xml
yarn.nodemanager.local-dirs=/data/1/yarn/nm-local-dir
yarn.nodemanager.log-dirs=/data/1/yarn/userlogs
# configured value of
yarn.nodemanager.linux-container-executor.group
yarn.nodemanager.linux-container-executor.group=yarn
# comma separated list of users who can not run applications
banned.users=hdfs,yarn,mapred,bin
# Prevent other super-users
min.user.id=500
```

Note: The min.user.id varies by Linux dist; for CentOS it is 500, RedHat is 1000.

2. Check the permissions on

```
\label{like-decomposition} $$ \user/lib/gphd/hadoop-yarn/bin/container-executor. $$ They should look like this:
```

```
---Sr-s--- 1 root yarn 364 Jun 11 00:08 container-executor If they do not, then set the owner, group and permissions as:
```

```
chown root:yarn container-executor
chmod 050 container-executor
chmod u+s container-executor
chmod g+s container-executor
```

3. Check the permissions on

/usr/lib/gphd/hadoop-yarn/etc/hadoop/container-executor.cfg. They should look like this:

```
-rw-r--r-- 1 root root 363 Jul 4 00:29 container-executor.cfg
```

If they do not, then set the owner, group and permissions as:

```
chown root:root container-executor.cfg
chmod 644 container-executor.cfg
```

Edit the Environment on the Datanodes

Important:

- At this point you should STOP the cluster if it is running
- You only need to do the steps below on the datanodes
- **1.** Uncomment the lines at the end of /etc/default/hadoop-hdfs-datanode:

```
# secure operation stuff
export HADOOP_SECURE_DN_USER=hdfs
export HADOOP_SECURE_DN_LOG_DIR=${HADOOP_LOG_DIR}/hdfs
export HADOOP_PID_DIR=/var/run/gphd/hadoop-hdfs/
export HADOOP_SECURE_DN_PID_DIR=${HADOOP_PID_DIR}
```

2. Set the JSVC variable:

If you are using the included jsvc the <code>JSVC_HOME</code> variable in /etc/default/hadoop should already be properly set:

```
export JSVC HOME=/usr/libexec/bigtop-utils
```

If you built or manually installed JSVC your <code>JSVC_HOME</code> will be <code>/usr/bin</code> and you must set it appropriately by modifying <code>/etc/default/hadoop</code> and setting the proper <code>JSVC_HOME</code>:

```
export JSVC HOME=/usr/bin
```

Important: Make sure JSVC HOME points to the correct jsvc binary

WARNING: As long as HADOOP_SECURE_DN_USER is set the datanode will try and start in secure mode.

Site XML Changes

Edit the Site XML

Finally you need to edit the site XML to turn on secure mode. Before proceeding it is good to understand who needs to talk to whom. By "talk" we mean use authenticated kerberos to initiate establishment of a communication channel. Doing this requires that you know your own principal to identify yourself and know the principal of the service you want to talk to. To be able to use its principal, a service needs to be able to login to Kerberos without a password using a keytab file.

- Each service needs to know its own principal name, of course
- Each running service on a node needs a service/host specific keytab file to start up
- Each datanode needs to talk to the namenode
- Each node manager needs to talk to the resource manager and the job history server
- Each client/gateway node needs to talk to the namenode, resource manager, and job history server

Remembering this helps when setting up and troubleshooting the site xml files.

Important:

- Redundant keytab files on some hosts do no harm and it makes management easier to have constant files. Remember though that the hostname MUST be correct for each entry.
- Before making changes, backup the current site xml files so that you can return to non-secure operation if needed.

Much of the changes can be consistent throughout the cluster site XML, but since datanode and node manager principals are hostname dependent (or more correctly the role for the yarn principal is set to the hostname), the yarn-site.xml for datanode and node manager principals will differ across the cluster.

1. Edit the /usr/lib/gphd/hadoop/etc/hadoop/core-site.xml as follows:

2. Edit the /usr/lib/gphd/hadoop/etc/hadoop/hdfs-site.xml as follows:

```
</property>
cproperty>
  <name>dfs.namenode.kerberos.http.principal
  <value>HTTP/namenode-hostname@REALM</value>
</property>
cproperty>
<name>dfs.namenode.kerberos.internal.spnego.principal/name>
  <value>HTTP/namenode-hostname@REALM</value>
</property>
<!-- ?(optional) secondary namenode secure configuration
info -->
cproperty>
  <name>dfs.secondary.namenode.keytab.file
  <value>/etc/security/keytab/hdfs.service.keytab</value>
</property>
cproperty>
  <name>dfs.secondary.namenode.kerberos.principal</name>
  <value>hdfs/secondary-namenode-hostname@REALM</value>
</property>
cproperty>
<name>dfs.secondary.namenode.kerberos.http.principal
  <value>HTTP/secondary-namenode-hostname@REALM</value>
</property>
cproperty>
<name>dfs.secondary.namenode.kerberos.internal.spnego.princi
pal</name>
  <value>HTTP/secondary-namenode-hostname@REALM</value>
</property>
<!-- datanode secure configuration info -->
cproperty>
  <name>dfs.datanode.data.dir.perm
```

```
<value>700</value>
</property>
<!-- these ports must be set < 1024 for secure operation -->
<!-- conversely they must be set back to > 1024 for
non-secure operation -->
cproperty>
  <name>dfs.datanode.address</name>
  <value>0.0.0.0:1004
</property>
cproperty>
  <name>dfs.datanode.http.address</name>
  <value>0.0.0.0:1006
</property>
<!-- remember the principal for the datanode is the
principal this hdfs-site.xml file is on -->
<!-- these (next three) need only be set on datanodes -->
cproperty>
  <name>dfs.datanode.kerberos.principal
  <value>hdfs/this-datanodes-hostname@REALM</value>
</property>
cproperty>
  <name>dfs.datanode.kerberos.http.principal
  <value>HTTP/this-datanodes-hostname@REALM</value>
</property>
cproperty>
  <name>dfs.datanode.keytab.file
  <value>/etc/security/keytab/hdfs.service.keytab</value>
</property>
<!-- OPTIONAL - set these to enable secure WebHDSF -->
<!-- on all HDFS cluster nodes -->
<!--(namenode, secondary namenode, datanode's) -->
property>
  <name>dfs.webhdfs.enabled
```

```
<value>true</value>
   </property>
   cproperty>
     <name>dfs.web.authentication.kerberos.principal</name>
     <value>HTTP/this-datanodes-hostname@REALM</value>
   </property>
   <!-- since we included the HTTP principal all -->
   <!-- keytabs we can use it here -->
   property>
     <name>dfs.web.authentication.kerberos.keytab
     <value>/etc/security/keytab/hdfs.service.keytab</value>
   </property>
   <!-- THE PROPERTIES BELOW ARE OPTIONAL AND -->
   <!-- REQUIRE RPC PRIVACY (core-site): -->
   <!-- THEY ENABLE ON WIRE HDFS BLOCK ENCRYPTION -->
   cproperty>
      <name>dfs.encrypt.data.transfer</name>
      <value>true</value>
   </property>
   cproperty>
      <name>dfs.encrypt.data.transfer.algorithm</name>
      <value>rc4</value>
      <description>may be "rc4" or "3des" - 3des has a
      significant performance impact</description>
   </property>
3. Edit the /usr/lib/gphd/hadoop/etc/hadoop/yarn-site.xml as follows:
   <!-- resource manager secure configuration info -->
   cproperty>
     <name>yarn.resourcemanager.principal</name>
     <value>yarn/resourcemgr-hostname@REALM</value>
   </property>
   cproperty>
```

```
<name>yarn.resourcemanager.keytab</name>
  <value>/etc/security/keytab/yarn.service.keytab</value>
</property>
<!-- remember the principal for the node manager is the
principal for the host this yarn-site.xml file is on -->
<!-- these (next four) need only be set on node manager nodes
-->
property>
  <name>yarn.nodemanager.principal
  <value>yarn/this-nodemgrs-hostname@REALM</value>
</property>
cproperty>
  <name>yarn.nodemanager.keytab</name>
  <value>/etc/security/keytab/yarn.service.keytab</value>
</property>
cproperty>
  <name>yarn.nodemanager.container-executor.class
<value>org.apache.hadoop.yarn.server.nodemanager.LinuxContai
nerExecutor</value>
</property>
cproperty>
<name>yarn.nodemanager.linux-container-executor.group
  <value>yarn</value>
</property>
edit /usr/lib/gphd/hadoop/etc/hadoop/mapred-site.xml:
<!-- job history server secure configuration info -->
property>
  <name>mapreduce.jobhistory.keytab</name>
  <value>/etc/security/keytab/mapred.service.keytab</value>
</property>
property>
```

```
<name>mapreduce.jobhistory.principal</name>
  <value>mapred/jobhistoryserver-hostname@REALM</value>
</property>
```

Complete the HDFS/YARN Secure Configuration

1. Start the cluster:

```
icm client start
```

- **2.** Check that all the processes start up. If not, refer to "Troubleshooting" on page 120.
 - **Control Processes**: namenode, resourcemanager, historyserver should all be running.
 - Cluster Worker Processes: datanode and namenode should be running.

 Note: Until you configure HBase security configuration, HBase will not start up on a secure cluster.
- **3.** Create a principal for a standard user (user must exist as a Linux user on all cluster hosts):

```
kadmin: addprinc testuser
Set the password when prompted.
```

- **4.** Login as that user on a client box (or any cluster box if you do not have specific client purposed systems).
- **5.** Get your kerberos TGT by running kinit and entering the password: kinit testuser.
- **6.** Test simple HDFS file list and directory create:

```
hadoop fs -ls
hadoop fs -mkdir testdir
```

If these do not work refer to "Troubleshooting" on page 120.

- **7.** Optional. Set the sticky bit on the /tmp directory (prevents non-super users from moving or deleting other users' files in /tmp)
 - **a.** Login as gpadmin on any HDFS service node (namenode, datanode)
 - **b.** Execute the following:

```
sudo -u hdfs kinit -k -t
/etc/security/keytab/hdfs.service.keytab
hdfs/this-hostname@REALM
```

c. Execute the following:

```
sudo -u hdfs hadoop fs -chmod 1777 /tmp
```

d. Run a simple MapReduce job such as the pi example:

```
hadoop jar
/usr/lib/gphd/hadoop-mapreduce/hadoop-mapreduce-examples-
2.0.2-alpha-gphd-2.0.1.0.jar pi 10 100
```

If this all works then you are ready to configure other services, if not refer to "Troubleshooting" on page 120.

Turning Secure Mode Off

To turn off secure mode:

1. Stop the cluster:

```
icm client stop
```

- **2.** Comment out HADOOP_SECURE_DN_USER in hadoop-env.sh and /etc/init.d/hadoop-hdfs-datanode on all datanodes
- **3.** Either:
 - **a.** If you made backups as suggested above: Restore the original site xml files.

or:

- **b.** If you do not have backups, then edit the site xml, as follows:
- Set Linux container executable to org.apache.hadoop.yarn.server.nodemanager.DefaultContainerEx ecutor on all datanodes.
- Set dfs.block.access.token.enable to false on all datanodes
- Return the datanode ports modified above so they are > 1024 again.
- Set hadoop.security.authentication to simple and hadoop.security.authorization to false in core-site.xml on all cluster nodes.
- Start the cluster: icm client start

Building and Installing JSVC

In order for the datanodes to start as root to get secure ports and then switch back to the hdfs user, JSVC must be installed. If the packaged jsvc binary is not working we recommend building jscv from source for your platform.

You only need to execute the make on one node, then the binary can be distributed to the others (assuming all systems are the same basic image):

1. Install gcc and make (you can remove them after this process if desired).

```
yum install gcc make
```

2. Download the Apache commons daemon; for example.

```
commons-daemon-1.0.15-src.zip was tested.
```

The demon is available here:

```
http://commons.apache.org/proper/commons-daemon/download daemon.cgi
```

- **3.** scp it to one of your datanode cluster systems.
- **4.** Uncompress it.

5. Change to the install directory:

```
cd commons-daemon-1.0.15-src/src/native/unix
```

6. If you are on a 64 bit machine and using a 64 bit JVM perform these exports before continuing to configure/make:

```
export CFLAGS=-m64
export LDFLAGS=-m64
```

7. Configure and make it:

```
./configure --with-java=/usr/java/default
make
```

8. Manually install it to the following location:

```
mv ./jsvc /usr/bin/jsvc
```

9. Check that the correct jsvc is found by doing running:

```
which jsvc
```

The correct output is:

/usr/bin/jsvc

10. Run:

```
jsvc -help
```

Look under the printed -jvm item in the output and you should see something like:

```
use a specific Java Virtual Machine. Available JVMs: 'server'
```

If the line under Available JVMs (where server appears above) is blank there is a problem as it cannot find the JVM; check that the JDK is installed properly in /usr/java/default.

Installing the MIT Kerberos 5 KDC

This section outlines a simple krb5 KDC set-up, mainly for test and developer purposes. These instructions were largely derived from "*Kerberos: The Definitive Guide*", James Garman, O'Reilly, pages 53-62.

- **1.** Install the Kerberos packages krb5-libs, krb5-workstation, and krb5-server on the KDC host.
- **2.** Define your REALM in /etc/krb5.conf.
 - For testing you can use the EXAMPLE.COM REALM.
 - Set the kdc and admin_server variables to the resolvable hostname of the KDC host.
 - Set the default domain to your REALM.

In the following example, REALM was changed to BIGDATA.COM and the KDC host is centos62-1.localdomain:

```
[logging]
```

```
default = FILE:/var/log/krb5libs.log
kdc = FILE:/var/log/krb5kdc.log
admin server = FILE:/var/log/kadmind.log
[libdefaults]
default realm = BIGDATA.COM
dns lookup realm = false
dns lookup kdc = false
ticket lifetime = 24h
renew lifetime = 7d
forwardable = true
[realms]
BIGDATA.COM = {
 kdc = centos62-1.localdomain:88
 admin server = centos62-1.localdomain:749
 default domain = BIGDATA.COM
}
[domain realm]
 .bigdata.com = BIGDATA.COM
bigdata.com = BIGDATA.COM
```

- **3.** Set up /var/kerberos/krb5kdc/kdc.conf.
 - If you want to use AES-256, uncomment the master key type line.
 - If you do not want to use AES-256, remove it from the supported enctypes line
 - Add a key stash file entry: /var/kerberos/krb5kdc/.k5.REALM.
 - Add the kadmind port entry: kadmind port = 749.

Important: The stash file lets the KDC server start up for root without a password being entered.

The result (using AES-256) for the above REALM is:

```
[kdcdefaults]
kdc_ports = 88
kdc_tcp_ports = 88

[realms]
BIGDATA.COM = {
  master_key_type = aes256-cts
  acl_file = /var/kerberos/krb5kdc/kadm5.acl
  dict file = /usr/share/dict/words
```

```
admin_keytab = /var/kerberos/krb5kdc/kadm5.keytab
key_stash_file = /var/kerberos/krb5kdc/.k5.BIGDATA.COM
kadmind_port = 749
supported_enctypes = aes256-cts:normal
aes128-cts:normal des3-hmac-sha1:normal
arcfour-hmac:normal des-hmac-sha1:normal
des-cbc-md5:normal des-cbc-crc:normal
```

4. Create the KDC master password:

```
run: kdb5 util create -s
```

DO NOT forget your password as this is the root KDC password.

5. Add an administrator account as username/admin@REALM:

```
start kadmin.local
kadmin.local: addprinc username/admin@REALM
```

Important: The KDC does not need to be running to add a principal.

6. Start the KDC by running:

```
service krb5kdc start
```

You should get an [OK] indication if it started without error.

7. Edit /var/kerberos/krb5kdc/kadm5.acl and change the admin permissions username from * to your admin.

You can add other admins with specific permissions if you want (man kadmind)

This is a sample ACL file:

```
joeit/admin@BIGDATA.COM
```

8. Use kadmin.local on the KDC to enable the administrator(s) remote access:

kadmind.local: ktadd -k /var/kerberos/krb5kdc/kadm5.keytab kadmin/admin kadmin/changepw

Important: kadmin.local is a KDC host only version of kadmin that can do things remote kadmin cannot (such as use the -norandkey option in ktadd)

9. Start kadmind:

```
service kadmin start
```

The KDC should now be done and ready to use, but you need to set up your clients first.

- **10.** Install krb5-libs and krb5-workstation on all cluster hosts, including any client/gateway hosts.
- **11.** Push your KDC /etc/krb5.conf to all workstation hosts.
- **12.** Do a simple test, as follows:
 - a. Login as the admin you created: kinit username/admin
 - b. Run kadmin and make sure you can login

If you get the message kinit: Cannot contact any KDC for realm 'REALM' while getting initial credentials, then the KDC is not running or the KDC host information in /etc/kdc.conf is incorrect.

You now should have a KDC that is functional for PHD secure cluster operations.

Zookeeper Secure Configuration

Zookeeper secure configuration for server is recommended for HBase.

Important: STOP cluster services before performing this configuration.

Zookeeper Servers

Create the Zookeeper Principals:

1. Create a principal for each Zookeeper Quorum Server host: kadmin: addprinc -randkey zookeeper/hostname@REALM

Create the Zookeeper Keytab Files

1. For each Zookeeper server host:

```
ktadd -norandkey -k
/etc/security/keytab/zookeeper-hostid.service.keytab
zookeeper/hostname@REALM
```

Distribute the Zookeeper Keytab Files

1. For each Zookeeper server host:

```
/etc/security/keytab directory, then run the following: chgrp hadoop zookeeper-hostid.service.keytab chown zookeeper zookeeper-hostid.service.keytab chmod 400 zookeeper-hostid.service.keytab ln -s zookeeper-hostid.service.keytab zookeeper.service.keytab
```

Move the appropriate keytab file for each host to that hosts'

Edit the Zookeeper Configuration

1. Add the following lines to /etc/gphd/zookeeper/conf/zoo.cfg: authProvider.1=org.apache.zookeeper.server.auth.SASLAuthenticationProvider jaasLoginRenew=3600000

2. Create a file in /etc/gphd/zookeeper/conf/jaas.conf and add to it:

```
Server {
   com.sun.security.auth.module.Krb5LoginModule required
   useKeyTab=true

   keyTab="/etc/security/keytab/zookeeper-hostid.service.key
   tab"
```

```
storeKey=true
useTicketCache=false
principal="zookeeper/hostname@REALM";
};
```

3. Add the following line to /etc/gphd/zookeeper/conf/java.env (create the file if it does not exist):

```
export
JVMFLAGS="-Djava.security.auth.login.config=/etc/gphd/zookee
per/conf/jaas.conf"
```

Verify the Zookeeper Configuration

1. Start up the cluster and connect using a client.

Note: You do not need to set up clients to use Kerberos but if you want this functionality see "Zookeeper Clients" below:

2. Connect as: zookeeper-client -server hostname:port

Note: the port is defined in /etc/gphd/zookeeper/conf/zoo.cfg, and is typically 2181.

3. Create a protected znode:

create /testznode testznodedata sasl:zkcli@REALM:cdwra

4. Verify the znode:

```
getAcl /testznode
```

You should see something like this:

```
'sasl,'zkcli@{{BIGDATA.COM%7D%7D
: cdrwa
```

Zookeeper Clients

Optional

1. Add a principal for the client on the client host:

```
kadmin.local: addprinc -randkey zclient/hostname@REALM
```

2. Add the keytab:

```
kadmin.local: ktadd -norandkey -k
/etc/security/keytab/zclient-hostid.client.keytab
zclient/hostname@REALM
```

3. Move the file to the /etc/security/keytab directory on the host and change the owner and group appropriately so that only users of the client can access the file:

chmod 400 /etc/security/keytab/zclient-hostid.client.keytab

4. Create a link:

```
ln -s zclient-hostid.client.keytab zclient.client.keytab
```

5. Add the following to the file /etc/gphd/zookeeper/conf/jaas.conf (creating the file if required):

```
Client {
  com.sun.security.auth.module.Krb5LoginModule required
  useKeyTab=true
  keyTab="/etc/security/keytab/zclient.client.keytab"
  storeKey=true
  useTicketCache=false
  principal="zclient/hostname@REALM";
};
```

If you get a failure message indicating a name lookup failure that indicates you should add a name service setting, add or edit the following line to

/etc/gphd/zookeeper/conf/java.env (create the file if it does not exist) to be:

```
export
JVMFLAGS="-Djava.security.auth.login.config=/etc/gphd/zookee
per/conf/jaas.conf
-Dsun.net.spi.nameservice.provider.1=dns,sun"
```

You cannot do this on a server node as the

-Dsun.net.spi.nameservice.provider.1=dns, sun line may cause the server to fail to start.

You should now be able to establish a secure session with zookeeper-client. Test this by staring zookeeper-client and insuring no errors occur while connecting.

You may have issues with addressing or be forced to use the actual server IP address with the -server option for zookeeper-client to handle incompatibilities between the settings needed to make the Kerberos lookups work

(-Dsun.net.spi.nameservice.provider.1=dns, sun) and what makes the Java host resolution work. This problem also may be encountered in trying to set up HBase to communicate with a secure Zookeeper, where it is more difficult to resolve.

HBase Secure Configuration

If you are running secure HBase you should also also run a secure Zookeeper (see "Zookeeper Secure Configuration"). You can, however, set the HBase master and region servers up to use Kerberos and test that they start without a secure Zookeeper. This section covers the basics of how to get HBase up and running in secure mode; for further information see the HBase documentation

(http://hbase.apache.org/book/security.html).

HBase Master and Regionservers

Create the HBase Principals

1. For the HBase master and each region server host run:

```
kadmin.local: addprinc -randkey hbase/hostname@REALM Where hostname refers to the service principal (master, regionserver) host.
```

Create the HBase Keytab files

1. For the HBase master and each region server host run:

```
kadmin.local: ktadd -norandkey -k
/etc/security/keytab/hbase-hostid.service.keytab
hbase/hostname@REALM
```

Distribute the HBase Keytab Files

1. For each host:

Move the appropriate keytab file to that hosts' /etc/security/keytab directory, then run the following:

```
chown hbase:hadoop hbase-hostid.service.keytab
chmod 400 hbase-hostid.service.keytab
ln -s hbase-hostid.service.keytab hbase.service.keytab
```

Edit the HBase Site XML

1. For each master and region server host add to

```
/etc/gphd/hbase/conf/hbase-site.xml:
cproperty>
  <name>hbase.security.authentication</name>
   <value>kerberos</value>
</property>
cproperty>
  <name>hbase.security.authorization
   <value>true</value>
</property>
cproperty>
   <name>hbase.coprocessor.region.classes</name>
  <value>org.apache.hadoop.hbase.security.token.TokenProvid
  er</value>
</property>
<!-- HBase secure region server configuration -->
cproperty>
  <name>hbase.regionserver.kerberos.principal
   <value>hbase/regionserver-hostname@REALM</value>
</property>
cproperty>
  <name>hbase.regionserver.keytab.file</name>
   <value>/etc/security/keytab/hbase.service.keytab</value>
```

Test HBase Start-Up

You can now test HBase start-up. Start the cluster services and check that the HBase Master and Regionservers start properly. If they do not look at the .log file in the /var/log/gphd/hbase/ directory for hints as to why. Make sure HDFS came up properly. As you fix issues you can run service hbase-master start or service hbase-regionserver start to check that the issue is resolved.

HBase Clients

1. Add the following to the hbase-site.xml file on every client host:

Enable Encrypted Communication

Optional

If you are running secure HBase you can enable encryption from clients to server by adding the following to hbase-site.xml on all clients:

This can also be set on a per-connection basis. Set it in the Configuration supplied to HTable:

```
Configuration conf = HBaseConfiguration.create();
conf.set("hbase.rpc.protection", "privacy");
HTable table = new HTable(conf, tablename);
```

The Apache HBase documentation indicates to expect a \sim 10% performance penalty when encryption is enabled.

Access Control

The version of HBase distributed with PHD supports access control. See the HBase documentation

(http://hbase.apache.org/book/hbase.accesscontrol.configuration.html) for instructions on configuring access controls.

REST Gateway

You can set up the REST Gateway to use Kerberos to authenticate itself as a principal to HBase. Note that all client access will use the REST Gateway's credentials set below, and have this user's privileges.

For every REST Gateway add the following to hbase-site.xml file:

You must also give the REST principal access privileges. Do this by adding the rest-principal-name to the acl table in HBase. Adding the permissions below are sufficient according to HBase documentation:

```
grant 'rest-principal-name', 'RWCA'
```

Thrift Client Configuration

See the HBase documentation (http://hbase.apache.org/book/security.html) for instructions on configuring Thrift clients.

HBase with Secure Zookeeper Configuration

For secure HBase you should also run a secure Zookeeper (see Zookeeper Secure Configuration above). If you do so you will need to execute the steps in this section. These steps must be done on the HBase master and all region servers.

1. Create a file /etc/gphd/hbase/conf/jaas.conf and add the following to it:

```
Client {
  com.sun.security.auth.module.Krb5LoginModule required
  useKeyTab=true
  useTicketCache=false
  keyTab="/etc/security/keytab/hbase.service.keytab"
  principal="hbase/hostname@REALM";
};
```

Make sure to replace *hostname@REALM* with the hostname of the server and the correct REALM.

2. Add the following near at the end of /etc/qphd/hbase/conf/hbase-env.sh:

```
export HBASE_OPTS="$HBASE_OPTS
-Djava.security.auth.login.config=/etc/gphd/hbase/conf/jaas.conf"
export HBASE MANAGES ZK=false
```

3. Edit the site xml and add the following:

4. Edit /etc/gphd/zookeeper/conf/zoo.cfg and add the following:

```
kerberos.removeHostFromPrincipal=true
kerberos.removeRealmFromPrincipal=true
```

Hive Secure Configuration

The Hive MetaStore supports Kerberos authentication for Thrift clients. You can configure a standalone Hive MetaStoreServer instance to force clients to authenticate with Kerberos by setting the property hive.metastore.sasl.enabled property in the hive-default.xml configuration file to true, as shown in the example below.

Add the Kerberos principals and their locations to the hive-default.xml or hive-site.xml (if you are user). For example:

```
property>
   <name>hive.metastore.sasl.enabled
  <value>true</value>
  <description>If true, the metastore thrift interface will be
   secured with SASL. Clients must authenticate with
  Kerberos.□</description>
</property>
cproperty>
   <name>hive.metastore.kerberos.keytab.file
   <value>/etc/****/hive.keytab</value>
  <description>The path to the Kerberos Keytab file containing the
  metastore thrift server's service principal.□</description>
   </property>
cproperty>
   <name>hive.metastore.kerberos.principal
   <value>hive-metastore/ HOST@EXAMPLE.COM</value>
   <description>The service principal for the metastore thrift
```

```
server. The special string _HOST will be replaced automatically
with the correct host name.</description>
</property>
```

USS Secure Configuration

USS version 0.4.0 adds support for secure Hadoop clusters. To work on a secure hadoop cluster, users have to add additional metadata about their cluster's security parameters to the USS catalog. In version 0.4.0, the metadata can be added using SQL commands. The USS CLI will be enhanced in an upcoming release to support secure hadoop clusters.

Securing USS

To enable security in USS using Keberos, add the following properties to the US configuration files. After doing so, restart the USS namenode:

uss-client-site.xml and uss-nn-site.xml

Property	Description
uss.security.authentication	This property indicates the authentication that USS must use. It is set to kerberos for secure hadoop clusters. For non-secure clusters, this property defaults to simple.
uss.namenode.principal	This property is set to the kerberos principal of the USS Namenode. The secure tools can generate the principal and set this property.

Accessing a secure HDFS cluster through USS

To access a secure HDFS cluster through USS, update the delegate_filesystem table in the USS catalog database by following the steps below:

1. Add a mount point as usual:

```
uss admin --add --mount-point-name "secure_mountpoint" -t
"hdfs://namenode:8020/user"
```

2. Login to the catalog on the USS Catalog host:

```
psql -h localhost -p 10432 -d usscatalog
```

3. Get the ID of the filesystem you just added by looking for the ID which has the base URI of the mount point you added.

```
SELECT base_uri, id from delegate_filesystem
where base uri like '%namenode:8020%';
```

4. Set security type to kerberos. The following statement assumes that you received the ID as 3 in the previous step:

```
update delegate_filesystem set security_type='kerberos'
where id = 3;
```

5. Set security info json by replacing \$PRIMARY_SERVER_PRINCIPAL and \$SECONDARY_SERVER_PRINCIPAL in the sql below with their values from hdfs-site.xml (dfs.namenode.kerberos.principal and dfs.secondary.namenode.kerberos.principal) update delegate_filesystem set security_info='{"primaryServerPrincipal": "\$PRIMARY_SERVER_PRINCIPAL", "secondaryServerSrincipal": "\$PRIMARY_SERVER_PRINCIPAL" }' where id = 3;

MapReduce Version 1 Configuration (MRv1)

To configure MRv1 follow the same steps as for HDFS and YARN except for the following differences.

MRv1 Kerberos Set-up Differences

For MRV1 do not create the yarn principals. Instead create the mapred principal for all task tracker and job tracker hosts.

MAPRED (JobTrackers, TaskTrackers): for each service host you need to do:
 addprinc -randkey mapred/hostname@REALM

You will not need to create and distribute yarn keytab files, but you will instead have to create and distribute the mapred keytab files for the JobTracker and TaskTracker hosts.

Container and Script Modifications Differences

1. Instead of editing the

/usr/lib/gphd/hadoop-yarn/etc/hadoop/container-executor.cfg file, edit the /usr/lib/gphd/hadoop-mr1/etc/hadoop/task-controller.cfg file to be:

```
# NOTE: the "hadoop.log.dir" should be the same value as what
the Hadoop daemons are using
hadoop.log.dir=/var/log/gphd/hadoop-mr1
mapreduce.tasktracker.group=mapred
banned.users=mapred,hdfs,bin
min.user.id=500
```

Note: The min.user.id varies by Linux dist; for CentOS it is 500, RedHat is 1000.

2. Check the permissions on

```
/usr/lib/gphd/hadoop-mr1/bin/task-controller. They should look like:
----Sr-s--- 1 root mapred 286 Jun 18 00:08 task-controller
If they do not then set the owner, group, and permissions as:
```

```
chown root:mapred task-controller
chmod 050 task-controller
chmod u+s task-controller
```

chmod g+s task-controller

MRv1 Site XML Differences

For MRv1 you do not edit the yarn-site.xml file. Instead add the following to the mapred-site.xml file:

```
<!-- JobTracker configuration info -->
<!-- JobTracker principal must be known to all cluster hosts
-->
cproperty>
  <name>mapreduce.jobtracker.kerberos.principal
  <value>mapred/jobtracker-hostname@REALM</value>
</property>
<!-- keytab only needs to be know to the JobTracker host -->
property>
  <name>mapreduce.jobtracker.keytab.file</name>
   <value>/var/local/hadoop/mapred.service.keytab</value>
</property>
<!-- TaskTracker configuration info -->
<!-- TaskTracker principal must be known to all cluster
hosts -->
cproperty>
  <name>mapreduce.tasktracker.kerberos.principal
   <value>mapred/this-tasktracker-hostname@REALM</value>
</property>
cproperty>
   <name>mapreduce.tasktracker.keytab.file
   <value>/var/local/hadoop/mapred.service.keytab</value>
</property>
<!-- TaskController configuration info -->
cproperty>
  <name>mapred.task.tracker.task-controller
  <value>org.apache.hadoop.mapred.LinuxTaskController</valu</pre>
  e>
</property>
cproperty>
  <name>mapreduce.tasktracker.group
   <value>mapred</value>
```

```
</property>
```

Auditing

You can enable auditing before deployment or re-configuration of a cluster.

To enable auditing:

- **1.** Locate your templates directory (by default ClusterConfigDir, this is created during initial instatallation, see *Pivotal HD Enterprise 1.0 Installation and Administrator Guide* for details).
- **2.** For HDFS and MapReduce, locate the hdfs subdirectory and edit the log4j.properties file as follows:

```
For HDFS change line:
```

```
hdfs.audit.logger=INFO,NullAppender
to:
hdfs.audit.logger=INFO,RFAAUDIT
```

For MapReduce change line:

```
mapred.audit.logger=INFO,NullAppender
to:
mapred.audit.logger=INFO,RFAAUDIT
```

For other components, locate the component sub-directory in the template and its corresponding log4j.properties file and make similar edits.

To specify auditing output location:

By default, log files and other auditing information is output to /var/log/gphd/hadoop/.

To set up logging to go to syslog, define the following:

```
# Configure syslog appender
log4j.appender.SYSLOG=org.apache.log4j.net.SyslogAppender
log4j.appender.SYSLOG.syslogHost=loghost
log4j.appender.SYSLOG.layout=org.apache.log4j.PatternLayout
log4j.appender.SYSLOG.layout.ConversionPattern=%d{ISO8601}
%p %c: %m%n
log4j.appender.SYSLOG.Facility=LOCAL1
```

You can now log audit information to syslog, for example:

```
hdfs.audit.logger=INFO,SYSLOG
```

You can also log to file and syslog, for example:

```
hdfs.audit.logger=INFO,RFAAUDIT,SYSLOG
```

Note that these changes only go into effect after deployment or re-configuration.

Auditing 119

Troubleshooting

- A good first step is to look for exceptions that may give you a hint in the logs (where hostname is the host on which the log file is located):
 - namenode

/var/log/gphd/hadoop/hadoop-hdfs-namenode-hostname.log

• resourcemanager:

/ var/log/gphd/hadoop-yarn/yarn-yarn-resource manager-host name .log

• historyserver:

/var/log/gphd/hadoop-mapreduce/mapred-mapred-historyserver-h ostname.log

datanode:

/var/log/gphd/hadoop-hdfs/hdfs/hadoop-hdfs-datanode-hostname .log

nodemanager:

/var/log/gphd//hadoop-yarn/yarn-yarn-nodemanager-hostname.log

You can enable debug level logging for the Java Kerberos classes by editing /etc/default/hadoop and setting the following:

HADOOP OPTS="\$HADOOP OPTS -Dsun.security.krb5.debug=true"

- Datanode will not start:
 - If you are getting a message about datanode requiring privileged resources to start check your port are < 1024 in yarn-site.xml.
 - Make sure you only changed the ports indicated in the instructions to be < 1024.
 - Make sure core-site.xml is configured to use kerberos.
 - Check keytab and principal entries in site xml, keytab directory owner/group is correct.
 - To inspect keytab files use: klist -e -k -t pathtokeytab.
 - Check that you modified hadoop-env.sh and /etc/init.d/hadoop-hdfs-datanode properly.
 - If the above checks are OK, run service hadoop-hdfs-datanode start and look at the output.
 - If there are no printed errors or you see an error stating that no VM can be found, it is a JSVC problem, see "Building and Installing JSVC" on page 105.
- Unable to find principal:
 - Check keytab and principal entries in site xml, keytab dir permissions.
- Unable to get password for username:
 - Check keytab and principal entries in site xml, keytab directory permissions. If these all look OK then run:

```
kinit -k -t
/etc/security/keytab/servicename.service.keytab
```

Troubleshooting 120

You should get no errors (just a prompt back). If there is an error check that the principal and keytab are correct.

- Check to make sure you used -norandkey when creating keytab files.
- Node manager will not start:
 - Login failure due to policy error exceptions in logs (typically seen as a remote exception to node manager for resource manager): check /usr/lib/gphd/hadoop/etc/hadoop/hadoop-policy.xml and replace any occurrences of \${HADOOP_HDFS_USER} with hdfs and \${HADOOP_YARN_USER} with yarn.

Troubleshooting 121

A. Creating a YUM EPEL Repository

Pivotal Command Center and Pivotal HD Enterprise expect some prerequisite packages to be pre-installed on each host, depending on the software that gets deployed on a particular host. In order to have a smoother installation it is recommended that each host would have yum access to an EPEL yum repository. If you have access to the Internet, then you can configure your hosts to have access to the external EPEL repositories. However, if your hosts do not have Internet access (or you are deploying onto a large cluster), then having a local yum EPEL repo would be highly recommended. This will also give you some control on the package versions you want deployed on your cluster.

Following are the steps to create a local yum repo:

- 1. Mount the RHEL/CentOS DVD on a machine that will act as the local yum repo
- **2.** Install a webserver on that machine (e.g. httpd), making sure that HTTP traffic can reach this machine
- **3.** Install the following packages on the machine:

```
yum-utils
createrep
```

4. Go to the directory where the DVD is mounted and run the following command: # createrepo .

5. Create a repo file on each host with a descriptive filename in the /etc/yum.repos.d/ directory of each host (for example, CentOS-6.1.repo) with the following contents:

```
[CentOS-6.1]
name=CentOS 6.1 local repo for OS RPMS
baseurl=http://172.254.51.221/centos/$releasever/os/
$basearch/
enabled=1
gpgcheck=1
gpgkey=http://172.254.51.221/centos/$releasever/os/$basearch/
RPM-GPG-KEY-CentOS-6
```

6. Validate that you can access the local yum repos by running the following command:

Yum list

A. Installing Binaries

This appendix provides instructions for installing and running the following Pivotal HD 1.0.3 components from downloaded tar files.

- Hadoop
- Zookeeper
- HBase
- Hive
- Pig
- Mahout
- Sqoop
- Flume
- Oozie

The installation instructions provided here are intended only as a Quick Start guide that start the services on one single host. Refer to Apache Hadoop documentation for information about other installation configurations.

```
http://hadoop.apache.org/docs/r2.0.5-alpha/
```

Notes:

- Pivotal HD and Pivotal HDMR1 should not be installed on the same cluster.
- Packages should come from same Distro (either PHD or PHDMR1)

Prerequisites

To install the Hadoop component (cluster install):

1. Add the hadoop user and then switch the login name to hadoop. All packages should be installed in the hadoop account.

```
useradd hadoop
passwd hadoop
su - hadoop
```

2. Make sure JAVA JDK 1.6, is installed on the system and append the following environment variables to ~/.bashrc

```
# export JAVA_HOME
export JAVA_HOME=/usr/java/default
```

3. Make sure the ~/.bashrc file takes effect:

```
$ source ~/.bashrc
```

4. SSH (both client and server) command is required for operations. Configure SSH to support connections that do not require passwords:

```
$ ssh-keygen -t rsa -P '' -f ~/.ssh/id_rsa
$ cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys
$ ssh hostname0

# copy authorized_keys to all hosts(hostname1, hostname2
etc.) in the cluster using scp
# NOTE: if an authorized_keys file already exists for # the user, rename your file authorized_keys2
$ scp /home/hadoop/.ssh/authorized_keys
hostname1:/home/hadoop/.ssh/
# Set the permissions on the file on all hosts
$ ssh hostname1
$ chmod 0600 ~/.ssh/authorized keys
```

Hadoop

1. Unpack the Hadoop tarball hadoop-2.0.5-alpha-gphd-2.0.3.0.tar.gz and and append the following environment variables to ~/.bashrc:

```
export HADOOP_HDFS_HOME=/path/to/hadoop
export PATH=$HADOOP_HOME/bin:$PATH
```

2. Make sure the ~/.bashrc file takes effect:

```
$ source ~/.bashrc
```

In the sections below, all the shell commands, unless explicitly specified, are all run from this $$HADOOP\ HOME$.

HDFS setup

1. Modify the file \$HADOOP_HOME/etc/hadoop/core-site.xml by adding the following to the configuration section:

2. Modify the file \$HADOOP_HOME/etc/hadoop/hdfs-site.xml by adding the following to the configuration section:

3. Format the HDFS namenode directory using default configurations:

```
$ bin/hdfs namenode -format
```

The default location for storing the namenode data is:

/tmp/hadoop-hadoop/dfs/name/

- **4.** Start the NameNode:
 - \$ sbin/hadoop-daemon.sh start namenode
- **5.** Start the DataNode:
 - \$ sbin/hadoop-daemon.sh start datanode
- **6.** Once the NameNode and DataNode are started, you can access the HDFS dashboard: http://localhost:50070/. You can also perform some tests from the command line:

```
$ bin/hdfs dfs -ls /
$ bin/hdfs dfs -mkdir -p /user/hadoop
# you can see a full list of hdfs dfs command options
$ bin/hdfs dfs
# put a local file to hdfs
$ bin/hdfs dfs -copyFromLocal /etc/passwd /user/hadoop/
```

- **7.** To stop the DataNode:
 - \$ sbin/hadoop-daemon.sh stop datanode
- **8.** To stop the NameNode:
 - \$ sbin/hadoop-daemon.sh stop namenode

Note: The DataNode or NameNode need to be started for the examples below.

YARN setup (Yarn distro only)

YARN

1. Modify the configuration file \$HADOOP_HOME/etc/hadoop/yarn-site.xml by adding the following to the configuration section:

```
<value>$HADOOP_CONF_DIR,$HADOOP_COMMON_HOME/share/hadoop/
common/*,$HADOOP_COMMON_HOME/share/hadoop/common/lib/*,$H
ADOOP_HDFS_HOME/share/hadoop/hdfs/*,$HADOOP_HDFS_HOME/sha
re/hadoop/hdfs/lib/*,$HADOOP_YARN_HOME/share/hadoop/yarn/
*,$HADOOP_YARN_HOME/share/hadoop/yarn/lib/*</value>
</property>
```

2. Start the YARN ResourceManager:

```
$ sbin/yarn-daemon.sh start resourcemanager
You can access the ResourceManager dashboard at: http://localhost:8088/
```

3. Start the YARN Node Manager:

```
$ sbin/yarn-daemon.sh start nodemanager
```

You can access the NodeManager dashboard at: http://localhost:8042/

You are now ready to perform YARN operations.

4. To stop YARN:

```
$ sbin/yarn-daemon.sh stop nodemanager
$ sbin/yarn-daemon.sh stop resourcemanager
```

YARN to run MapReduce

1. Modify the file \$HADOOP_HOME/etc/hadoop/mapred-site.xml (you can copy it from \$HADOOP_HOME/etc/hadoop/mapred-site.xml.template) by adding the following to the configuration section:

- **2.** Make sure the YARN ResourceManager and NodeManager are running (if not, refer to the ""YARN setup (Yarn distro only)" on page 125").
- **3.** If you want to track mapreduce history, start the MapReduce history daemon:

```
$ sbin/mr-jobhistory-daemon.sh start historyserver
You can access the server dashboard at: http://localhost:19888/
```

4. Run some MapReduce examples:

Example One:

```
$ cd $HADOOP_HOME
$ bin/hadoop jar
share/hadoop/mapreduce/hadoop-mapreduce-examples-2.0.5-alpha
```

```
-gphd-2.0.3.0.jar pi 2 10000
```

This command will submit the MapReduce task to calculate the Pi value.

Example Two:

```
$ bin/hadoop jar
share/hadoop/mapreduce/hadoop-mapreduce-examples-2.0.5-alpha
-gphd-2.0.3.0.jar
```

This shows you a list of example programs you can run.

When the job is running, you can view the application progress at the ResourceManager dashboard.

MapReduce V1 (MR1) Setup and Run (MR1 Distro only)

1. Unpack the MR1 tarball hadoop-mr1-1.0.3-gphd-2.0.3.0.tar.gzand and append the following environment variables to ~/.bashrc:

```
# Add HADOOP_HOME to the path
export HADOOP_HOME=/path/to/hadoop-mr1
PATH=$HADOOP_HOME/bin:$PATH
```

2. Make sure the ~/.bashrc file takes effect:

```
$ source ~/.bashrc
```

Note: When you trying to start HDFS service, you need to specify the full path of hadoop-daemon.sh, for example: \$HOME HDFS HOME/sbin/hadoop-daemon.sh.

3. Edit the files under

hadoop-mr1/tar/hadoop-mr1-1.0.3-gphd-2.0.3.0/conf/ directory and setup HDFS according to the Prerequisites and HDFS setup sections.

4. Modify the file \$HADOOP HOME/conf/mapred-site.xml.

- **5.** Ensure you have already started HDFS services by checking you can access http://localhost:50070 to see HDFS NameNode page, if you cannot, see the HDFS setup section for how to start the services.
- **6.** Create basic directory on HDFS:

```
$ $HADOOP_HDFS_HOME/bin/hdfs dfs -mkdir /tmp
$ $HADOOP_HDFS_HOME/bin/hdfs dfs -chmod -R 1777 /tmp
$ $HADOOP_HDFS_HOME/bin/hdfs dfs -mkdir -p /user/hadoop
$ $HADOOP_HDFS_HOME/bin/hdfs dfs -mkdir -p /mapred/system
```

7. Start Job Tracker and Task Tracker:

```
$ $HADOOP_HOME/bin/hadoop-daemon.sh start jobtracker
$ $HADOOP HOME/bin/hadoop-daemon.sh start tasktracker
```

8. You can access the map/reduce administration page here: http://localhost:50030.

9. Run an example Map/Reduce job to ensure that MR1 is working:

```
$ cd $HADOOP_HOME
$ bin/hadoop jar hadoop-examples-1.0.3-gphd-2.0.3.0.jar pi 2
100
<you should see job succeeded>
$ exit
```

If you can see that the job finished successfully, you have setup PHDMR1 correctly, otherwise, check your configuration files and ensure all HDFS services, job tracker and task tracker all started successfully.

Zookeeper

1. Unpack the Zookeeper tarball zookeeper-3.4.5-gphd-2.0.3.0.tar.gz and and append the following environment variables to ~/.bashrc:

```
export ZK_HOME=/path/to/zookeeper
PATH=$PATH:$ZK_HOME/bin
```

2. Make sure the ~/.bashrc file takes effect:

```
$ source ~/.bashrc
```

3. Go to the folder \$ZK HOME/conf, then run:

```
$ cd $ZK_HOME/conf
$ cp zoo_sample.cfg zoo.cfg
```

Note that as we are running Zookeeper in one node, there is no need to change the configuration file.

4. Start Zookeeper:

```
$ cd $ZK_HOME
$ bin/zkServer.sh start
```

5. Confirm that Zookeeper is running properly by running the following test:

```
$ cd $ZK_HOME
$ bin/zkCli.sh
> create /zk_test my_data
> get /zk_test
> quit
```

6. Stop the ZooKeeper server:

```
$ cd $ZK_HOME
$ bin/zkServer.sh stop
```

HBase

Following is an example of installing an instance of HBase that is running in pseudo-distributed mode. There is also an option to install a standalone or fully distributed HBase. Refer to Apache HBase documentation for information about other installation configurations. http://hbase.apache.org/book/book.html.

To install HBase

- 1. Unpack the HBase tar file hbase-0.94.8-gphd-2.0.3.0.tar.gz, the extracted folder is referred as \$HBASE HOME.
- **2.** Edit the following parameter values in \$HBASE_HOME/conf/hbase-site.xml, as follows:

```
<configuration>
cproperty>
  <name>hbase.rootdir</name>
   <value>hdfs://localhost:8020/hbase</value>
</property>
cproperty>
   <name>hbase.cluster.distributed
  <value>true</value>
  <description>mode: fully distributed, not to manage
   zookeeper</description>
</property>
cproperty>
   <name>hbase.zookeeper.quorum</name>
  <value>localhost</value>
</property>
</configuration>
```

3. Edit hbase-env.sh to turn off the HBase management.

```
HBASE MANAGES ZK=false
```

4. Give HBase access to HDFS.

HBase has the hadoop jars in the \$HBASE_HOME/lib directory. If you are using the 2.0.5-alpha-gphd-2.0.3.0 version of Hadoop, you can omit this step. If not, do the following:

- a. Delete the \$HBASE HOME/lib/hadoop-*.jar file.
- **b.** Copy the \$HADOOP HOME/*/hadoop-.jar file to \$HBASE HOME/lib/.
- **5.** Before starting HBase, make sure ZooKeeper server is running.
- **6.** Start HBase:

```
$HBASE HOME/bin/start-hbase.sh
```

Note: You can check the status of HBase at the following location:

http://localhost:60010

7. Confirm that HBase is installed and running properly by conducting the following test:

```
$ cd $HBASE_HOME
$ bin/hbase shell
hbase(main):003:0> create 'test', 'cf'
hbase(main):003:0> list 'test'
hbase(main):004:0> put 'test', 'row1', 'cf:a', 'value1'
hbase(main):005:0> put 'test', 'row2', 'cf:b', 'value2'
hbase(main):006:0> put 'test', 'row3', 'cf:c', 'value3'
hbase(main):007:0> scan 'test'
hbase(main):008:0> get 'test', 'row1'
hbase(main):012:0> disable 'test'
hbase(main):013:0> drop 'test'
```

8. Stop HBase:

\$HBASE HOME/bin/stop-hbase.sh

Hive

1. Unpack the Hive tarball hive-0.11.0-gphd-2.0.3.0.tar.gz and append the following environment variables to ~/.bashrc:

```
export Hive_HOME=/path/to/hive
export PATH=$Hive_HOME/bin:$PATH
export CLASSPATH=$Hive HOME/lib:$CLASSPATH
```

2. Make sure the ~/.bashrc file takes effect:

```
$ source ~/.bashrc
```

3. Create /tmp and /user/hive/warehouse (also known as hive.metastore.warehouse.dir) and set permissions to: chmod g+w in HDFS:

```
$ $HADOOP_HOME/bin/hadoop fs -mkdir /tmp
$ $HADOOP_HOME/bin/hadoop fs -mkdir /user/hive/warehouse
$ $HADOOP HOME/bin/hadoop fs -chmod g+w /tmp
```

```
$ $HADOOP HOME/bin/hadoop fs -chmod g+w /user/hive/warehouse
```

4. Test Hive.

```
hive> CREATE TABLE pokes (foo INT, bar STRING);
hive> LOAD DATA LOCAL INPATH './examples/files/kv1.txt'
OVERWRITE INTO TABLE pokes;
hive> SELECT a.* FROM pokes a where foo=400;
```

HCatalog

HCatalog is contained in the same tarball as Hive.

1. After you extracted tarball hive-0.11.0-gphd-2.0.3.0.tar.gz, append the following environment variables to ~/.bashrc:

```
export HCAT_HOME=$HIVE_HOME/hcatalog
export HCAT_PREFIX=$HCAT_HOME
export HIVE_CONF_DIR=$HIVE_HOME/conf
export HADOOP LIBEXEC DIR=$HADOOP HOME/libexec
```

2. Make sure the ~/.bashrc file take effect:

```
$ source ~/.bashrc
```

3. Run some HCatalog commands to verify your setup is correct. You should see similar output as shown below. Note that we have ommitted some trivial output for clarity:

```
$ cd $HCAT HOME
$ bin/hcat -e "create table pokes (foo int, bar string)"
Time taken: 9.625 seconds
$ bin/hcat -e "show tables"
OK
pokes
Time taken: 7.783 seconds
$ bin/hcat -e "describe pokes"
OK
foo
                      int
                                             None
bar
                         string
                                                  None
Time taken: 7.301 seconds
$ bin/hcat -e "alter table pokes add columns (new col int)"
\cap K
Time taken: 7.003 seconds
$ bin/hcat -e "describe pokes"
OK
                      int
foo
                                             None
                                             None
bar
                      string
new col
                      int
                                             None
```

```
Time taken: 7.014 seconds

$ bin/hcat -e "drop table pokes"

OK

Time taken: 9.78 seconds

$ exit
```

WebCatalog (OPTIONAL)

1. After you have installed HCatalog, manually copy the configuration file:

```
$ cp $HCAT_HOME/etc/webhcat/webhcat-default.xml
$HIVE_CONF_DIR/webhcat-site.xml
```

2. Edit the file you just copied (webhcat-site.xml) as follows:

Note that the "..." in above snipped refers to the original value of the property. You need to append two more variable name to value of this property.

3. Start WebCatalog service:

```
$ cd $HCAT_HOME
$ sbin/webhcat server.sh start
```

Note that starting WebCatalog service will write something under the current directory, so ensure the current user has permission to write in current directory.

4. Test WebCatalog:

```
$ curl
http://localhost:50111/templeton/v1/ddl/database/?user.name=
hadoop
```

5. Stop WebCatalog service:

```
$ cd $HCAT_HOME
$ sbin/webhcat server.sh stop
```

Pig

1. Unpack the Hive tarball pig-0.10.1-gphd-2.0.3.0.tar.gz and append the following environment variables to ~/.bashrc:

```
export PIG_HOME=/path/to/pig
export PATH=$PIG_HOME/bin:$PATH
export CLASSPATH=$PIG HOME/lib:$CLASSPATH
```

2. Make sure the ~/.bashrc file take effect:

```
$ source ~/.bashrc
```

3. Test Pig

```
[hadoop@localhost ~]$ pig
2013-09-10 02:00:05,753 [main] INFO org.apache.pig.Main -
Apache Pig version 0.10.1-qphd-2.0.3.0 (r: unknown) compiled
Sep 04 2013, 15:00:02
2013-09-10 02:00:05,753 [main] INFO org.apache.pig.Main -
Logging error messages to:
/home/hadoop/pig 1378792805751.log
2013-09-10 02:00:06,184 [main] INFO
org.apache.pig.backend.hadoop.executionengine.HExecutionEngi
ne - Connecting to hadoop file system at:
hdfs://localhost:8020
2013-09-10 02:00:06,185 [main] WARN
org.apache.hadoop.conf.Configuration -
mapred.used.genericoptionsparser is deprecated. Instead, use
mapreduce.client.genericoptionsparser.used
2013-09-10 02:00:06,191 [main] WARN
org.apache.hadoop.conf.Configuration - fs.default.name is
deprecated. Instead, use fs.defaultFS
2013-09-10 02:00:07,115 [main] WARN
org.apache.hadoop.conf.Configuration - fs.default.name is
deprecated. Instead, use fs.defaultFS
grunt>
```

Mahout

1. Unpack the Mahout mahout-distribution-0.7-gphd-2.0.3.0.tar.gz and append the following environment variables to ~/.bashrc:

```
export MAHOUT_HOME=/path/to/mahout
export PATH=$MAHOUT_HOME/bin:$PATH
export CLASSPATH=$MAHOUT_HOME/lib:$CLASSPATH
```

2. Make sure the ~/.bashrc file takes effect:

```
$ source ~/.bashrc
```

3. Test Mahout:

Note: Make sure hdfs and mapreduce services are running.

```
1. wget
http://archive.ics.uci.edu/ml/databases/synthetic_control/synthetic_control.data
2. $HADOOP_HDFS_HOME/bin/hdfs dfs -mkdir -p
/user/hadoop/testdata
3. $HADOOP_HDFS_HOME/bin/hdfs dfs -put
synthetic_control.data testdata
4. $MAHOUT_HOME/bin/mahout
org.apache.mahout.clustering.syntheticcontrol.kmeans.Job
5. $HADOOP_HDFS_HOME/bin/hdfs dfs -ls -R output
```

Sqoop

- **1.** Install and Deploy MySQL.
- **2.** Unpack the Sqoop

```
sqoop-1.4.2-gphd-2.0.3.0.bin__hadoop-2.0.5-alpha-gphd-2.0.3.0.t
ar.gz and append the following environment variables to ~/.bashrc:
export SQOOP_HOME=/path/to/sqoop
export PATH=$SQOOP_HOME/bin:$PATH
export CLASSPATH=$SQOOP_HOME/lib:$CLASSPATH
```

3. Make sure the ~/.bashrc file takes effect:

```
$ source ~/.bashrc
```

4. Move mysql-connector-java.jar to

/usr/share/java/mysql-connector-java.jar and add a softlink to sqoop's lib folder:

```
$ ln -sf /usr/share/java/mysql-connector-java.jar
$SQOOP HOME/lib/mysql-connector-java.jar
```

5. Create user hadoop in mysql, and grant all privileges to that user.

```
$ mysql -u root [-p]
mysql> insert into mysql.user(Host,User,Password)
values("%","hadoop",password("hadoop"));
mysql> GRANT ALL PRIVILEGES ON *.* TO 'hadoop'@'%'
identified by 'hadoop';
mysql> flush privileges;
```

6. Start Service:

```
$ service mysqld start
```

Flume

1. Unpack the Mahout mahout-distribution-0.7-gphd-2.0.3.0.tar.gz and append the following environment variables to ~/.bashrc:

```
export FLUME HOME=/path/to/flume
```

2. Make sure the ~/.bashrc file takes effect:

```
$ source ~/.bashrc
```

3. Set up an example configuration

```
# example.conf: A single-node Flume configuration
# Name the components on this agent
al.sources = r1
al.sinks = k1
al.channels = c1
```

Describe/configure the source

```
al.sources.rl.type = netcat
al.sources.rl.bind = localhost
al.sources.rl.port = 44444

# Describe the sink
al.sinks.kl.type = logger

# Use a channel which buffers events in memory
al.channels.cl.type = memory
al.channels.cl.capacity = 1000
al.channels.cl.transactionCapacity = 100

# Bind the source and sink to the channel
al.sources.rl.channels = cl
al.sinks.kl.channel = cl
```

4. Run the example:

```
$ cd $FLUME_HOME
$ bin/flume-ng agent --conf-file example.conf --name al
-Dflume.root.logger=INFO,console
(note: on the above command, "al" refers to the agent name
set in file example.conf)
```

Oozie

- **1.** Download and unpack Apache Tomcat 6.0.37 package.
- **2.** Download ext-2.2 package from http://extjs.com/deploy/ext-2.2.zip and unzip.
- **3.** Unpack the Oozie oozie-3.3.2-gphd-2.0.3.0-distro.tar.gz and append the following environment variables to ~/.bashrc:

```
export CATALINA_HOME=/path/to/tomcat
export OOZIE_HOME=/path/to/oozie
```

4. Make sure the ~/.bashrc file take effect:

```
$ source ~/.bashrc
```

5. Setup Oozie

```
$ $00ZIE_HOME/bin/oozie-setup.sh -hadoop 2.x $HADOOP_HOME
-extjs /path/to/ext2.2
```

6. Create the Oozie Database

```
$ $00ZIE_HOME/bin/ooziedb.sh create -sqlfile oozie.sql -run
Validate DB Connection
```

7. Add the following configuration to the Hadoop core-site.xml, then restart hadoop:

8. Copy the hadoop configuration to Ooze:

\$ cp \$HADOOP_HOME/etc/hadoop/* \$OOZIE_HOME/conf/hadoop-conf

9. Start Oozie:

\$ \$00ZIE HOME/bin/oozied.sh start

10. Run an Oozie example

```
$ cd $OOZIE_HOME
$ tar xvf oozie-examples.tar.gz
$ sed -e
's/jobTracker=localhost:8021/jobTracker=localhost:8032/'
examples/apps/map-reduce/job.properties > temp; cp temp
examples/apps/map-reduce/job.properties
$ hdfs dfs -mkdir -p /user/hadoop/examples
$ hdfs dfs -put examples/ /user/hadoop/
$ ./bin/oozie job -oozie http://localhost:11000/oozie
-config examples/apps/ssh/job.properties -run
```