

Pivotal Extension Framework

Version 2.1

Installation and User Guide

Rev: Ao2

© 2013 GoPivotal, Inc.

Table of Contents

| 1 | Pivo | otal Extension Framework Installation and Administration | 5 |
|---|------|---|----|
| | 1.1 | About PXF | 5 |
| | | 1.1.1 Prerequisites | |
| | | 1.1.2 Upgrading from GPXF | 5 |
| | 1.2 | Installing PXF | 6 |
| | | 1.2.1 Installing the PXF JAR File | 6 |
| | | 1.2.2 Setting up the Java Classpath | 6 |
| | | 1.2.3 Enabling the REST Service | 7 |
| | | 1.2.4 Restarting the Cluster | 7 |
| | | 1.2.5 Secure PXF | 7 |
| | | 1.2.6 Built-in Profiles | 9 |
| | 1.3 | Accessing HDFS File Data with PXF | 10 |
| | | 1.3.1 Syntax | 10 |
| | | 1.3.2 FORMAT clause | 11 |
| | | 1.3.3 Fragmenter | 11 |
| | | 1.3.4 Accessor | 11 |
| | | 1.3.5 Resolver | 12 |
| | | 1.3.6 Additional Options | 13 |
| | | 1.3.7 About the Public Directory | 14 |
| | | 1.3.8 Record key in key-value file formats | 14 |
| | | 1.3.9 Customized Writable Schema File Guidelines | 15 |
| | 1.4 | Accessing Hive Data with PXF | 16 |
| | | 1.4.1 Syntax | 16 |
| | | 1.4.2 Hive Command Line | 16 |
| | | 1.4.3 Mapping Hive Collection Types | 17 |
| | | 1.4.4 Partition Filtering | 17 |
| | 1.5 | Accessing HBase Data with PXF | 19 |
| | | 1.5.1 Syntax | 19 |
| | | 1.5.2 Column Mapping | 20 |
| | | 1.5.3 Row Key | 20 |
| | 1.6 | Accessing GemFire Data with PXF | 21 |
| | | 1.6.1 Syntax | 22 |
| | 1.7 | Troubleshooting | 22 |
| 2 | Pivo | otal Extension Framework External Table and API Reference | 24 |
| | 2.1 | | |
| | 2.2 | About the Java Class Services and Formats | 24 |
| | | 2.2.1 Fragmenter | 28 |
| | | 2.2.2 Accessor | 29 |
| | | 2.2.3 Resolver | 31 |
| | | 2.2.4 Analyzer | 34 |
| | 2.3 | About Custom Profiles | 36 |
| | 2.4 | About Query Filter Push-Down | 37 |
| | | | |

| | 2/11 | Filter Availability and Ordering | 37 |
|-----|-------|----------------------------------|----|
| | | | |
| | 2.4.2 | Creating a Filter Builder Class | 38 |
| | 2.4.3 | Filter Operations | 38 |
| | | Sample Implementation | |
| | 2.4.5 | Using Filters | 43 |
| 2.5 | Refer | ence | 44 |
| | 2.5.1 | External Table Examples | 44 |
| | 2.5.2 | Plugin Examples | 46 |
| | 2.5.3 | Configuration Files | 52 |



Copyright © 2013 GoPivotal, Inc. All rights reserved.

GoPivotal, Inc. believes the information in this publication is accurate as of its publication date. The information is subject to change without notice. THE INFORMATION IN THIS PUBLICATION IS PROVIDED "AS IS." GOPIVOTAL, INC. ("Pivotal") MAKES NO REPRESENTATIONS OR WARRANTIES OF ANY KIND WITH RESPECT TO THE INFORMATION IN THIS PUBLICATION, AND SPECIFICALLY DISCLAIMS IMPLIED WARRANTIES OF MERCHANTABILITY OR FITNESS FOR A PARTICULAR PURPOSE.

Use, copying, and distribution of any Pivotal software described in this publication requires an applicable software license.

All trademarks used herein are the property of Pivotal or their respective owners.

Use of Open Source

This product may be distributed with open source code, licensed to you in accordance with the applicable open source license. If you would like a copy of any such source code, Pivotal will provide a copy of the source code that is required to be made available in accordance with the applicable open source license. Pivotal may charge reasonable shipping and handling charges for such distribution.



1 Pivotal Extension Framework Installation and Administration

1.1 About PXF

PXF is an extensible framework that allows HAWQ to query external system data. PXF includes built-in connectors for accessing data that exists inside HDFS files, Hive tables, and HBase tables. Users can also create their own connectors to other parallel data stores or processing engines. To create these connectors using JAVA plugins, see the *Pivotal Extension Framework API and Reference Guide*.

1.1.1 Prerequisites

Check that the following systems are installed and running before you install PXF:

- HAWQ
- Pivotal Hadoop (PHD)
- Hadoop File System (HDFS) with REST service enabled on the Namenode and all the Datanodes.
- Hive: Check that the Hive Metastore service is available and running everywhere and that you have set the property *hive.matastore.uri* in the *hive-site.xml* file on the Namenode.
- HBase
- When configuring Secure (Kerberized) HDFS, NameNode port must be 8020 for PXF to function.
 Limitation will be removed next release.

1.1.2 Upgrading from GPXF

If you have a previous version of HAWQ, or GPXF installed, see the Pivotal ADS 1.1.3 Release Notes, for instructions about how to upgrade to HAWQ 1.1.x.

Please note the following:

- 1. DROP the tables created using GPXF and CREATE them again using PXF.
- 2. When you CREATE tables for PXF, remember to perform the following:
 - 1. Change the protocol name in the LOCATION clause from gpxf to pxf.
 - 2. Ensure that Fragmenter, Accessor, and Resolver are always specified for the table.
 - 3. Check that you have the new names for the Fragmenter, Accessor, and Resolver classes. See the *Pivotal Extension Framework API and Reference Guide*.
- 3. Check that you are using the correct gucs for PXF:

```
gpxf_enable_filter_pushdown -> pxf_enable_filter_pushdown
gpxf_enable_stat_collection -> pxf_enable_stat_collection
gpxf_enable_locality_optimizations -> pxf_enable_locality_optimizations
```



1.2 Installing PXF



Note

This topic describes how you can install PXF as a separate component, if you did not install it using the Pivotal Hadoop (HD) Enterprise Command Line Interface.

This topic contains the following information:

- Installing the PXF JAR File
- Setting up the Java Classpath
- Enabling REST service
- Restarting the Hadoop Cluster
- Testing PXF with HDFS Files
- Installing PXF to Work with Hive
- Testing Hive and PXF Integration

The following steps are required to install and configure PXF on a PHD 1.1.1 cluster.



Notes

- HBase steps are only required for PXF using HBase.
- Hive steps are only required for PXF for Hive.
- All steps must be performed on all nodes, unless noted differently.

1.2.1 Installing the PXF JAR File

Install the PXF JAR files on all nodes in the cluster:

```
sudo rpm -i pxf-2.1.1-x.rpm
```

- The pxf-2.1.1-x.rpm resides in the Pivotal ADS/HAWQ stack file. For example, PADS-1.1.1-x/pxf-2.1.1-x.noarch.rpm
- The script installs the JAR files at the default location at /usr/lib/gphd/pxf-2.1.1. A softlink will be created in /usr/lib/gphd/pxf

1.2.2 Setting up the Java Classpath

Append the following lines to the /etc/gphd/hadoop/conf/hadoop-env.sh on all nodes in the cluster:



```
export GPHD_ROOT=/usr/lib/gphd
export HADOOP_CLASSPATH=\
$GPHD_ROOT/pxf/pxf.jar:\
$GPHD_ROOT/publicstage:\
$GPHD_ROOT/pxf/avro-1.5.4.jar:\
$GPHD_ROOT/pxf/avro-mapred-1.5.4.jar:\
/usr/lib/gphd/zookeeper/zookeeper.jar:\
/usr/lib/gphd/hbase/hbase.jar:\
/etc/gphd/hbase/conf:\
$GPHD_ROOT/hive/lib/hive-service-0.11.0-gphd-2.1.0.0.jar:\
$GPHD_ROOT/hive/lib/libthrift-0.9.0.jar:\
$GPHD_ROOT/hive/lib/hive-metastore-0.11.0-gphd-2.1.0.0.jar:\
$GPHD_ROOT/hive/lib/libfb303-0.9.0.jar:\
$GPHD_ROOT/hive/lib/hive-common-0.11.0-gphd-2.1.0.0.jar:\
$GPHD_ROOT/hive/lib/hive-common-0.11.0-gphd-2.1.0.0.jar:\
$GPHD_ROOT/hive/lib/hive-common-0.11.0-gphd-2.1.0.0.jar:\
$GPHD_ROOT/hive/lib/hive-exec-0.11.0-gphd-2.1.0.0.jar:\
```

This adds the following to the PHD 1.1.1 Java classpath:

- PXF JAR, and staging dir /usr/lib/gphd/publicstage (see "About the Public Directory" below).
- HBase JAR, configuration dir, zookeeper.
- Hive: hive-service, libthrift, hive-metastore, libfb303, hive-common, hive-exec.
- Avro: avro, avro-mapred JARs for Avro files and serialization support.

1.2.3 Enabling the REST Service

To enable REST service, edit the /etc/gphd/hadoop/conf/hdfs-site.xm/files on all the nodes:

1.2.4 Restarting the Cluster

- 1. From the Admin node, use the Pivotal Hadoop (HD) Manager Command Line Interface to start the cluster.
- 2. Optionally, you can examine the classpath with the following command:

```
massh /tmp/clientnodes verbose "hadoop classpath"
```

1.2.5 Secure PXF

You can use PXF HDFS tables on secure HDFS cluster.



Read, write and analyze of PXF tables on HDFS files are enabled.

No changes are required to existing PXF tables.

Requirements

- Both HDFS and YARN principals have to be created and properly configured.
- HDFS must use port 8020.
- Hawq should be configured to work in secure mode properly.

Limitations

HDFS Namenode port MUST be 8020. This is a limitation that will be fixed on next version.

HBase/Hive setups requiring authentication are currently not supported.

Common Errors

You might see this error when YARN isn't properly configured for Kerberized HDFS

```
ERROR: remote component error (500) from '...': Problem accessing /gpdb/v<X>/Fragmenter/getFragments. Reason: Server Error Caused by: java.io.IOException: Can't get Master Kerberos principal for use as renewer
```

You might see this error when NameNode port is not 8020

```
ERROR: fail to get filesystem credential for uri hdfs://<namenode>:8020/
```

Reading and Writing Data with PXF

PXF comes with a number of built-in connectors for reading data that exists inside HDFS files, Hive tables, and HBase tables and for writing data into HDFS files. These built-in connectors use the PXF extensible API. You can also use the extensible API to create your own connectors to any other type of parallel data store or processing engine. See Pivotal Extension Framework External Table and API Reference for more information about the API.

This topic contains the following information:

- Accessing HDFS File Data with PXF (Read + Write)
- Accessing HIVE Data with PXF (Read only)
- Accessing HBASE Data with PXF (Read only)
- Accessing GemFireXD Data with PXF (Read only)



1.2.6 Built-in Profiles

A profile is a collection of common metadata attributes. Use the convenient and simplified PXF syntax.

PXF comes with a number of built-in profiles which group together a collection of metadata attributes to achieve a common goal:

| Profile | Description | Fragmenter/Accessor/Resolver |
|----------------|---|---|
| HdfsTextSimple | Read or write delimited single line records from or to plain text files on HDFS. | HdfsDataFragmenterLineBreakAccessorStringPassResolver |
| HdfsTextMulti | Read delimited single or multi line records (with quoted linefeeds) from plain text files on HDFS. It is not splittable (non parallel) and therefore reading is slower than reading with HdfsTextSimple. | HdfsDataFragmenterQuotedLineBreakAccessorStringPassResolver |
| Hive | Use this when connected to Hive. | HiveDataFragmenterHiveAccessorHiveResolver |
| HBase | Use this when connected to an HBase data store engine. | HBaseDataFragmenterHBaseAccessorHBaseResolver |
| Avro | Reading Avro files (i.e fileName.avro). | HdfsDataFragmenterAvroFileAccessorAvroResolver |
| GemFireXD | Use this when connected connected to GemFireXD | GemFireXDFragmenterGemFireXDAccessor |

Adding and Updating Profiles

Administrators can add new profiles or edit the built-in profiles inside *pxf-profiles.xml* (and apply with the Pivotal Hadoop (HD) Enterprise Command Line Interface). You can use the all the profiles in *pxf-profiles.xml*

.



Each profile has a mandatory unique name and an optional description.

In addition, each profile contains a set of plugins that are an extensible set of metadata attributes.

Custom Profile Example

1.3 Accessing HDFS File Data with PXF



Notes

- Pivotal recommends that you test PXF on HDFS before connecting to Hive or HBase.
- PXF on secure HDFS clusters requires NameNode to be configured on port 8020
- HBase/Hive configurations requiring user authentication are not supported

The syntax for accessing an HDFS file is as follows:

1.3.1 Syntax

```
CREATE [READABLE|WRITABLE] EXTERNAL TABLE <tbl name> (<attr list>)

LOCATION ('pxf://<name node hostname:50070>/<path to file or directory>?Profile=<chosen
profile>[&<additional options>=<value>]')

FORMAT '[TEXT | CSV | CUSTOM]' (<formatting properties>);

SELECT ... FROM <tbl name>; --to read from hdfs with READABLE table.

INSERT INTO <tbl name> ...; --to write to hdfs with WRITABLE table.
```

To read the data in the files or to write based on the existing format, you need to select the FORMAT, Profile, or one of the classes.

This topic describes the following:



- FORMAT clause
- Fragmenter
- Accessor
- Resolver



Note

For more details about the API and classes, see the *Pivotal Extension Framework API and Reference Guide*.

1.3.2 FORMAT clause

To read data, use the following formats with any PXF connector:

- FORMAT 'TEXT': Use with plain delimited text files on HDFS.
- FORMAT 'CSV': Use with comma-separated value files on HDFS.
- **FORMAT 'CUSTOM'**: Use with other files, such as binary formats. Must always be used with built-in formatter '*pxfwritable_import*' (for read) or '*pxfwritable_export*' (for write).

1.3.3 Fragmenter

Always use either [HdfsTextSimple | HdfsTextMulti] Profile or an HdfsDataFragmenter for HDFS file data.



Note

For read tables, you must include a Profile or a Fragmenter in the table definition.

1.3.4 Accessor

The choice of an Accessor depends on the HDFS data file type.

Note: You must include a Profile or an Accessor in the table definition.

| File Type | Accessor | FORMAT clause | Comments |
|----------------------|-------------------|---|--------------|
| Plain Text delimited | LineBreakAccessor | FORMAT 'TEXT' (<format list="" param="">)</format> | Read + Write |



| File Type | Accessor | FORMAT clause | Comments |
|-------------------|-------------------------|--|---|
| Plain Text CSV | LineBreakAccessor | FORMAT 'CSV' (<format list="" param="">)</format> | LineBreakAccessor is parallel and faster. Use if each logical data row is a physical data line. Read + Write |
| | QuotedLineBreakAccessor | | QuotedLineBreakAccessor is slower and non parallel. Use if the data includes embedded (quoted) linefeed characters. Read Only |
| SequenceFile | SequenceFileAccessor | FORMAT 'CUSTOM' (formatter='pxfwritable_import') | Read + Write (use formatter='pxfwritable_export' for write) |
| AvroFile | AvroFileAccessor | FORMAT 'CUSTOM' (formatter='pxfwritable_import') | Read Only |

1.3.5 Resolver

Choose the Resolver format if data records are serialized in the HDFS file.

Note: You must include a Profile or a Resolver in the table definition.

| Record Serialization | Resolver | Comments |
|-------------------------|--------------|--|
| Avro | AvroResolver | Avro files include the record schema, Avro serialization can be used in other file types (e.g, Sequence File). For Avro serialized records outside an Avro file, include a schema file name (.avsc) in the url under the optional <i>Schema-Data</i> option. The schema file name must exist in the public stage directory. Deserialize Only (Read) |



| Record Serialization | Resolver | Comments |
|-------------------------|--------------------|---|
| Java Writable | WritableResolver | Include the name of the Java class that uses Writable serialization in the url under the optional <i>Schema-Data</i>. The class file must exist in the public stage directory (or in Hadoop's class path). Deserialize and Serialize (Read + Write) See Customized Writable Schema File Guidelines |
| None (plain text) | StringPassResolver | Does not serialize plain text records. The database parses plain records. Passes records as they are. Deserialize and Serialize (Read + Write) |

1.3.6 Additional Options

| Option Name | Description |
|-------------------|--|
| COMPRESSION_CODEC | Useful for WRITABLE PXF tables. Specifies the c ompression codec class name for compressing the written data. The class must implement the org.apache.hadoop.io.compress.CompressionCodec interface. Some valid values are org.apache.hadoop.io.compress.DefaultCodec, org.apache.hadoop.io.compress.GzipCodec, org.apache.hadoop.io.compress.BZip2Codec. Note: org.apache.hadoop.io.compress.BZip2Codec is runs in a single thread and can be slow. This option has no default value. When the option is not defined, no compression will be done. |
| COMPRESSION_TYPE | Useful WRITABLE PXF tables with SequenceFileAccessor. Ignored when COMPRESSION_CODEC is not defined. Specifies the compression type for sequence file. Valid options are: RECORD - only the value part of each row is compressed. BLOCK - both keys and values are collected in 'blocks' separately and compressed. Default value: RECORD. |



| Option Name | Description |
|-------------------|--|
| SCHEMA-DATA | The data schema file used to create and read the HDFS file. For example, you could create an avsc (for Avro), or a java class (for Writable Serialization) file. Check that the file exists on the public directory (see About the Public Directory). This option has no default value. |
| THREAD-SAFE | Determines if the table query can run in multithread mode or not. When set to FALSE requests will be handled in a single thread. Should be set when a plugin or other element that are not thread safe are used (e.g. compression codec). Allowed values: TRUE, FALSE. Default value is TRUE - requests can run in multithread mode. |
| <custom></custom> | Any option that is desired to add to the pxf URI string will be accepted and passed along with its value to the Fragmenter, Accessor, Analyzer and Resolver implementations |

1.3.7 About the Public Directory

PXF provides a space to store your customized serializers and schema files on the filesystem. You must add schema files on all the datanodes and restart the cluster. The RPM creates the directory at the default location: /usr/lib/gphd/publicstage.

Ensure that all HDFS users have read permissions to HDFS services and limit write permissions to specific users.

1.3.8 Record key in key-value file formats

For sequence file and other file formats that store rows in a key-value format, the key value can be accessed through HAWQ using the saved keyword '*recordkey*' as a field name.

The field type must correspond to the key type, the same as the other fields must match the HDFS data.

WritableResolver supports read and write of recordkey, which can be of the following Writable hadoop types: BooleanWritable, ByteWritable, IntWritable, DoubleWritable, FloatWritable, LongWritable, Text.

If the *recordkey* field is not defined, the key is ignored in read, and a default value (segment id as LongWritable) is written in write.

Example

Let's say we have a data schema Babies.class containing 3 fields: (name text, birthday text, weight float).



An external table must include these three fields, and can either include or ignore the recordkey.

```
-- writable table with recordkey

CREATE WRITABLE EXTERNAL TABLE babies_registry (recordkey int, name text, birthday text, weight float)

LOCATION

('pxf://namenode_host:50070/babies_1940s?ACCESSOR=SequenceFileAccessor&RESOLVER=WritableResolver&D.
'CUSTOM' (formatter='pxfwritable_export');

INSERT INTO babies_registry VALUES (123456, "James Paul McCartney", "June 18, 1942", 3.800);

-- writable table without recordkey

CREATE WRITABLE EXTERNAL TABLE babies_registry2 (name text, birthday text, weight float)

LOCATION

('pxf://namenode_host:50070/babies_1940s?ACCESSOR=SequenceFileAccessor&RESOLVER=WritableResolver&D.
'CUSTOM' (formatter='pxfwritable_export');

INSERT INTO babies_registry VALUES ("Richard Starkey", "July 7, 1940", 4.0); -- this record's key will have some default value
```

The same goes for reading data from an existing file of a key-value format, e.g. a Sequence file.

```
-- readable table with recordkey

CREATE EXTERNAL TABLE babies_1940 (recordkey int, name text, birthday text, weight float)

LOCATION

('pxf://namenode_host:50070/babies_1940s?ACCESSOR=SequenceFileAccessor&RESOLVER=WritableResolver&D'CUSTOM' (formatter='pxfwritable_import');

SELECT * FROM babies_1940; -- retrieves each record's key
-- readable table without recordkey

CREATE EXTERNAL TABLE babies_1940_2 (name text, birthday text, weight float)

LOCATION

('pxf://namenode_host:50070/babies_1940s?ACCESSOR=SequenceFileAccessor&RESOLVER=WritableResolver&D'CUSTOM' (formatter='pxfwritable_import');

SELECT * FROM babies_1940_2; -- ignores the records' key
```

1.3.9 Customized Writable Schema File Guidelines

When using a WritableResolver, a schema file needs to be defined. The file needs to be a java class file and must be on the class path of PXF.

The class file must follow the following requirements:

- 1. Must implement org.apache.hadoop.io.Writable interface.
- 2. WritableResolver uses reflection to recreate the schema and populate its fields (for both read and write). Then it uses the Writable interface functions to read/write.
 - Therefore fields must be public to enable access to them. Private fields will be ignored.
- 3. Fields are accessed and populated by the order in which they are declared in the class file.
- 4. Supported field types:
 - String, int, double, float, long, short, boolean, byte array.
 - Array of any of the above types is supported, but the constructor must define the array size so the reflection will work.



1.4 Accessing Hive Data with PXF



NOTE

Check the following before adding PXF support on Hive:

- You are running the Hive Metastore service on a machine in your cluster.
- Check that you have set the hive.metastore.uris property in the hive-site.xml on the Namenode.
- The Hive JAR files are installed on the cluster nodes.

See Setting up the Java Classpath.

1.4.1 Syntax

Hive tables are always defined in a specific way in PXF, regardless of the underlying file storage format. The PXF Hive plugins automatically detect source tables:

- Text based
- SequenceFile
- RCFile
- ORCFile

The source table can also be a combination of these types. The PXF Hive plugin uses this information to query the data in runtime. The following PXF table definition is valid for any file storage type.

```
CREATE EXTERNAL TABLE hivetest(id int, newid int)

LOCATION ('pxf://<NN host>:50070/<hive table name>?PROFILE=Hive')

FORMAT 'custom' (formatter='pxfwritable_import');

SELECT * FROM hivetest;
```

Note: 50070 as noted in the example above is the REST server port on the HDFS NameNode. If a different port is assigned in your installation, use that port.

1.4.2 Hive Command Line

To start the Hive command line and work directly on a Hive table:

```
/>${HIVE_HOME}/bin/hive
```

Here's an example of how to create and load data into a sample Hive table from an existing file.



```
Hive> CREATE TABLE test (name string, type string, supplier_key int, full_price double) row format delimited fields terminated by ',';
Hive> LOAD DATA local inpath '/local/path/data.txt' into table test;
```

1.4.3 Mapping Hive Collection Types

PXF supports Hive data types that are not primitive types. For example :

```
CREATE TABLE sales_collections (
   item STRING,
   price FLOAT,
   properties ARRAY<STRING>,
   hash MAP<STRING, FLOAT>,
   delivery_address STRUCT<street:STRING, city:STRING, state:STRING, zip:INT>
)

ROW FORMAT DELIMITED FIELDS

TERMINATED BY '\001' COLLECTION ITEMS TERMINATED BY '\002' MAP KEYS TERMINATED BY '\003' LINES

TERMINATED BY '\n' STORED AS TEXTFILE;

LOAD DATA LOCAL INPATH '/local/path/<some data file>' INTO TABLE sales_collection;
```

To query a Hive table schema similar to the one in the example you need to define the PXF external table with attributes corresponding to members in the Hive table array and map fields. For example:

```
CREATE EXTERNAL TABLE gp_sales_collections(
 item_name TEXT,
 item_price REAL,
  property_type TEXT,
 property_color TEXT,
 property_material TEXT,
 hash_key1 TEXT,
 hash_val1 REAL,
 hash_key2 TEXT,
 hash_val3 REAL,
  delivery_street TEXT,
  delivery_city TEXT,
 delivery_state TEXT,
 delivery_zip INTEGER
LOCATION ('pxf://<namenode_host>:50070/sales_collections?PROFILE=Hive')
FORMAT 'custom' (FORMATTER='pxfwritable_import');
```

1.4.4 Partition Filtering

The PXF Hive plugin uses the Hive partitioning feature and directory structure. This enables partition exclusion on HDFS files that contain the Hive table. To use the Partition Filtering feature to reduce network traffic and I/O, run a PXF query using a WHERE clause that refers to a specific partition in the partitioned Hive table.

To take advantage of PXF Partition filtering push down, name the partition fields in the external table. These names must be the same as those stored in the Hive table. Otherwise, PXF ignores Partition filtering and the filtering is performed on the HAWQ side, impacting performance.



NOTE

The Hive plugin only filters on partition columns but not on other table attributes.

Example

Create a Hive table sales_part with 2 partition columns - delivery_state and delivery_city:

```
CREATE TABLE sales_part (name STRING, type STRING, supplier_key INT, price DOUBLE)

PARTITIONED BY (delivery_state STRING, delivery_city STRING)

ROW FORMAT DELIMITED FIELDS TERMINATED BY ',';
```

Load data into this Hive table and add some partitions:

```
LOAD DATA LOCAL INPATH '/local/path/data1.txt' INTO TABLE sales_part PARTITION(delivery_state = 'CALIFORNIA', delivery_city = 'San Francisco');

LOAD DATA LOCAL INPATH '/local/path/data2.txt' INTO TABLE sales_part PARTITION(delivery_state = 'CALIFORNIA', delivery_city = 'Sacramento');

LOAD DATA LOCAL INPATH '/local/path/data3.txt' INTO TABLE sales_part PARTITION(delivery_state = 'NEVADA' , delivery_city = 'Reno');

LOAD DATA LOCAL INPATH '/local/path/data4.txt' INTO TABLE sales_part PARTITION(delivery_state = 'NEVADA' , delivery_city = 'Las Vegas');
```

The Hive storage directory should appears as follows:

```
/hive/warehouse/sales_part/delivery_state=CALIFORNIA/delivery_city='San Francisco'/data1.txt
/hive/warehouse/sales_part/delivery_state=CALIFORNIA/delivery_city=Sacramento/data2.txt
/hive/warehouse/sales_part/delivery_state=NEVADA/delivery_city=Reno/data3.txt
/hive/warehouse/sales_part/delivery_state=NEVADA/delivery_city='Las Vegas'/data4.txt
```

To define a PXF table to read this Hive table and take advantage of partition filter push down, define the fields corresponding to the Hive partition fields at the end of the attribute list.

When defining an external table, check that the fields corresponding to the Hive partition fields are at the end of the column list. In HiveQL, issuing a *select** statement on a partitioned table shows the partition fields at the end of the record.



```
CREATE EXTERNAL TABLE pxf_sales_part(
  item_name TEXT,
  item_type TEXT,
  supplier_key INTEGER,
  item_price DOUBLE PRECISION,
  delivery_state TEXT,
  delivery_city TEXT
)
LOCATION ('pxf://namenode_host:50070/sales_part?Profile=Hive')
FORMAT 'custom' (FORMATTER='pxfwritable_import');
SELECT * FROM pxf_sales_part;
```

Example

In the following example the HAWQ query filters the *delivery_city* partition *Sacramento*. The filter on *item_name* is not pushed down since it is not a partition column. It is performed on the HAWQ side after all the data on *Sacramento* is transferred for processing.

```
SELECT * FROM pxf_sales_part WHERE delivery_city = 'Sacramento' AND item_name = 'shirt';
```

Example

The following HAWQ query reads all the data under *delivery_city* partition *CALIFORNIA*, regardless of the city partition.

```
SELECT * FROM pxf_sales_part WHERE delivery_state = 'CALIFORNIA'
```

1.5 Accessing HBase Data with PXF



NOTE

Before using PXF HBase plugin, verify the following:

- That you have set the Hadoop-env.sh on the Namenode.
- All Datanodes have the *hbase.jar*, *zookeeper.jar* and the *HBase conf* directory set in the *HADOOP_CLASSPATH*.

See Installing PXF for more information.

1.5.1 Syntax



To query an HBase table use the following syntax:

```
CREATE EXTERNAL TABLE <pxf tblname> (<col list - see details below>)

LOCATION ('pxf://<NN REST host>:<NN REST port>/<HBase table name>?PROFILE=HBase')

FORMAT 'CUSTOM' (FORMATTER='pxfwritable_import');

SELECT * FROM <pxf tblname>;
```

1.5.2 Column Mapping

Most HAWQ external tables (PXF or others) require that the HAWQ table attributes match the source data record layout, and include all the available attributes. However, use the PXF HBase plugin to specify the subset of HBase qualifiers to define the HAWQ PXF table. To set up a clear mapping between each attribute in the PXF table and a specific qualifier in the HBase table, you can use either:

- Directmapping
- Indirect mapping

In addition, the HBase row key is handled in a special way.

1.5.3 Row Key

You can use the HBase table row key in several ways. For example, you can see them using query results, or you can run a WHERE clause filter on a range of row key values. To use the row key in the HAWQ query, define the HAWQ table with the reserved PXF attribute *recordkey*. This attribute name tells PXF to return the record key in any key-value based system and in HBase.



NOTE

Since HBase is byte and not character-based Pivotal recommends that you define the *recordkey* as type *bytea*. This may result in better ability to filter data and increase performance.

```
CREATE EXTERNAL TABLE <tname> (recordkey bytea, ... ) LOCATION ('pxf:// ...')
```

Direct Mapping

Use Direct Mapping to map HAWQ table attributes to HBase qualifiers. You can specify the HBase qualifier names of interest, with column family names included, as quoted values.

For example, you have defined an HBase table called *hbase_sales* with multiple column families and many qualifiers. To see the following in the resulting attribute section of the CREATE EXTERNAL TABLE:

- rowkey
- qualifier saleid in the column family cf1



qualifier comments in the column family cf8

```
CREATE EXTERNAL TABLE hbase_sales (
recordkey bytea,
"cfl:saleid" int,
"cf8:comments" varchar
) ...
```

The PXF HBase plugin uses these attribute names as-is and returns the values of these HBase qualifiers.

Indirect Mapping (via Lookup Table)

Direct mapping method is fast and intuitive, but using indirect mapping helps to reconcile HBase qualifier names with HAWQ behavior:

- HBase qualifier names that are longer than 32 characters. HAWQ has a 32 character limit on attribute name size.
- HBase qualifier names can be binary or non-printable. HAWQ attribute names are character based.

In either case, Indirect Mapping uses a lookup table on HBase. You can create the lookup table to store all necessary lookup information. This works as a template for any future queries. The name of the lookup table must be *pxflookup* and must include the column family named *mapping*.

Using the sales example in Direct Mapping, if our *rowkey* represents the HBase table name and the *mapping* column family includes the actual attribute mapping in the key value form of *<hawq attr name>=<hbase cf:qualifier>*.

Example

| (row key) | mapping |
|-----------|-------------------|
| sales | id=cf1:saleid |
| sales | cmts=cf8:comments |

Note: The mapping assigned new names for each qualifier. You can use these names in your HAWQ table definition:

```
CREATE EXTERNAL TABLE hbase_sales (
recordkey bytea
id int,
cmts varchar
) ...
```

PXF automatically matches HAWQ to HBase column names when a pxflookup table exists in HBase.

1.6 Accessing GemFire Data with PXF



NOTE

Before using PXF GemFire plugin, verify the following:

- That you have installed the *gfxd rpm* on the Namenode and on the Datanodes.
- The namenode and all Datanodes have the sqlfire.jar set in the HADOOP_CLASSPATH.

See Installing PXF for more information.

1.6.1 Syntax

To query an GemFire table use the following syntax:

```
CREATE EXTERNAL TABLE <pxf tblname> (<col list>)

LOCATION ('pxf://<NN REST host>:<NN REST port>/<GemFire table name>?Profile=GemFireXD')

FORMAT 'CUSTOM' (FORMATTER='pxfwritable_import');

SELECT * FROM <pxf tblname>;
```

1.7 Troubleshooting

The following table describes some common errors while using PXF:

Table: PXF Errors and Explanation

Error

ERROR: invalid URI pxf://localhost:50070/demo/file1: missing options section

ERROR: protocol "pxf" does not exist

ERROR: remote component error: 0

DETAIL: There is no pxf servlet listening on the host and port specified in the external table url.

ERROR: Missing FRAGMENTER option in the pxf uri: pxf://localhost:50070/demo/file1?a=a

ERROR: remote component error: 500

DETAIL: Problem accessing /gpdb/v<X>/Fragmenter/getFragments. Reason: org.apache.hadoop.mapred.ln



ERROR: remote component error: 500

DETAIL: Problem accessing /gpdb/v<X>/Fragmenter/getFragments. Reason: org.apache.hadoop.mapred.lr.

files

ERROR: remote component error: 500 PXF not correctly installed in CLASSPATH

ERROR: GPHD component not found

ERROR: remote component error (500)

Problem accessing /gpdb/v<X>/Fragmenter/getFragments. Reason: java.io.IOExceptic

ERROR: fail to get filesystem credential for uri hdfs://<namenode>:8020/

HBase Specific Errors

ERROR: remote component error: 500

DETAIL: Problem accessing /gpdb/v<X>/Fragmenter/getFragments. Reason: org.apache.hadoop.hbase.clie

ERROR: remote component error: 500

DETAIL: Problem accessing /gpdb/v<X>/Fragmenter/getFragments. Reason: org.apache.hadoop.hbase.Tai

ERROR: remote component error: 500 (/HTTP/1.1 500 Internal Server Error

ERROR: remote component error: 500 (/HTTP/1.1 500 org/apache/zookeeper/KeeperException

ERROR: remote component error: 500 (/HTTP/1.1 500 Illegal HBase column name name

Hive Specific Errors

ERROR: remote component error: 500

DETAIL: Problem accessing /gpdb/v<X>/Fragmenter/getFragments. Reason: java.net.ConnectException: ConnectException: Conne

ERROR: remote component error: 500

DETAIL: Problem accessing /gpdb/v<X>/Fragmenter/getFragments. Reason: NoSuchObjectException(mes.



2 Pivotal Extension Framework External Table and API Reference

You can extend PXF functionality and add new services and formats using the Java API without changing HAWQ. The API includes the four classes Fragmenter, Accessor, Resolver, Analyzer. The Fragmenter, Accessor, and Resolver classes must be implemented to add a new service. The Analyzer class is optional.

2.1 Creating an External Table

Syntax for an EXTERNAL TABLE that uses the PXF protocol is as follows:.

```
CREATE EXTERNAL TABLE ext_table <attr list, ...>

LOCATION('pxf://<namenode>:<port>/path/to/data?FRAGMENTER=FragmenterForX&ACCESSOR=AccessorForX&REScustom user options>=<value>')FORMAT 'custom'(formatter='pxfwritable_import');
```

Where:

Table: Parameter values and description

| Parameter | Value and description |
|-------------------|--|
| namenode | The current host of the PXF service is HDFS Namenode port. |
| REST | Port for Namenode, 50070 by default. |
| path/to/data | A directory, file name, wildcard pattern, table name, etc |
| FRAGMENTER | The plugin (java class) to use for fragmenting data. Used in READABLE external tables only. |
| ACCESSOR | The plugin (java class) to use for accessing the data. Used in READABLE and WRITABLE tables. |
| RESOLVER | The plugin (java class) to use for serializing and deserializing the data. Used in READABLE and WRITABLE tables. |
| Custom Options | Anything else that is desired to add. Will be passed in runtime to the plugins indicated above. |

For more information about this example, see "About the Java Class Services and Formats" and the Pivotal Extension Installation and Administrator Guide.

2.2 About the Java Class Services and Formats

The Java class names you must include in the PXF URI are Fragmenter, Accessor, and Resolver. The Fragmenter class is mandatory for READABLE tables, and not supported for WRITABLE tables. Pivotal recommends that you reuse a previously-defined Accessor or Resolver data format.

All the attributes are passed from HAWQ as headers to the PXF Java service. The Java service retrieves the source data and converts it to a HAWQ-readable format. You can pass any additional information to the user-implemented services.

The example in "Creating an External Table" shows the available keys and associated values. The example also contains attributes that are passed in from the HAWQ side. The available keys and associated values are as follows:

```
FRAGMENTER: `FragmenterForX'

ACCESSOR: `AccessorForX'

RESOLVER: `ResolverForX'
```

These three Java plugins and the optional plugin, Analyzer, extend the com.pivotal.pxf.utilities.Plugin class.

The Java classes can be described as follows:

```
/*

* Base class for all plugin types (Accessor, Resolver, Fragmenter, Analyzer)

* Holds InputData as well (the meta data information used by all plugin types)

*/
public class Plugin
{

protected InputData inputData;

/*

 * C'tor

 */
public Plugin(InputData input);

/**

 * Checks if the plugin is thread safe or not, based on inputData.

 *

 * @return if plugin is thread safe or not

 */
public boolean isThreadSafe() {
 return true;
}

}
```

Attributes are available through the *com.pivotal.pxf.utilities.InputData* class. The following example shows how *inputData.getProperty('USERINFO1')* returns *optional_info*.

```
/*

* Common configuration of all MetaData classes

* Provides read-only access to common parameters supplied using system properties

*/
public class InputData
```

```
{
    /* Constructor of InputData
    * Parses greenplum.* configuration variables
   public InputData(Map<String, String> paramsMap);
     \star Expose the parameters map
   public Map<String, String> getParametersMap();
    \ /* \ {\tt Copy \ contructor \ of \ InputData}
    * Used to create from an extending class
   public InputData(InputData copy);
    * Returns a property as a string type
   public String getProperty(String property);
    * returns the number of segments in GP
   public int totalSegments();
    * returns the current segment ID
   public int segmentId();
    / \, ^{\star} returns the current outputFormat
    * currently either text or gpdbwritable
    * /
   public OutputFormat outputFormat();
     * returns the server name providing the service
   public String serverName();
    * returns the server port providing the service
   public int serverPort();
    * Returns true if there is a filter string to parse
    * /
    public boolean hasFilter();
    * The filter string
   public String filterString();
    * returns the number of columns in Tuple Description
```

```
public int columns();
* returns column index from Tuple Description
public ColumnDescriptor getColumn(int index);
 * returns a data fragment
public int getDataFragment();
\mbox{\scriptsize \star} returns the column descriptor of the recordkey column.
 * If the recordkey column was not specified by the user in the create table statement,
 * then getRecordkeyColumn will return null.
public ColumnDescriptor getRecordkeyColumn();
/* returns the path to the resource required
 * (might be a file path or a table name)
public String path();
/* returns the path of the schema used for various deserializers
* e.g, Avro file name, Java object file name.
public String srlzSchemaName() throws FileNotFoundException, IllegalArgumentException;
* returns the ClassName for the java class that handles the file access
public String accessor();
 * returns the ClassName for the java class that handles the record deserialization
public String resolver();
 ^{\star} The avroSchema fetched by the AvroResolver and used in case of Avro File
 {}^{*} In case of avro records inside a sequence file this variable will be null
 \mbox{\scriptsize \star} and the AvroResolver will not use it.
public Schema GetAvroFileSchema();
* Returns table name
* /
public String tableName();
* The avroSchema is set from the outside by the AvroFileAccessor
public void SetAvroFileSchema(Schema theAvroSchema);
 * Returns the compression codec (can be null - means no compression)
```



```
*/
public String compressCodec();
}
```

2.2.1 Fragmenter



Note

The Fragmenter Plugin reads data into HAWQ. Such tables are called READABLE PXF tables. The Fragmenter Plugin cannot write data out of HAWQ. Such tables are called WRITABLE tables.

The Fragmenter is responsible for passing datasource metadata back to HAWQ. It also returns a list of data fragments to the Accessor or Resolver. Each data fragment describes some part of the requested data set. It contains the datasource name, such as the file or table name, including the hostname where it is located. For example, if the source is a HDFS file, the Fragmenter returns a list of data fragments containing a HDFS file block. Each fragment includes the location of the block. If the source data is an HBase table, the Fragmenter returns information about table regions, including their locations.

The following implementations are shipped with PXF 1.x and higher:

```
HdfsDataFragmenter
HBaseDataFragmenter
HiveDataFragmenter
```

The fragmenter uses com.pivotal.pxf.fragmenters.FragmentsOutput to return information:

```
/*
  * Class holding information about fragments (FragmentInfo)
  */
public class FragmentsOutput
{
   public FragmentsOutput();
   public void addFragment(String sourceName, String[] hosts);
   public void addFragment(String sourceName, String[] hosts, String userData);
   public List<FragmentInfo> getFragments();
}
```

Any Fragmenter class needs to extend *com.pivotal.pxf.fragmenters.Fragmenter*.

com.pivotal.pxf.fragmenters.Fragmenter

```
/*
 * Interface that defines the splitting of a data resource into fragments that can be processed
in parallel
 * GetFragments returns the fragments information of a given path (source name and location of
each fragment).
 * Used to get fragments of data that could be read in parallel from the different segments.
 */
public abstract class Fragmenter extends Plugin
{
    protected FragmentsOutput fragments;

    public Fragmenter(InputData metaData)
    {
        super(metaData);
        fragments = new FragmentsOutput();
    }

    /*
        * returns the data fragments
        */
        public abstract FragmentsOutput GetFragments() throws Exception;
}
```

Class Description

getFragments() returns a string in a JSON format of the retrieved fragment. For example, if the input path is a HDFS directory, the source name for each fragment should include the file name including the path for the fragment.

2.2.2 Accessor

The Accessor retrieves specific fragments and passes records back to the Resolver. For example, the Accessor creates a *FileInputFormat* and a Record Reader for an HDFS file and sends this to the Resolver. In the case of HBase or Hive files, the Accessor returns single rows from an HBase or Hive table. PXF 1.x or higher contains the following implementations:

Table: Accessor base classes

| Accessor class | Description |
|------------------------|---|
| HdfsAtomicDataAccessor | Base class for accessing datasources which cannot be split. These will be accessed by a single HAWQ segment. QuotedLineBreakAccessor - Accessor for TEXT files that has records with embedded linebreaks |
| | |

| HdfsSplittableDataAccessor | Base class for accessing HDFS files using RecordReaders: | |
|----------------------------|---|--|
| | LineBreakAccessor - Accessor for TEXT files (replaced the deprecated TextFileAccessor, LineReaderAccessor) AvroFileAccessor - Accessor for Avro files | |
| HiveAccessor | Accessor for Hive tables | |
| HBaseAccessor | Accessor for HBase tables | |

The class needs to extend the com.pivotal.pxf.Plugin class, and implement one or both interface s:

- com.pivotal.pxf.accessors.lReadAccessor
- · com.pivotal.pxf.accessors.IWriteAccessor

```
* Internal interface that defines the access to data on the source
 * data store (e.g, a file on HDFS, a region of an HBase table, etc).
 * All classes that implement actual access to such data sources must
 * respect this interface
 * /
public interface IReadAccessor
   public boolean openForRead() throws Exception;
   public OneRow readNextObject() throws Exception;
   public void closeForRead() throws Exception;
}
 * An interface for writing data into a data store
 * (e.g, a sequence file on HDFS).
 * All classes that implement actual access to such data sources must
 * respect this interface
 * /
public interface IWriteAccessor
   public boolean openForWrite() throws Exception;
   public OneRow writeNextObject(OneRow onerow) throws Exception;
   public void closeForWrite() throws Exception;
}
}
```

The Accessor calls openForRead() to read existing data. After reading the data, it calls closeForRead(). readNextObject() and returns one of the following:

- a single record, encapsulated in a OneRow object
- null if it reaches EOF

The Accessor calls *openForWrite()* to write data out. After writing the data, it writes a *OneRow* object with *writeNextObject()*, and when done calls *closeForWrite()*. *OneRow* represents a key-value item.



com.pivotal.pxf.format.OneRow:

```
/*
\mbox{\scriptsize \star} Represents one row in the external system data store. Supports
* the general case where one row contains both a record and a
 * separate key like in the HDFS key/value model for MapReduce
* (Example: HDFS sequence file)
* /
public class OneRow
{
    * Default constructor
    public OneRow()
    * Constructor sets key and data
    public OneRow(Object inKey, Object inData)
    * Copy constructor
    public OneRow(OneRow copy)
     * Setter for key
    public void setKey(Object inKey)
    * Setter for data
    public void setData(Object inData)
    * Accessor for key
     * /
    public Object getKey()
     * Accessor for data
    public Object getData()
    * Show content
    public String toString()
}
```

2.2.3 Resolver



The Resolver deserializes records in the *OneRow* format and serializes them to a list of *OneField* objects. PXF converts a *OneField* object to a HAWQ-readable *GPDBWritable* format. PXF 1.x or higher contains the following implementations:

Table: Resolver base classes

| Resolver class | Description | |
|--------------------|---|--|
| StringPassResolver | Supports: | |
| | GPBWritable VARCHAR | |
| | StringPassResolver replaced the deprecated TextResolver. It passes whole records (composed of any data types) as strings without parsing them | |
| WritableResolver | Resolver for custom Hadoop Writable implementations. Custom class can be specified with the schema {{{}}} and supports the following: | |
| | GPDBWritable.BOOLEAN GPDBWritable.INTEGER GPDBWritable.BIGINT GPDBWritable.REAL GPDBWritable.FLOAT8 GPDBWritable.VARCHAR GPDBWritable.BYTEA | |
| AvroResolver | Supports the same field objects as WritableResolver. | |
| HBaseResolver | Supports the same field objects as <i>WritableResolver</i> and also supports the following: | |
| | GPDBWritable.SMALLINT GPDBWritable.NUMERIC GPDBWritable.TEXT GPDBWritable.BPCHAR GPDBWritable.TIMESTAMP | |
| HiveResolver | Supports the same field objects as <i>WritableResolver</i> and also supports the following: | |
| | GPDBWritable.SMALLINT GPDBWritable.TEXT GPDBWritable.TIMESTAMP | |

The class needs to extend the *com.pivotal.pxf.resolvers.Plugin class*, and implement one or both interface s:

- com.pxf.resolvers.lReadResolver
- com.pxf.resolvers.IWriteResolver

```
/*
 * Interface that defines the deserialization of one record brought from
 * the data Accessor. Every implementation of a deserialization method
 * (e.g, Writable, Avro, ...) must implement this interface.
 */
public interface IReadResolver
{
    public List<OneField> getFields(OneRow row) throws Exception;
}

/*
 * Interface that defines the serialization of data read from the DB
 * into a OneRow object.
 * Every implementation of a serialization method
 * (e.g, Writable, Avro, ...) must implement this interface.
 */
public interface IWriteResolver
{
    public OneRow setFields(DataInputStream inputStream) throws Exception;
}
```

Notes

- getFields should return a list of com.pivotal.pxf.format.OneField, each representing a single field.
- setFields should return a single OneRow object, given an input stream.

com.pivotal.pxf.format.OneField

```
/*
 * Defines one field on a deserialized record.
 * 'type' is in OID values recognized by GPDBWritable
 * 'val' is the actual field value
 */
public class OneField
{
    public OneField() {}
    public OneField(int Type, Object Val)
    {
        type = Type;
        val = Val;
    }
    public int type;
    public Object val;
}
```

The value of *Type* should follow the *com.pivotal.pxf.hadoop.io.GPDBWritable* type *enums*. *Val* is the appropriate Java class. Supported types are as follows:

Table: Resolver supported types

| GPDBWritable recognized OID | Field value |
|-----------------------------|-------------------------------------|
| GPDBWritable.SMALLINT | Short |
| GPDBWritable.INTEGER | Integer |
| GPDBWritable.BIGINT | Long |
| GPDBWritable.REAL | Float |
| GPDBWritable.FLOAT8 | Double |
| GPDBWritable.NUMERIC | String ("651687465135468432168421") |
| GPDBWritable.BOOLEAN | Boolean |
| GPDBWritable.VARCHAR | GPDBWritable.BPCHAR |
| | GPDBWritableTEXT - String |
| GPDBWritable.BYTEA | byte [] |
| GPDBWritable.TIMESTAMP | Timestamp |

2.2.4 Analyzer



The Analyzer provides PXF statistical data for the HAWQ query optimizer. For a detailed explanation about HAWQ statistical data gathering, see *ANAL YZE* in the *Pivotal ADS Administrator Guide*. Implement the PXF Analyzer for the HDFS text, sequence, and AVRO files. For HBase tables and Hive tables, the Analyzer returns default values.

Notes:

- The new *boolean guc pxf_enable_stat_collection* requests statistics. The default value is *on*. When you turn it off, the statistics collected reflect default values.
- Pivotal recommends that you implement the Analyzer to return an estimated result as fast as possible.

The class needs to extend com.pivotal.pxf.analyzers.Analyzer.

com.pivotal.pxf.analyzers.Analyzer

```
* Abstract class that defines getting statistics for ANALYZE.
 * GetEstimatedStats returns statistics for a given path
 * (block size, number of blocks, number of tuples (rows)).
 * Used when calling ANALYZE on a PXF external table, to get
 * table's statistics that are used by the optimizer to plan queries.
public abstract class Analyzer extends Plugin
    public Analyzer(InputData metaData)
        super(metaData);
    }
     * path is a data source name (e.g, file, dir, wildcard, table name)
     * returns the data statistics in json format
     * NOTE: It is highly recommended to implement an extremely fast logic
     * that returns *estimated* statistics. Scanning all the data for exact
     * statistics is considered bad practice.
    public String GetEstimatedStats(String data) throws Exception
        /* Return default values */
        return DataSourceStatsInfo.dataToJSON(new DataSourceStatsInfo());
}
```

GetEstimatedStats creates a DataSourceStatsInfo class, and returns the result DataSourceStatsInfo.dataToJSON.

com.pivotal.pxf.analyzers.DataSourceStatsInfo

```
* DataSourceStatsInfo is a public class that represents the size
 * information of given path.
public class DataSourceStatsInfo
    public DataSourceStatsInfo(long blockSize, long numberOfBlocks, long numberOfTuples);
     * Default values
    public DataSourceStatsInfo();
     * Given a FragmentsSizeInfo, serialize it in JSON to be used as the result string for Hawq.
An example result is as
     * follows:
     * {"PXFFilesSize":{"blockSize":67108864,"numberOfBlocks":1,"numberOfTuples":5}}
    public static String dataToJSON(DataSourceStatsInfo info) throws IOException;
     * Given a size info structure, convert it to be readable. Intended
     * for debugging purposes only. 'datapath' is the data path
     * part of the original URI (e.g., table name, *.csv, etc).
    public static String dataToString(DataSourceStatsInfo info, String datapath);
    public long getBlockSize();
    public long getNumberOfBlocks();
   public long getNumberOfTuples();
}
```

2.3 About Custom Profiles

Administrators can add new profiles or edit the built-in profiles in

Each profile mandatory unique name, and an optional description.

In addition, each profile contains a set of plugins that are an extensible set of metadata attributes.



2.4 About Query Filter Push-Down

If a query includes a number of WHERE clause filters, HAWQ may push all or some queries to PXF. If pushed to PXF, the Accessor can use the filtering information when accessing the data source to fetch tuples. These filters only return records that pass filter evaluation conditions. This reduces data processing and reduces network traffic from the SQL engine.

This topic includes the following information:

- Filter Availability and Ordering
- Creating a Filter Builder class
- Filter Operations
- Sample Implementation
- Using Filters

2.4.1 Filter Availability and Ordering

PXF allows push-down filtering if the following rules are met:

- Only single expressions or a group of AND'ed expressions no OR'ed expressions.
- Only expressions of supported data types and operators. See the *Pivotal Extension Framework Installation and Administration Guide* for more information.

FilterParser scans the pushed down filter list and uses the user's build() implementation to build the filter.

- For simple expressions (e.g, a >= 5) FilterParser places column objects on the left of the expression and therefore constants on the right.
- For compound expressions (e.g <expression> AND <expression>) it handles three cases in the build() function:
- 1. Simple Expression: <Column Index> <Operation> <Constant>
- 2. Compound Expression: <Filter Object> AND <Filter Object>
- 3. Compound Expression: <List of Filter Objects> AND <Filter Object>

2.4.2 Creating a Filter Builder Class

To check if a filter queried PXF, call the InputData hasFilter() function:

```
/*
 * Returns true if there is a filter string to parse
 */
public boolean hasFilter()
{
   return filterStringValid;
}
```

If hasFilter() returns false, there is no filter information. If it returns true, PXF parses the serialized filter string into a meaningful filter object to use later. To do so, create a filter builder class that implements the FilterParser.IFilterBuilder interface:

```
/*
 * Interface a user of FilterParser should implement
 * This is used to let the user build filter expressions in the manner she
 * sees fit
 *
 * When an operator is parsed, this function is called to let the user decide
 * what to do with it operands.
 */
interface IFilterBuilder
{
   public Object build(Operation operation, Object left, Object right) throws Exception;
}
```

While PXF parses the serialized filter string from the incoming HAWQ query, it calls the *build() interface* function. PXF calls this function for each condition or filter pushed down to PXF. Implementing this function returns some Filter object or representation that the Accessor or Resolver uses in runtime to filter out records. The *build()* function accepts an Operation as input, and left and right operands.

2.4.3 Filter Operations

```
/*
 * Operations supported by the parser
 */
public enum Operation
{
    HDOP_LT, //less than
    HDOP_GT, //greater than
    HDOP_LE, //less than or equal
    HDOP_GE, //greater than or equal
    HDOP_EQ, //equal
    HDOP_NE, //not equal
    HDOP_NE, //not equal
    HDOP_AND //AND'ed conditions
};
```

Filter Operands

There are three types of operands:

- Column Index
- Constant
- Filter Object

Column Index

```
/*
 * The class represents a column index
 * It used to know the type of an operand in the stack
 */
public class ColumnIndex
{
    private int index;

    public ColumnIndex(int idx)
    {
        index = idx;
    }

    public int index()
    {
        return index;
    }
}
```

Constant

```
/*
 * The class represents a constant object (String, Long, ...)
 * It used to know the type of an operand in the stack
 */
public class Constant
{
    private Object constant;

    public Constant(Object obj)
    {
        constant = obj;
    }

    public Object constant()
    {
        return constant;
    }
}
```

Filter Object

Filter Objects can be internal, such as those you define; or external, those that the remote system uses. For example, for HBase, you define the HBase *Filter* class (*org.apache.hadoop.hbase.filter.Filter*), while for Hive, you use an internal default representation created by the PXF framework, called *BasicFilter*. You can decide the filter object to use, including writing a new one. *BasicFilter* is the most common:

```
* Basic filter provided for cases where the target storage system does not provide it's own
  * For example: Hbase storage provides it's own filter but for a Writable based record in a
SequenceFile
 \mbox{\scriptsize *} there is no filter provided and so we need to have a default
static public class BasicFilter
    private Operation oper;
    private ColumnIndex column;
    private Constant constant;
      * C'tor
     public BasicFilter(Operation inOper, ColumnIndex inColumn, Constant inConstant)
         oper = inOper;
         column = inColumn;
         constant = inConstant;
     }
      * returns oper field
     public Operation getOperation()
     {
         return oper;
      * returns column field
     public ColumnIndex getColumn()
     {
         return column;
      * returns constant field
     public Constant getConstant()
         return constant;
     }
 }
```

2.4.4 Sample Implementation

Let's look at the following sample implementation of the filter builder class and its *build()* function that handles all 3 cases. Let's assume that BasicFilter was used to hold our filter operations

```
public class MyDemoFilterBuilder implements FilterParser.IFilterBuilder
   private InputData inputData;
   public MyDataFilterBuilder(InputData input)
       inputData = input;
    }
     * Translates a filterString into a FilterParser.BasicFilter or a list of such filters
   public Object getFilterObject(String filterString) throws Exception
       FilterParser parser = new FilterParser(this);
       Object result = parser.parse(filterString);
       if (!(result instanceof FilterParser.BasicFilter) && !(result instanceof List))
           throw new Exception("String " + filterString + " resolved to no filter");
       return result;
    }
    public Object build(FilterParser.Operation opId,
                       Object leftOperand,
                       Object rightOperand) throws Exception
       if (leftOperand instanceof FilterParser.BasicFilter)
           //sanity check
           if (opId != FilterParser.Operation.HDOP_AND || !(rightOperand instanceof
FilterParser.BasicFilter))
               throw new Exception("Only AND is allowed between compound expressions");
           //case 3
           if (leftOperand instanceof List)
               return handleCompoundOperations((List<FilterParser.BasicFilter>)leftOperand,
(FilterParser.BasicFilter)rightOperand);
           //case 2
           else
               return handleCompoundOperations((FilterParser.BasicFilter)leftOperand,
(FilterParser.BasicFilter)rightOperand);
       }
       //sanity check
       if (!(rightOperand instanceof FilterParser.Constant))
           throw new Exception("expressions of column-op-column are not supported");
       //case 1 (assume column is on the left)
       return handleSimpleOperations(opId, (FilterParser.ColumnIndex)leftOperand,
(FilterParser.Constant)rightOperand);
   FilterParser.ColumnIndex column,
                                                         FilterParser.Constant constant)
```

Here is an example of creating a filter builder class to implement the Filter interface, implement the *build()* function, and generate the Filter object. To do this, use either the Accessor, Resolver, or both to call the *getFilterObject* function:

```
if (inputData.hasFilter())
{
    String filterStr = inputData.filterString();
    DemoFilterBuilder demobuilder = new DemoFilterBuilder(inputData);
    Object filter = demobuilder.getFilterObject(filterStr);
    ...
}
```

2.4.5 Using Filters

Once you have built the Filter object(s), you can use them to read data and filter out records that do not meet the filter conditions:

- 1. Check whether you have a single or multiple filters.
- 2. Evaluate each filter and iterate over each filter in the list. Disqualify the record if a filter conditions fail.

```
if (filter instanceof List)
{
    for (Object f : (List)filter)
        <evaluate f>; //may want to break if evaluation results in negative answer for any
filter.
}
else
{
    <evaluate filter>;
}
```

Example of evaluating a single filter:

```
//Get our BasicFilter Object
FilterParser.BasicFilter bFilter = (FilterParser.BasicFilter)filter;

//Get operation and operator values
FilterParser.Operation op = bFilter.getOperation();
int colIdx = bFilter.getColumn().index();
String val = bFilter.getConstant().constant().toString();

//Get more info about the column if desired
ColumnDescriptor col = input.getColumn(colIdx);
String colName = filterColumn.columnName();

//Now evaluate it against the actual column value in the record...
```

2.5 Reference

This section contains the following information:

- External Table Samples
- Plugin Examples
- Configuration Files

2.5.1 External Table Examples

Example 1

Shows an external table that can analyze all *Sequencefiles* that are populated *Writable* serialized records and exist inside the hdfs directory *sales/2012/01*. *SaleItem.class* is a java class that implements the *Writable* interface and describes a java record that includes three class members.

Note: In this example the class member names do not necessarily match the database attribute names, but the types match. *SaleItem.class* must exist in the classpath of every Datanode.



```
CREATE EXTERNAL TABLE jan_2012_sales (id int, total int, comments varchar)
LOCATION
('pxf://10.76.72.26:50070/sales/2012/01/*.seq?fragmenter=HdfsDataFragmenter&accessor=SequenceFileA'custom' (formatter='pxfwritable_import');
```

Example 2

Shows an external table that can analyze an HBase table called *sales*. It has 10 column families *(cf1 - cf10)* and many qualifier names in each family. This example focuses on the *rowkey*, the qualifier *saleid* inside column family *cf1*, and the qualifier *comments* inside column family *cf8* and uses Direct Mapping:

```
CREATE EXTERNAL TABLE hbase_sales (hbaserowkey text, "cf1:saleid" int, "cf8:comments" varchar)

LOCATION
('pxf://10.76.72.26:50070/sales?fragmenter=HBaseDataFragmenter&accessor=HBaseAccessor&resolver=HBaseAccessor*);
```

Example 3

This example uses Indirect Mapping. Note how the attribute name changes and how they correspond to the HBase lookup table. Executing a *SELECT from my_hbase_sales*, the attribute names automatically convert to their HBase correspondents.

```
CREATE EXTERNAL TABLE my_hbase_sales (hbaserowkey text, id int, cmts varchar)

LOCATION
('pxf://10.76.72.26:8080/sales?fragmenter=HBaseDataFragmenter&accessor=HBaseAccessor&resolver=HBase
'custom' (formatter='pxfwritable_import');
```

Example 4

Shows an example for writable table of compressed data.

```
CREATE WRITABLE EXTERNAL TABLE sales_aggregated_2012 (id int, total int, comments varchar)
LOCATION
('pxf://10.76.72.26:8080/sales/2012/aggregated?accessor=LineBreakAccessor&resolver=StringPassResol
'TEXT';
```

Example 5

Shows an example for writable table into sequence file, using schema file. Note that for write, the formatter is *pxfwritable_export*.



```
CREATE WRITABLE EXTERNAL TABLE sales_max_2012 (id int, total int, comments varchar)
LOCATION
('pxf://10.76.72.26:8080/sales/2012/max?fragmenter=HdfsDataFragmenter&accessor=SequenceFileAccesso
'custom' (formatter='pxfwritable_export');
```

2.5.2 Plugin Examples

This section contains sample dummy implantations of all four plug-ins. It also contains a usage example.

Dummy Fragmenter

```
package com.pivotal.pxf.examples;
import java.net.InetAddress;
import com.pivotal.pxf.utilities.InputData;
import com.pivotal.pxf.utilities.Plugin;
import com.pivotal.pxf.fragmenters.Fragmenter;
import com.pivotal.pxf.fragmenters.FragmentsOutput;
public class DummyFragmenter extends Fragmenter
public DummyFragmenter(InputData metaData)
super(metaData);
}
/*
* path is a data source name (e.g, file, dir, wildcard, table name).
* returns the data fragments
public FragmentsOutput GetFragments() throws Exception
String localhostname = java.net.InetAddress.getLocalHost().getHostName();
fragments.addFragment(this.inputData.path() + ".1" /* source name */,
new String[]{localhostname} /*
available hosts list
fragments.addFragment(this.inputData.path() + ".2" /* source name */,
new String[]{localhostname} /*
available hosts list
*/);
fragments.addFragment(this.inputData.path() + ".3" /* source name */,
new String[]{localhostname} /*
available hosts list
*/);
return fragments;
}
}
```

Dummy Accessor

```
package com.pivotal.pxf.examples;

import com.pivotal.pxf.format.OneRow;
import com.pivotal.pxf.utilities.InputData;
import com.pivotal.pxf.utilities.Plugin;
import com.pivotal.pxf.accessors.IReadAccessor;

import com.pivotal.pxf.accessors.IWriteAccessor;
```

```
* Internal interface that defines the access to a file on HDFS. All classes
 ^{\star} that implement actual access to an HDFS file (sequence file, avro file,...)
 * must respect this interface
* Dummy implementation, for documentation
* /
public class DummyAccessor extends Plugin implements IReadAccessor
,IWriteAccessor
{
   private int rowNumber;
   private int fragmentNumber;
   private Log Log;
   public DummyAccessor(InputData metaData)
    super(metaData);
       rowNumber = 0;
       fragmentNumber = 0;
       Log = LogFactory.getLog(DummyAccessor.class);
    }
   public boolean openForRead() throws Exception
       /* fopen or similar */
       return true;
   public OneRow readNextObject() throws Exception
       /* check for EOF */
       if (fragmentNumber > 0)
           return null; /* signal EOF, close will be called */
       int fragment = this.inputData.getDataFragment();
       /* generate row */
       OneRow row = new OneRow(Integer.toString(fragment) + "." + Integer.toString(rowNumber),
/* key */
                              Integer.toString(rowNumber) + ",text," +
Integer.toString(fragment) /* value */);
       /* advance */
       rowNumber += 1;
       if (rowNumber == 2) {
           rowNumber = 0;
           fragmentNumber += 1;
       /* return data */
       return row;
   public void closeForRead() throws Exception
       /* fclose or similar */
    }
```

```
@Override
public boolean openForWrite() throws Exception {
    /* fopen or similar */
    return true;
    }

    @Override
    public boolean writeNextObject(OneRow onerow) throws Exception {
    Log.info(onerow.getData());
    return true;
    }

    @Override
    public void closeForWrite() throws Exception {
    /* fclose or similar */
    }
}
```

Dummy Resolver

```
package com.pivotal.pxf.examples;
import java.util.List;
import java.util.LinkedList;
import com.pivotal.pxf.format.OneField;
import com.pivotal.pxf.format.OneRow;
import com.pivotal.pxf.utilities.InputData;
import com.pivotal.pxf.utilities.Plugin;
import com.pivotal.pxf.resolvers.IReadResolver;
import com.pivotal.pxf.resolvers.IWriteResolver;
import com.pivotal.pxf.hadoop.io.GPDBWritable;
 ^{\star} Class that defines the deserializtion of one record brought from the external input data.
 * Every implementation of a deserialization method (Writable, Avro, BP, Thrift, ...)
 * must inherit this abstract class
 * Dummy implementation, for documentation
public class DummyResolver extends Plugin implements IReadResolver, IWriteResolver
   private int rowNumber;
    public DummyResolver(InputData metaData)
       super(metaData);
       rowNumber = 0;
    }
    public List<OneField> getFields(OneRow row) throws Exception
        /* break up the row into fields */
        List<OneField> output = new LinkedList<OneField>();
        String[] fields = ((String)row.getData()).split(",");
        output.add(new OneField(GPDBWritable.INTEGER /* type */,Integer.parseInt(fields[0]) /*
value */));
        output.add(new OneField(GPDBWritable.VARCHAR ,fields[1]));
        output.add(new OneField(GPDBWritable.INTEGER ,Integer.parseInt(fields[2])));
       return output;
    }
    @Override
    public OneRow setFields(DataInputStream inputStream) throws Exception {
       /* should read inputStream row by row */
       if (rowNumber > 5)
          return null;
       return new OneRow(null, new String("row number " + rowNumber++));
    }
```



Dummy Analyzer

```
package com.pivotal.pxf.examples;
import com.pivotal.pxf.utilities.InputData;
import com.pivotal.pxf.utilities.Plugin;
import com.pivotal.pxf.analyzers.Analyzer;
import com.pivotal.pxf.analyzers.DataSourceStatsInfo;
\mbox{*} Abstract class that defines getting statistics for ANALYZE.
* GetEstimatedStats returns statistics for a given path
* (block size, number of blocks, number of tuples).
\mbox{*} Used when calling ANALYZE on a PXF external table,
\mbox{\scriptsize \star} to get table's statistics that are used by the optimizer to plan queries.
* Dummy implementation, for documentation
public class DummyAnalyzer extends Analyzer
public DummyAnalyzer(InputData metaData)
super(metaData);
}
* path is a data source name (e.g, file, dir, wildcard, table name).
* returns the data statistics in json format
public String GetEstimatedStats(String data) throws Exception
return DataSourceStatsInfo.dataToJSON(new DataSourceStatsInfo(16 /* disk block size in bytes */,
3 /* number of disk blocks */,
2 /* number of rows in disk block (avg) */));
}
```

Usage Example



```
psql=# CREATE EXTERNAL TABLE dummy_tbl (intl integer, word text, int2 integer)
('pxf://localhost:50070/eran_location?FRAGMENTER=com.pivotal.pxf.examples.DummyFragmenter&ACCESSOR
format'custom' (formatter = 'pxfwritable_import');
CREATE EXTERNAL TABLE
psql=# SELECT * FROM dummy_tbl;
int1 | word | int2
0 | text | 0
1 | text | 0
0 | text | 0
1 | text | 0
0 | text | 0
1 | text | 0
(6 rows)
psql=# CREATE WRITABLE EXTERNAL TABLE dummy_tbl_write (int1 integer, word text, int2 integer)
('pxf://localhost:50070/eran_location?ACCESSOR=com.pivotal.pxf.examples.DummyAccessor&RESOLVER=com
format'custom' (formatter = 'pxfwritable_import');
CREATE EXTERNAL TABLE
psql=# INSERT INTO dummy_tbl_write VALUES (1, 'a', 11), (2, 'b', 22);
INSERT 0 2
```

2.5.3 Configuration Files

This section contains sample environment variable files for HDFS, HIVE, and HBase:

hadoop-env.sh

You can use this file to configure the following types of configurations:

- HDFS only
- HDFS and HBase
- · HDFS, HBase, and Hive

HDFS only



```
export GPHD_ROOT=/usr/lib/gphd
export HADOOP_CLASSPATH=\
$HADOOP_CLASSPATH:\
$GPHD_ROOT/pxf/pxf.jar:\
$GPHD_ROOT/publicstage:\
$GPHD_ROOT/publicstage:\
$GPHD_ROOT/pxf/avro-1.5.4.jar:\
$GPHD_ROOT/pxf/avro-mapred-1.5.4.jar:\
$GPHD_ROOT/pxf/avro-mapred-2.0.2_alpha_gphd_2_0_1_0/
hadoop-mapreduce-client-core-2.0.2-alpha-gphd-2.0.1.0-SNAPSHOT.jar:\
```

HDFS and **HBase**

```
export GPHD_ROOT=/usr/lib/gphd
export GPHD_CONF=/etc/gphd
export HADOOP_CLASSPATH=\
$HADOOP_CLASSPATH:\
$GPHD_ROOT/pxf/pxf.jar:\
$GPHD_ROOT/publicstage:\
$GPHD_ROOT/pxf/avro-1.5.4.jar:\
$GPHD_ROOT/pxf/avro-mapred-1.5.4.jar:\
$GPHD_ROOT/pxf/avro-mapred-1.5.4.jar:\
$GPHD_ROOT/zookeeper/zookeeper.jar:\
$GPHD_ROOT/hbase/hbase.jar:\
$GPHD_ROOT/hbase/conf:\
$GPHD_ROOT/hadoop-mapreduce-2.0.2_alpha_gphd_2_0_1_0/hadoop-mapreduce-client-core-2.0.2-alpha-gphd
```

HDFS, HBase, and Hive

```
export GPHD_ROOT=/usr/lib/gphd
export GPHD_CONF=/etc/gphd
export HIVELIB_ROOT=$GPHD_ROOT/hive/lib
export HADOOP_CLASSPATH=\
$HADOOP_CLASSPATH:\
$GPHD_ROOT/pxf/pxf.jar:\
$GPHD_ROOT/publicstage:\
$GPHD_ROOT/pxf/avro-1.5.4.jar:\
$GPHD_ROOT/pxf/avro-mapred-1.5.4.jar:\
$GPHD_ROOT/zookeeper/zookeeper.jar:\
$GPHD_ROOT/hbase/hbase.jar:\
$GPHD_CONF/hbase/conf:\
$HIVELIB_ROOT/hive-service-0.9.1-gphd-2.0.1.jar:\
$HIVELIB_ROOT/libthrift-0.7.0.jar:\
$HIVELIB_ROOT/hive-metastore-0.9.1-gphd-2.0.1.jar:\
$HIVELIB_ROOT/libfb303-0.7.0.jar:\
$HIVELIB_ROOT/hive-common-0.9.1-gphd-2.0.1.jar:\
$HIVELIB_ROOT/hive-exec-0.9.1-gphd-2.0.1.jar:\
$GPHD_ROOT/hadoop-mapreduce-2.0.2_alpha_gphd_2_0_1_0/hadoop-mapreduce-client-core-2.0.2-alpha-gphd
```

hbase-site.xml



You can use this file to configure the following types of configurations:

- HBase
- · HDFS, HBase and Hive

HBase

The Java Class path requires the PXF JAR file. *hbase-site.xml* needs to be configured to match the hbase settings on all nodes (Namenode and Datanodes).

```
export GPHD_ROOT=/usr/lib/gphd
export GPHD_CONF=/etc/gphd
export HADOOP_CLASSPATH=\
$HADOOP_CLASSPATH:\
$GPHD_ROOT/pxf/pxf.jar:\
$GPHD_ROOT/publicstage:\
$GPHD_ROOT/publicstage:\
$GPHD_ROOT/pxf/avro-1.5.4.jar:\
$GPHD_ROOT/pxf/avro-mapred-1.5.4.jar:\
$GPHD_ROOT/zookeeper/zookeeper.jar:\
$GPHD_ROOT/hbase/hbase.jar:\
```

HDFS, HBase and Hive

```
export GPHD_ROOT=/usr/lib/gphd
export HIVELIB_ROOT=$GPHD_ROOT/hive/lib
export HADOOP_CLASSPATH=\
$GPHD_ROOT/pxf/pxf.jar:\
$GPHD_ROOT/publicstage:\
$GPHD_ROOT/publicstage:\
$GPHD_ROOT/pxf/avro-1.5.4.jar:\
$GPHD_ROOT/pxf/avro-mapred-1.5.4.jar:\
$GPHD_ROOT/zookeeper/zookeeper.jar:\
$GPHD_ROOT/hbase/hbase.jar:\
```