

Pivotal Command Center

Version 2.2

User Guide

Rev: 01

© 2014 GoPivotal, Inc.

Table of Contents

1	1 PCC Overview	5
	1.1 About Pivotal Command Center	5
	1.1.1 Pivotal Command Center UI and CLI	6
	1.1.2 Performance Monitor (nmon)	6
	1.1.3 PostgreSQL Database	6
	1.2 Architectural Overview	7
2	2 Installing PCC	8
	2.1 Quick Start	8
	2.2 Supported Platforms and Browsers	10
	2.2.1 Platforms	10
	2.2.2 Browsers	10
	2.3 Requirements	11
	2.3.1 Prerequisites	11
	2.3.2 Product Downloads	12
	2.3.3 Package Accessibility	12
	2.4 Installation Instructions	13
	2.4.1 Installing Pivotal Command Center	13
	2.4.2 Copying PHD Service Files	15
	2.4.3 Importing PHD Services	15
	2.4.4 Launching Pivotal Command Center	16
	2.5 Uninstalling Pivotal Command Center	16
	2.6 Upgrading Pivotal Command Center	16
3	3 Using PCC	18
	3.1 Overview	18
	3.1.1 Status indicators	18
	3.2 Logging In	
	3.2.1 Login Screen	19
	3.3 Settings	20
	3.4 User Management	20
	3.4.1 Users	21
	3.4.2 Profiles	22
	3.4.3 Passwords	
	3.5 Cluster Status Page	22
	3.6 Configuring and Deploying a Cluster	
	3.6.1 Adding a Cluster	24
	3.6.2 Starting the Cluster	
	3.6.3 Initializing and Configuring HAWQ	28
	3.7 Starting Stopping and Uninstalling the Cluster	
	3.7.1 Starting	29
	3.7.2 Stopping	
	3.7.3 Uninstalling	
	3.8 Dashboard	

	3.9 Cluster Analysis	32
	3.10 MapReduce Job Monitor	33
	3.10.1 Job Details	34
	3.11 YARN App Monitor	
	3.12 HAWQ Query Monitor	
	3.13 Topology	
	3.13.1 Topology Actions	
	3.14 Logs	
	3.14.1 Filtering Logs	
	3.14.2 Viewing Logs	
4	Creating a YUM EPEL Repository	41
	Command Line Reference	
	5.1 Backup and Restore	
	5.1.1 Backup	42
	5.1.2 Restore	42



Copyright © 2014 Pivotal Software, Inc. All rights reserved.

GoPivotal, Inc. believes the information in this publication is accurate as of its publication date. The information is subject to change without notice. THE INFORMATION IN THIS PUBLICATION IS PROVIDED "AS IS." PIVOTAL SOFTWARE, INC. ("Pivotal") MAKES NO REPRESENTATIONS OR WARRANTIES OF ANY KIND WITH RESPECT TO THE INFORMATION IN THIS PUBLICATION, AND SPECIFICALLY DISCLAIMS IMPLIED WARRANTIES OF MERCHANTABILITY OR FITNESS FOR A PARTICULAR PURPOSE.

Use, copying, and distribution of any Pivotal software described in this publication requires an applicable software license.

All trademarks used herein are the property of Pivotal or their respective owners.

Use of Open Source

This product may be distributed with open source code, licensed to you in accordance with the applicable open source license. If you would like a copy of any such source code, Pivotal will provide a copy of the source code that is required to be made available in accordance with the applicable open source license. Pivotal may charge reasonable shipping and handling charges for such distribution.



1 PCC Overview

This document is a User Guide for the Pivotal Command Center (PCC) User Interface.

This section provides a brief overview of Pivotal Command Center, including:

- About Pivotal Command Center
 - Pivotal Command Center UI and CLI
 - Performance Monitor (nmon)
 - PostgreSQL Database
- Architectural Overview

1.1 About Pivotal Command Center

Pivotal Command Center (PCC) allows an administrative user to configure, deploy, monitor, and manage one or more Pivotal HD clusters. The Command Center has both a graphical user interface and command-line tools to deploy and configure, monitor, and administer Pivotal HD clusters.

- For UI operations, see Using PCC.
- For command line operations, see the Pivotal HD Enterprise Installation and Administrator Guide.



This release of Command Center allows administering and monitoring of only Pivotal HD Enterprise 2.x clusters.

PCC provides complete life cycle management for Pivotal HD Clusters by performing the following two main groups of functions:

- · Cluster configuration and deployment
- Cluster monitoring and management

These functions are served through a set of RESTful web services that run as a web application on Jetty server on the Command Center admin host. This is called . This web application stores its metadata and cluster configuration for Pivotal HD cluster nodes and services in the Pivotal Command Center PostgreSQL database. It makes use of a Puppet Server to perform most of its HD cluster installation and configuration. It also has a polling service that retrieves Hadoop metrics from the cluster and stores them in the Command Center PostgreSQL Database at periodic intervals.



1.1.1 Pivotal Command Center UI and CLI

The PCC UI provides the user with a single web-based graphical user interface to configure and deploy, monitor and manage one or more Pivotal HD cluster. This web application is hosted on a Ruby-on-Rails application which presents the status and metrics of the clusters. The system metrics data is gathered by the Performance Monitor (nmon) component. The Command Center UI invokes the APIs to retrieve all Hadoop-specific cluster metrics and status information. This includes the Hadoop metrics that was previously retrieved by the polling service.

PCC provides a command-line interface (CLI) for more advanced users to perform installation, configuration and uninstalls. This tool invokes the APIs to install and configure the various Pivotal HD services. The CLI also provides a way to perform other administrative actions such as starting and stopping clusters. For how to use this CLI, please refer to the *Pivotal HD Enterprise Installation and Administrator Guide*.

1.1.2 Performance Monitor (nmon)

Pivotal Command Center comes with a Performance Monitor called (for node monitor). This makes use of a highly scalable message passing architecture to gather performance metrics from each node that Command Center monitors. This consists of a master daemon that runs on the Command Center admin host and an daemon that runs on all the cluster nodes that report system metric information to the master. This includes metrics such as CPU, memory, disk I/O and network usage information.

The master on the admin host dumps the system metrics it receives from the agents on the cluster nodes into a PostgreSQL DB. This is then queried by the Command Center UI application to display its cluster analysis graphs.

The agents hosts are deployed throughout the cluster during Pivotal HD cluster deployment itself (see for details).

The agents are deployed as services on each host, including on the Pivotal Command Center admin host.

To stop or start the service run the following as root:

```
# service nmon stop
# service nmon start
```

1.1.3 PostgreSQL Database

Pivotal Command Center makes use of a PostgreSQL Database to store the following:

- Cluster configurations
- Hadoop cluster metrics
- · System metrics of the cluster

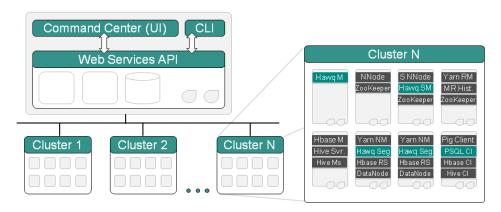


Pivotal Command Center Metadata

1.2 Architectural Overview

For more details about Pivotal HD Enterprise, refer to the Pivotal HD Installation and Administrator Guide.

Pivotal Command Center - Architecture





2 Installing PCC

This section describes how to install Pivotal Command Center and then use PCC to deploy and configure a Pivotal HD cluster.

- Quick Start
- Supported Platforms and Browsers
- Requirements
- Installation Instructions
- Uninstalling Pivotal Command Center
- Upgrading Pivotal Command Center

2.1 Quick Start

Below is a *Quick Start* table intended for users who are familiar with the product and various prerequisites; more detailed instructions are provided later in this section.

Note that shaded rows depict operations that you perform via a Browser/PCC UI.

Task	Sub-Task	
Requirements	Check JDK version is 1.7.0_15 or later	
See Requirements for more details.	(as root)	
	# java -version	
	If not, download it from Oracle.	
	Check Yum accessibility	
	(as root)	
	Verify that all hosts have yum access to an EPEL yum repository.	
	# yum list <list of="" packages=""></list>	

Install Pivotal Command Center

See Install PCC for more details.

(as root)

- 1. Create a directory (phd) for your PCC installation:
 - # mkdir phd
- 2. Copy tar file to your specified directory on the admin node, for example:

```
# scp ./PCC-2.2.x. version.build.os .x86_64.tar.gz
host:/root/phd/
```

- 3. Login as root and untar to that directory:
 - # cd /root/phd

```
# tar --no-same-owner -zxvf PCC-2.2.x.
```

version.build.os .x86_64.tar.gz

- 4. Run the installation script from the directory where it was extracted:
 - # ./install
- 5. As the rest of the installation is done as the <code>gpadmin</code> user, change to that user:

```
# su - gpadmin
```

6. Enable Secure Connections

Copy the PHD Services packages to the Admin Node

See Copying PHD
Services for more details.

(as gpadmin)

- Copy the Pivotal HD services (PHD, ADS (HAWQ)) tarballs from the initial download location to the gpadmin home directory (home/user/gpadmin).
- 2. Change the owner of the packages to gpadmin then untar the tarballs.

For example: If the file is a tar.gz or tgz, use:

tar -zxf packagename.tgz

If the file is a tar, use:

tar -xf packagename.tar

Import the PHD services

(as gpadmin)

See Importing PHD Services for more details.

For each service (PHD, ADS) you are importing, run the following:

icm_client import -s < PATH TO EXTRACTED PHD SERVICE
TAR BALL >



Browser:	Launch a browser and enter the host on which you installed PCC:
Launch Pivotal Command Center UI	https://CommandCenterHost:5443 The Command Center login page is launched in your browser. The default
See Launching PCC for more details.	username/password is gpadmin/Gpadmin1 (case sensitive).
Browser:	After logging into PCC, click Add Cluster from the Cluster Status page.
Configure and deploy a cluster	Navigate through the Add Cluster Wizard to configure and deploy a cluster.
See Configuring and Deploying a Cluster for more details.	
Browser:	Return to the PCC UI and start the cluster from the Cluster Status page.
Start the cluster	
See Starting the Cluster for more details.	
Initialize and Configure HAWQ	(as gpadmin) ssh to the HAWQ master, then run the following:
See Initializing and Configuring HAWQ for	<pre># source /usr/local/hawq/greenplum_path.sh</pre>
more details.	# /etc/init.d/hawq init
	Add the IP address of your instance of Command Center to HAWQ's pg_hba.conf file

2.2 Supported Platforms and Browsers

2.2.1 Platforms

- RHEL 6.2 64-bit, 64 64-bit
- CentOS 6.2 64-bit, 64 64-bit

2.2.2 Browsers

(minimum screen resolution: 1280 x 800)

• Firefox 23



- IE 9, 10 (Chrome Tab not supported)
- Chrome 33.0.1750.146

2.3 Requirements

2.3.1 Prerequisites

- Installation of Pivotal Command Center assumes the user has a working knowledge of the following:
 - Yum: Enables you to install or update software from the command line. See http://yum.baseurl.org/.
 - RPM (Redhat Package Manager).
 - NTP. See http://www.ntp.org
 - SSH (Secure Shell Protocol). See http://www.linuxproblem.org/art_9.html
- DNS lookup. Verify that the admin host is be able to reach every cluster node using its hostname and IP address. Verify that every cluster node is able to reach every other cluster node using its hostname and IP address:

```
ping -c 2 myhost.mycompany.com
The return code should be 0

ping -c 2 192.168.1.2
The return code should be 0
```

• JAVA JDK. Ensure that you are running Oracle JAVA JDK version 1.7.0_xx (minimum 1.7.0.15)

As root:

```
# java -version
```

If you are not running the correct JDK, you can download a supported version from the Oracle site here:

http://www.oracle.com/technetwork/java/javase/downloads/index.html

Install the JDK on the admin node and add it to alternatives as follows:

```
# sudo /usr/sbin/alternatives --install "/usr/bin/java" "java"
"/usr/java/jdk1.7.0_xx/bin/java" 3
# sudo /usr/sbin/alternatives --install "/usr/bin/javac" "javac"
"/usr/java/jdk1.7.0_xx/bin/javac" 3
# sudo /usr/sbin/alternatives --config java
```

Once you have installed Pivotal Command Center, you can use the following command to import a downloaded JDK:



As gpadmin, run:

```
# icm client import -r <PATH TO JDK>
```

Note that the -r option refers to rpm files

- YUM. Verify that all hosts have yum access to an EPEL yum repository. See Package Accessibility for more details.
- Internet Connectivity. You will need an active internet connection to install PCC. This is required to pull the software dependencies for the UI.

2.3.2 Product Downloads

The following packages are required:

```
PCC-2.2.x.*.version_build_OS.x86_64.tar.gz
```

2.3.3 Package Accessibility

Pivotal Command Center and Pivotal HD Enterprise expect some prerequisite packages to be pre-installed on each host, depending on the software that gets deployed on a particular host. In order to have a smoother installation it is recommended that each host would have yum access to an EPEL yum repository. If you have access to the Internet, then you can configure your hosts to have access to the external EPEL repositories. However, if your hosts do not have Internet access (or you are deploying onto a large cluster), then having a local yum EPEL repo is highly recommended. This will also give you some control on the package versions you want deployed on your cluster. See for instructions on how to setup a local yum repository or point your hosts to an EPEL repository.

For Pivotal Command Center 2.2.x, here is a list of prerequisites that need to either already be installed on the Command Center admin host or on an accessible yum repository:

- httpd
- \bullet mod_ssl
- postgresql
- postgresql-devel
- postgresql-server
- compat-readline5
- createrepo
- sigar
- sudo

Run the following command on the admin node to make sure that you are able to install the prerequisite packages during installation.

```
$ yum list <LIST OF PACKAGES>
```



For example:

```
$ yum list httpd mod_ssl postgresql postgresql-devel postgresql-server
compat-readline5 createrepo sigar sudo
```

If any of them are not available or not already installed, then you may have not added the repository correctly to your admin host.

For the cluster hosts (where you plan to install the cluster), the prerequisite packages depend on the software you will eventually install there, but you may want to verify that the following two packages are installed or accessible by yum on all hosts:

- nc
- postgresql-devel

For example:

```
$ yum list nc postgresql-devel
```

2.4 Installation Instructions

Once you have met the prerequisites, you are ready to begin the installation.

Perform the following installation steps as a root user.



Avoid using hostnames that contain capital letters because Puppet has an issue generating certificates for domains with capital letters.

Also avoid using underscores as they are invalid characters in hostnames.

If you are upgrading from an earlier version of Pivotal Command Center, see Upgrading PCC.

2.4.1 Installing Pivotal Command Center

1. Create a directory (phd) for your PCC installation:

```
# mkdir phd
```

2. Copy the Command Center tar file to your host. For example:

```
# scp ./PCC-2.2.x.version.build.os.x86_64.tar.gz host:/root/phd/
```

3. Log into the Command Center admin host as root user. cd to the directory where the Command Center tar files are located and untar them. For example:

```
# cd /root/phd
# tar --no-same-owner -zxvf PCC-2.2.x.version.build.os.x86_64.tar.gz
```

4. Still as root user, run the installation script. This installs the required packages, configures Pivotal Command Center, and starts services.



Important

You must run the installation script from the directory where it was extracted; for example: For example: PCC-2.2.x.version

For example:

```
# 1s
PCC-2.2.x.version
PCC-2.2.x.version.build.os.x86_64.tar.gz
# cd PCC-2.2.x.version
# ./install
```

You will see installation progress information on the screen. Once the installation successfully completes, you will receive an installation success message on your screen.

Once you have configured and deployed a cluster, you can view your cluster status here:

https://<CommandCenterHost>:5443/status

5. Enable Secure Connections:

Pivotal Command Center uses HTTPS to secure data transmission between the client browser and the server. By default, the installation script generates a self-signed certificate.

Alternatively you can provide your own Certificate and Key by following these steps:

a. Edit /etc/httpd/conf.d/pcc- vhost.conf file and change following two directives to point to location of ssl certificate and key:

```
SSLCertificateFile: /usr/local/greenplum-cc/ssl/sfo-wl.ic.cert
SSLCertificateKeyFile: /usr/local/greenplum-cc/ssl/sfo-wl.ic.key
```

- b. Set the ownership of the certificate file and key file to gpadmin.
- c. Change the permission to owner read-only (mode 400).
- d. Restart PCC with the following command:

```
service commander restart
```

6. Verify that your PCC instance is running by executing the following command:

```
$ service commander status
```

7. You can switch to the gpadmin user for the rest of the installation process.

```
su - gpadmin
```

Next Steps: Installing PHD Services, then Enabling PHD Services.

Starting, Stopping, and Restarting Command Center Services

To stop or restart Command Center services, as root run the following commands on the Pivotal Command Center admin host:



```
$ service commander stop
$ service commander start
$ service commander restart
```



These commands need to be run as root.

2.4.2 Copying PHD Service Files

- Copy the Pivotal HD, ADS tarballs from the initial download location to the gpadmin home directory (home/user/gpadmin).
- 2. Change the owner of the packages to gpadmin then untar the tarballs. For example:

```
For PHD, if the file is a tar.gz or tgz, use:
```

```
tar -zxf PHD-2.0.x-x.tgz

If the file is a tar, use:
tar -xf PHD-2.0.x-x.tar

For Pivotal ADS, if the file is a tar.gz or tgz, use
tar -zxf PADS-1.2.x-x.tgz

If the file is a tar, use:
tar -xf PADS-1.2.x-x.tar
```

2.4.3 Importing PHD Services

```
1. As gpadmin, extract the following tarball for Pivotal HD:
```

```
For example:

# icm_client import -s PHD-2.0.x-x/

2. Optional for HAWQ/PXF: As gpadmin, extract the following tar ball for HAWQ and PXF: #
icm_client import -s <PATH TO EXTRACTED ADS TAR BALL>
For example:

# icm_client import -s PADS-1.2.x-x/
For more information, see the log file located at:
```

/var/log/gphd/gphdmgr/gphdmgr-import.log

You are now ready to configure and deploy a cluster from the Pivotal Command Center UI.

icm_client import -s <PATH TO EXTRACTED PHD TAR BALL>



2.4.4 Launching Pivotal Command Center

Launch a browser and navigate to the host on which you installed Command Center. For example:

https://CommandCenterHost:5443

The Command Center login page is launched in your browser. The default username/password is <code>gpadmin/gpadmin1</code> (case sensitive).

Next Steps

See Using PCC for details about using the application, including how to change the default password and how to deploy and configure a HD cluster via the Command Center UI.

See the *Pivotal HD Enterprise Installation and Administrator Guide* for instructions for using the command-line interface (CLI) of Pivotal Command Center to deploy and configure a HD cluster.

2.5 Uninstalling Pivotal Command Center

Follow the steps below to uninstall Pivotal Command Center and the Pivotal HD cluster:

- 1. Stop services on all your clusters (See the *Pivotal HD Enterprise Installation and Administrator Guide* for detailed steps).
- 2. Uninstall all your clusters (See the *Pivotal HD Enterprise Installation and Administrator Guide* for detailed steps).
- 3. From the directory where you untarred the Pivotal Command Center, run the uninstall script as root:
 - # cd /root/phd/PCC-2.2.x.version/
 - # ./uninstall

2.6 Upgrading Pivotal Command Center

The following instructions are for upgrading Pivotal Command Center from version 2.1.x to 2.2.



Upgrade Notes

- If you are upgrading to a new version of Pivotal Command Center, make sure you are also upgrading to compatible versions of Pivotal HD and Pivotal ADS (optional).
- See the latest version of the Pivotal Command Center Release notes for Pivotal Interoperability Matrix.
- We recommend that you always back up any critical data before performing any upgrades.
- As of PCC 2.1.1, there is a new default password, Gpadmin1, and new password rules are being enforced (minimum 8 characters, case-sensitive, 1 uppercase letter, one number). However if you are upgrading to this version of PCC, your original password (by default gpadmin) is maintained.
- JDK 1.7 is a requirement for PCC version 2.2. If you were using JDK 1.6 with an earlier version of PCC, you must install JDK 1.7 before upgrading to PCC 2.2.

Follow the steps below to upgrade your Pivotal CC to a newer version:

1. [Optional] As gpadmin, stop the cluster:

```
# icm_client stop -l CLUSTERNAME
```

2. As root, stop the Command Center services:

```
# service commander stop
```

- 3. Download the new PCC tarball and untar.
- 4. Run the installer from the new PCC location:
 - # ./installNote that your installation will fail if you are do not have the required JDK 1.7 installed.
- 5. Enable Secure Connections.
 - Pivotal Command Center uses HTTPS to secure data transmission between the client browser and the server. By default, the installation script generates a self-signed certificate. Alternatively you can provide your own Certificate and Key; to do this, follow the instructions provided in the installation instructions, here: Enable Secure Connections.
- 6. Check nmon status:

```
# service nmon status
```

If the status anything other than running, stop and restart nmon as follows:

```
# service nmon stop
```

service nmon start



3 Using PCC

This section provides an overview of the Pivotal Command Center 2.2 user interface and details about using the application to configure and deploy a Pivotal HD Cluster:

- Overview
- User Management
- Cluster Status Page
- Configuring and Deploying a Cluster
- Dashboard
- Cluster Analysis
- MapReduce Job Monitor
- YARN App Monitor
- HAWQ Queries
- Topology
- Logs

3.1 Overview

Pivotal Command Center UI is a browser-based application for configuring, deploying, administering, and monitoring Pivotal HD clusters. At a high level, the screens consist of:

- Cluster Status Page—Provides status information about any clusters you have configured and deployed. Also provides access to the Add Cluster Wizard that allows you to configure and deploy clusters from the UI. See Configuring and Deploying a Cluster for more details.
- Dashboard—Provides an overview of your Pivotal HD cluster. This screen shows at one glance the
 most important states and metrics that an administrator needs to know about the Pivotal HD cluster.
- Cluster Analysis—Provides detailed information about various metrics of your Pivotal HD cluster. This provides cluster-wide metrics all the way down to host-level metrics.
- MapReduce Job Monitor—Provides details about all, or a filtered set of MapReduce jobs.
- YARN App Monitor—Provides details about all, or a filtered set of YARN applications.
- HAWQ Queries—When HAWQ (a revolutionary MPP database on Hadoop solution) is deployed on the cluster, Command Center can show the progress of all actively running queries on HAWQ.
- Topology—This screen shows you what roles have been installed on each host. You can also add and remove slaves to/from the cluster from this screen.
- Logs—This screen displays system logs based on filter criteria you select such as log levels, host name, time period, and so on.

3.1.1 Status indicators

Throughout the user interface the following indicators are used to indicate the status of nodes:

· Green: Succeeded



• Blue: Running

Grey: Stopped/Pending

• Red: Killed/Failed

3.2 Logging In

Launch a browser and navigate to the host on which you installed Command Center. For example:

https://CommandCenterHost:5443

The Command Center login page is launched in your browser. The default username/password is <code>gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/Gpadmin/G</code>

To change the default port (5443), update the port settings in the following file:

/usr/local/greenplum-cc/config/app.yml

3.2.1 Login Screen

The first time you launch the Command Center UI, a login screen appears showing the hostname of this instance of Pivotal Command Center.



The default admin username/password is <code>gpadmin/Gpadmin1</code> (case-sensitive). You can change this password via the Settings menu.

Passwords are case-sensitive and must be at least 8 letters long and contain 1 upper-case letter and 1 number.

Once you have entered a valid username/password, click the **Login** button to launch the Command Center UI.



3.3 Settings

Once you have logged in, you can click the gear icon in the upper right corner of the screen from any PCC page to display the Settings menu.



From the settings menu you can select one of:

- About. Select this to display version information about this instance of PCC
- Cluster Status. Select this option to go back to the Cluster Status page to view the list of available
- Users. Select this option to add/edit user information. There are two types of users, administrative and non-administrative (read-only); you will only see this **Users** option if you are an administrative user. See Users below.
- Change Password. Click this to change your password. See Passwords below.
- Edit Profile. Click this to change your profile. See Users below
- Logout. Select this option to logout from this instance of PCC.

3.4 User Management

There are two types of users, administrative (super user) and non-administrative (read-only).

If you are an administrative user, you can add users, edit your own or other users' profiles, change your own or other users' passwords, and delete users.



Non-administrative users are read-only users; their actions are limited to read-only actions.



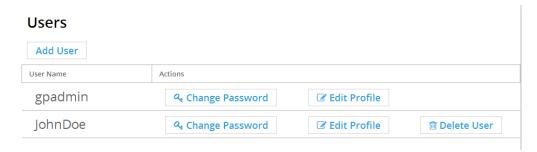
For many user management actions, you will be asked to Re-authenticate to Continue. In all cases, this means you need to enter the password of the current user performing the action to continue.



3.4.1 Users

You will only see the **Users** option in the **Settings** menu if you are an administrative user.

Select **Users** from the **Settings** menu to see a screen listing all current users:



Adding a User

1. Click Add User to create a new user.

A New User form appears. Enter the following information about the new user:

- First Name.
- Last Name.
- User Name. This must be a unique name.
- Email. Must be a valid email address.
- Administrator. Check this box if the new user is to have administrative privileges.
- Password. Enter a password for the new user. This must meet the minimum password requirements. Passwords are case-sensitive and must be at least 8 letters long and contain 1 upper-case letter and 1 number.
- Confirm Password. Confirm the new user's password.
- 2. Click OK.

A dialog box appears asking you to **Re-authenticate to Continue**.

- 3. Enter the password of the current user (the administrative user who is creating the new user).
- 4. Click OK.

You are returned to the Users screen where you should now see your new user(s) in the list.

Deleting a User

Only administrative users can delete users and the default gpadmin user can never be deleted. To delete a user:

- 1. Select **Users** from the **Settings** menu
- 2. Click the **Delete User** link adjacent to the user you want to delete. A dialog box appears asking you to Re-authenticate to Continue.
- 3. Enter the password of the current user (the administrative user who is deleting the user).
- 4. Click OK.

You are returned to the Users screen where the user you just deleted should no longer appear in the list.



3.4.2 Profiles

Your profile includes: First Name, Last Name, User Name, Email Address, and User Type (administrative or not); it does not include your password.

All users can edit their own profiles by:

• Selecting Edit Profile from the Settings menu.

Administrative users can edit their own passwords and also the profiles of other users by:

• Selecting Users from the Settings, then clicking the Edit Profile link adjacent to the user whose profile they wish to edit.

After making edits to any profile, a dialog box appears asking you to Re-authenticate to Continue. Enter the password of the current user making the edits to continue, then click **OK**.

3.4.3 Passwords

Passwords are case-sensitive and must be at least 8 letters long and contain 1 upper-case letter and 1 number.

All users can change their own passwords by:

Selecting Change Password from the Settings menu.

Administrative users can change their own passwords and also the passwords of other users by:

 Selecting Users from the Settings menu, then clicking the Change Password link adjacent to the user whose password you wish to change.

After changing any password, a dialog box appears asking you to Re-authenticate to Continue. Enter the password of the current user making the edits to continue, then click OK.

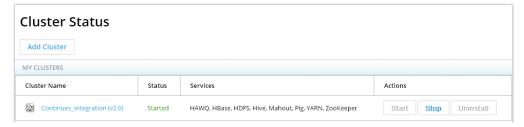
3.5 Cluster Status Page

Once you have launched Command Center, the initial screen you see is the Cluster Status screen. This displays a list of available clusters to monitor, the status of each cluster (started, stopped), and a list of services running on that cluster (Hive, Mahout, and so on).



The Add Cluster and Actions buttons are only visible to administrative users.





From this page you can:

- Administrative users only: Click Add Cluster to launch the Add Cluster Wizard.
- Click the cluster name in the table to view the Dashboard for that cluster.
- From any point within Command Center UI, you can always select a different cluster by using the **Select Cluster** drop-down menu in the upper right corner of the screen.
- Administrative users only: You can either Start, Stop, or Uninstall a cluster. Depending on the state of the cluster, some of these buttons will be enabled while others are disabled.

3.6 Configuring and Deploying a Cluster



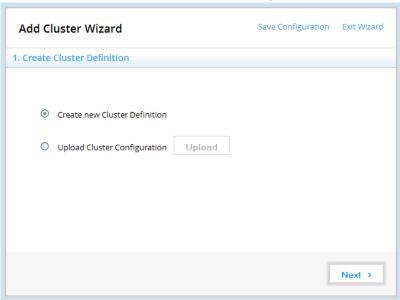
Before you can configure and deploy a cluster, make sure you have already installed and enabled the PHD Services (see Install PHD Services and Enable the PHD Services).

After you have logged in to Pivotal Command Center, the Cluster Status page appears. From here, if you are an administrative user, you are able to launch the Add Cluster Wizard that enables you to configure and deploy a Pivotal HD Cluster.

As you move through the wizard, the right hand pane displays where you are in the deployment process:

3.6.1 Adding a Cluster

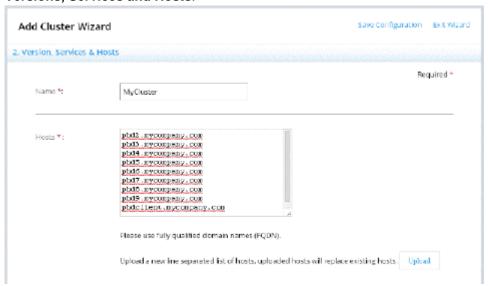
1. Click Add Cluster. The Add Cluster Wizard opens:



The Wizard allows you to create a new configuration from scratch or upload and edit any existing configuration. The Summary panel along the right shows you the progress of your configuration and deployment.

2. Create Cluster Definition. Either:

- a. If you are configuring a new cluster, select Create a new Cluster Definition then click Next.
- b. If you want to edit an existing cluster; select **Upload Cluster Configuration**, click **Upload**, then navigate to the clusterConfig.xml file that you wish to edit; then click **Next**. In this case, the following fields in the Wizard will be populated with the cluster definition properties of that clusterConfig.xml file you just uploaded. Follow the instructions below to edit those values.
- 3. Versions, Services and Hosts:



A Notes

- Hosts can be entered individually, newline-separated; or can be expressed in a range, for example host[1-5].yourdomain.com. They can also be expressed in multiple ranges, for example host [1-3].subdomain[1-2].yourdomain.com. Any hosts expressed in ranges are expanded during host verification. Hosts that do not exist within a specified range will be ignored, so you can specify a wide range and only those hosts that are available within that range will be added.
- If you are editing an existing configuration, some if not all of these fields will be pre-populated. Edit where appropriate.
- You need to scroll down to view all the fields on this screen. The Next button will not be active until you have entered all the required fields.

Enter the following information:

- Name: Required. Enter a name for this cluster. Special characters are not supported.
- Hosts: Required. Enter a new line-separated list of FQDN host names. You can also click **Upload** to use a text file containing a new line-separated list of host names.
- Root Password: Required. Enter the root password.
- GP Admin Password: Required. Enter the gpadmin user password. PCC creates this user on all nodes.
- **JDK Path**: Enter the JDK filename (not the absolute path). For example: jdk-7u51-linux-x64.rpm.

Note: JDK 1.7.0 15 (min) is a prerequisite.

If not already installed, you should download it then install using icm_client import -r (see Installing PCC for more details)

- Setup NTP: Check this box if you want to set up NTP (Network Time Protocol).
- Disable SELinux: Check this box if you want to disable SELinux. Recommended.
- Disable IPTables: Check this box if you want to disable IPTables. Recommended.
- Run ScanHosts: Leave this box checked if you want to run scanhosts. The scanhosts command verifies the prerequisites for the cluster node have been met and provides a detailed report of any missing prerequisites. Running this command ensures that clusters are deployed smoothly and is strongly recommended.

Click Next.

4. Host Verification:

The Host Verification page opens. This step may take a few minutes, it verifies connections to the hosts you just set up. Once the Eligibilty field changes from Pending, to Eligible for all hosts, you can click **Next**. You will see any error and informational messages displayed in the comments fields.



If you specified hosts using ranges, they will be expanded at this point.

5. Topology:

This is the section where you specify the roles to be installed on the hosts. For example, you can specify where your hadoop namenode, data node and so on, should be installed. Note that all mandatory roles should have at least one host allocated.

Each service has its own section on this page; you can use the top menu options as shortcuts to those sections on the page, or simply scroll down to each section.



Notes

You need to click Enter or Tab before each field is accepted. Once you enter the text
and click Enter or Tab, the text will change appearance and appear enclosed in a
box, as shown in the figure below. The entry on the left has been accepted, the entry
on the right has not.

centos62-2 * invalidhost.com

- Hosts can be specified in ranges, see the notes for Versions, Services and Hosts, for more information.
- At any point during this stage you can click **Save Configuration** at the top right of the page. This saves the configuration file and downloads it. Once saved, a link to the configuration file appears at the bottom of the page. Click that link to open and view the clusterConfig.xml file. You cannot edit this xml file directly.

These are the roles that need to have installation nodes defined:

- CLIENT: ICM installs Pig, Hive, HBase, and Mahout libraries on this host.
- HDFS: Name Node, Secondary Name Node, Data Nodes
- YARN: Resource Manager, History Server, Node Managers
- Zookeeper: Zookeeper Server
- HBase: Hbase Master, HBase Region Servers.
- Hive: Hive Master, Hive Metastore
- HAWQ: Primary Node, Secondary Node, HAWQ Segment Nodes
- USS: Name Node and Catalog
- PXF: No hosts to configure. Installed on the client host.
- Mahout: No hosts to configure. Installed on the client host.
- Pig: No hosts to configure. Installed on the client host.

Click **Next** once you have finished role-mapping.

6. Cluster Configuration:

This page displays a list of all configuration files that define this cluster; the clusterConfig.xml (to edit service configuration global values) as well as the service specific configuration files.

All these configuration files are already populated with the values you have already entered; or with default values.

Click any file name to open that configuration file in an editor and enter/edit values.

If you make any changes, click **Save** to return to the Cluster Configuration page.

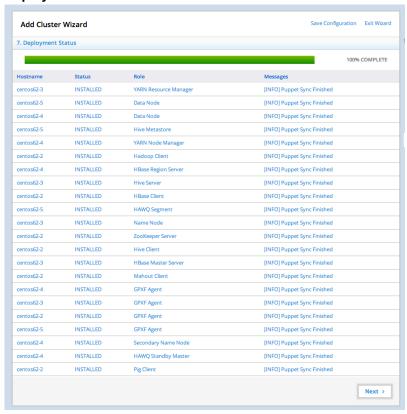
Once you have completed all your edits, click Next.

7. Validation:

If the configuration has errors they will be displayed here; otherwise you will see post-deployment instructions.

Click Deploy

8. Deployment Status:



This screen shows the progression of the deployment. Information displayed includes:

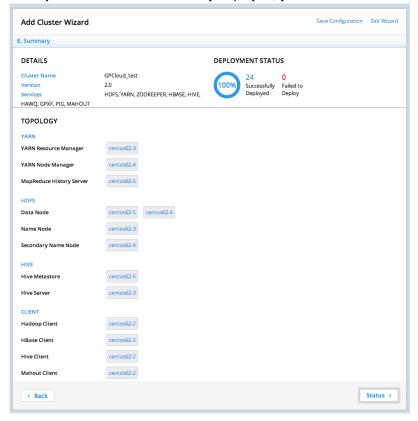
- Hostname
- Status
- Role
- Messsages

Once the deployment is complete, click Next.



9. Summary

Once your cluster has successfully deployed, you can view a summary of the cluster, as shown here:



10. Return to the Cluster Status page:

Once you have reviewed this summary, click **Status**, to return to the Cluster Status page. Your new cluster will be listed on this page, with the status of **installed**.

If you are deploying HAWQ you need to initialize and configure HAWQ after you start the cluster. See Initializing and Configuring HAWQ.

Next steps: Starting the Cluster then Initializing and Configuring HAWQ.

3.6.2 Starting the Cluster

To start your cluster; click **Actions**: **Start** on the Cluster Status page.

3.6.3 Initializing and Configuring HAWQ

For HAWQ users:

ssh to the HAWQ master, then as gpadmin, run the following:

```
# su - gpadmin
```

source /usr/local/hawq/greenplum_path.sh



/etc/init.d/hawq init

Add the IP address of your instance of Command Center to HAWQ's $pg_hba.conf$ file, for example:

vi /data1/master/gpseg-1/pg_hba.conf

host all gpadmin <Command_center host IP>/24 trust

Then restart HAWQ:

/etc/init.d/hawq restart

You have now completed your cluster configuration and deployment.

See the following sections of this document for details about using the PCC UI to administer and monitor your cluster.

3.7 Starting Stopping and Uninstalling the Cluster

These functions are only available to administrative users, also, depending on the state of the cluster, some of these buttons will be enabled while others are disabled.

3.7.1 Starting

To start a cluster:

From the Cluster Status page locate the cluster you want to start, then click the Actions: Start button.

3.7.2 Stopping

To stop a running cluster:

From the Cluster Status page locate the cluster you want to stop, then click the **Actions**: **Stop** button.

3.7.3 Uninstalling

Only stopped clusters can be uninstalled.

To uninstall a cluster:

- 1. From the Cluster Status page locate the cluster you want to uninstall, then click the **Actions**: **Uninstall** button.
- 2. An Uninstall confirmation dialog appears. Click **Yes** to confirm you want to uninstall this cluster; **No** to cancel the operation and return to the Cluster Status page.



3. Once you confirm you want to continue with the uninstall operation, a Preserve Cluster State dialog appears. You need to specify whether you want to preserve the data for this cluster. Your options are: **Cancel**. This cancels the operation and returns you to the Cluster Status page.

Yes. This preserves the data for this cluster before uninstalling the cluster and returning you to the Cluster Status page.

No. This uninstalls the cluster without preserving the data; then returns you to the Cluster Status page.

3.8 Dashboard

The dashboard gives you a high level view of a cluster at a glance. You are able to view the status of the most important cluster services, such as HDFS and YARN, and allows you to start and stop each service individually. It also shows you how the most important cluster metrics are trending in a visual way.

The graphs provide a unified view of the state of your system. They are also useful in detecting outliers and pinpointing specific problems that may be present in your system.



The right side of the Dashboard displays the state of the following services, provided they have been deployed for this cluster:

HDFS

For HDFS, the dashboard provides the following information/functionality:

- The status of HDFS. You can use the **Actions** dropdown menu to **Start/Stop HDFS** depending on its status.
 - When the last NameNode checkpoint occurred.
 - The percentage of cluster storage being used by HDFS and how much is free.
 - The number of DataNodes that are up and whether they are running normally or with problems.
 - The Actions dropdown menu allows you to Rebalance (redistribute your data across the cluster) your cluster, and to View Rebalancer Log.



The only Action option available to non-administrative users is to View Rebalancer Log.



If High Availability (HA) is enabled for your cluster, you will see the status of two NameNodes here.

YARN

For YARN, the dashboard provides the following information:

- The status of YARN. You can use the **Actions** dropdown menu to **Start/Stop YARN** depending on its status (not available for non-administrative users).
 - Whether or not the YARN History Server is running.

Note: The History Server stores a history of the mapreduce jobs run on the cluster.

The number of NodeManagers that are running.

HBase

For HBase, the dashboard provides the following information:

 The status of the HBase master. You can use the Actions dropdown menu to Start/Stop HBase depending on its status (not available for non-administrative users).

Zookeeper

For Zookeeper, the dashboard provides the following information:

- The status of Zookeeper. You can use the Actions dropdown menu to Start/Stop Zookeeper depending on its status (not available for non-administrative users).
 - The Status of the Zookeeper server host.

Hive

For Hive, the dashboard provides the following information:

 The status of Hive. You can use the **Actions** dropdown menu to **Start/Stop Hive** depending on its status (not available for non-administrative users).

The Dashboard also provides metrics about:



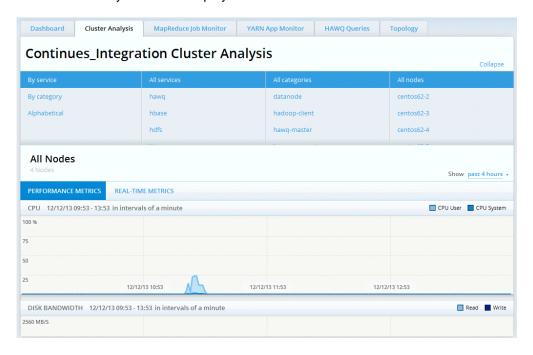
- Mapreduce Slot Utilization
- Namenode RPC Times
- Hadoop Datanodes Average CPU
- Hadoop Datanodes Average Bandwidth
- Namenode Operations Per Second
- Hadoop Datanodes Average Disk Bandwidth
- Hadoop Datanodes Average Memory
- Mapreduce Jobs By Status

3.9 Cluster Analysis

The Cluster Analysis screen provides detailed metrics on your Pivotal HD cluster.

It provides cluster-wide metrics all the way down to host-level metrics. It provides Hadoop-specific metrics, as well as system metrics that you can drill down to if needed.

The Cluster Analysis screen displays the same data that is shown in the dashboard but in greater detail.



By default the Cluster Analysis screen displays the metrics for all services, all categories, and all nodes. You can filter the information displayed by combinations of the following filters:

By Service. Metrics can be filtered by services such as HDFS, YARN, or HAWQ.

By Category. Metrics can be filtered by categories such as:



- namenode
 - secondarynamenode
 - datanode
 - yarn-resourcemanager
 - yarn-nodemanager
 - mapreduce-historyserver
 - hawq-master
 - hawq-segment

Alphabetically. Metrics can be filtered alphabetically.

Based on the filters you select, the lower part of the Cluster Analysis screen provides detailed graphs that display data related to:

- CPU
- Disk Bandwidth
- Network Bandwidth
- Memory
- Load
- Swap Usage
- Swap I/O
- Network Operations
- Disk Operations

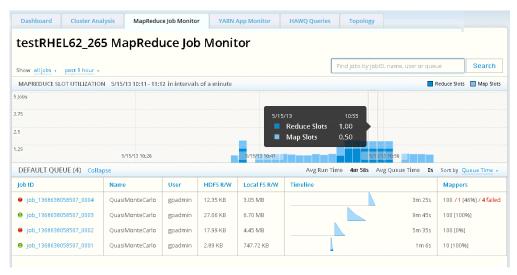
You can view either the **Performance Metrics**, which show the cluster/node utilization over-time, the **Real-time Metrics** which show the current metrics in real-time, or **Storage Metrics**, which show metrics about cluster storage.

If you select Cluster Analysis for **All Nodes** (the default), the Trending Metrics graph for the cluster is displayed.

3.10 MapReduce Job Monitor

The Job Monitor screen tracks the MapReduce jobs that are executed in the Pivotal HD cluster when the YARN MapReduce service is running. It provides details about all, or a filtered set of MapReduce jobs.





The MapReduce jobs displayed can be filtered by state and/or time range.

By state:

- all jobs (set by default)
 - · currently pending jobs
 - · currently running jobs
 - succeeded jobs
 - failed jobs
 - killed jobs
 - · error state jobs

By time range: By selecting a preset time range in hours, weeks, months, year, or by specifying a custom time range.

The MapReduce jobs can also be filtered by searching for values for the following:

- jobID
- name
- user
- queue

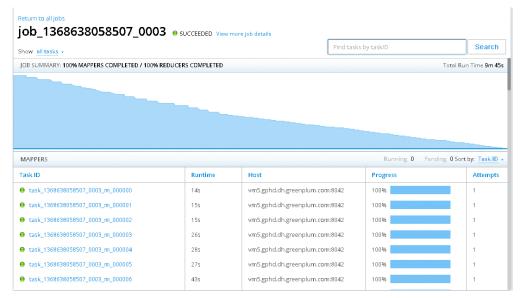
Enter your search value in the search bar in the following format: searchKey=searchValue, where searchKey is one of **jobID**, **name**, **user**, or **queue**.

These are substring searches. For example: **jobID=1363920466130** will locate a job with **jobID=job_1363920466130_0002**

3.10.1 Job Details

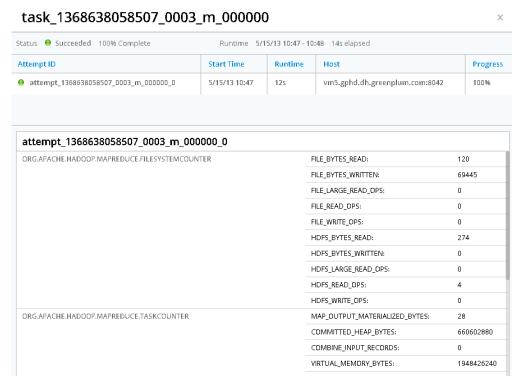
When you click on any of the jobs in the Job Monitor more details of the job are shown.





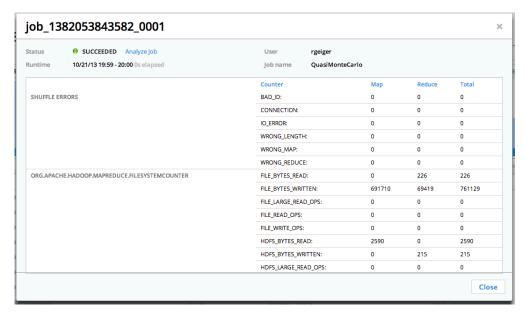
This screen displays all the tasks that are have been allocated for the selected job and their progress. You can see the mapper and the reducer tasks separately. In the above screen capture, the bars in the JOB SUMMARY section represent the two Mapper tasks that have run, one took 19 seconds, the other, 20 seconds.

Clicking on each task ID will show even more details about that particular task. You can also filter on a particular task ID in the search bar.



To see job related counters click on View more job details next to the job ID:

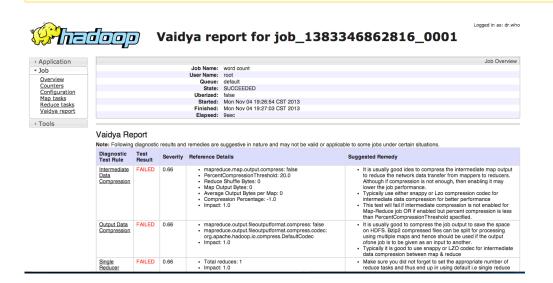




Click the **Analyze Job** link adjacent to the Status field to open a Vaidya report about the selected job, as shown below:



This capability is beta and will be improved in coming releases.



About Vaidya

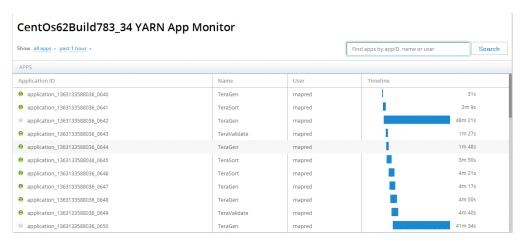
Vaidya is a diagnostic tool installed with PHD for Map/Reduce jobs. After a job is executed successfully, it uses a job history log and job configuration information to identify any performance or scalability problems with the job. Upon execution, it provides a job analysis report indicating specific problems with the job along with the remedy to correct them.

For more information about Vaidya, see the PHD Enterprise Stack and Tool Reference Guide.



3.11 YARN App Monitor

The YARN App Monitor screen tracks YARN applications that are executed in the Pivotal HD Cluster.



The YARN applications displayed can be filtered by category and/or time range:

- By Category:
 - · all apps (set by default)
 - · currently pending apps
 - · currently running apps
 - succeeded apps
 - · failed apps
 - killed apps
 - error state jobs
- By Time Range: By selecting a preset time range in hours, weeks, months, year, or by specifying a custom time range.

The YARN applications can also be filtered by the following fields by entering it in the search bar in the following format: searchKey=searchValue:

- appID
- name
- user

These are substring searches. For example: appID=1363920466130 will locate the application with appID=application_1363920466130_0002

3.12 HAWQ Query Monitor

The HAWQ Query monitor is only displayed when HAWQ is installed on the cluster.

This screen displays all active queries running on the HAWQ cluster:



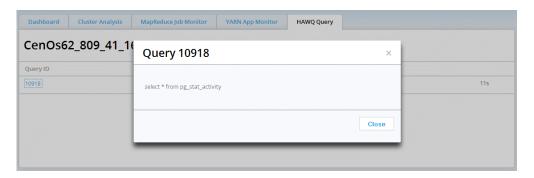


In this release, this screen only displays active queries as can be seen when you run:

SELECT * FROM pg_stat_activity;

on the HAWQ cluster.

Click on a Query ID to get the syntax of that query:



3.13 Topology

This screen shows you what roles have been installed on each host. You can also add and remove slaves to/from the cluster via this screen.



A node is considered a slave if it contains the following services: datanode, hbase-regionserver, yarn-nodemanager.



3.13.1 Topology Actions



This functionality is restricted to administrative users.



Adding Slaves to the Cluster

Click on the **Add Slaves to the Cluster** option from the **Topology Actions** menu. An **Add Slaves** dialog appears .

This dialog lists all current nodes. Enter the slave nodes you want to add either individually, or as ranges, for example Node [1-9], provide the Root Password, then click **Add Slaves to Cluster**.

You are warned that new slaves added to the cluster are not started automatically. To start a slave node, return to the Dashboard and navigate to the role defined for that node, then select **Action > Start** <service name>, for example **Action > Start Hive**.



- Newly added slaves are not started automatically.
- If you type a node name twice, that name will flash yellow until one is deleted.
- If you type the name of a node that already exists, that name becomes highlighted and you are shown an error message warning you that:
 - Highlighted hosts are duplicates and will not be added.
- If you close the dialog before the operation is finished, it continues in the background. Refresh the Topology page to see if it has successfully completed.

Removing Slaves from a Cluster

Click on the **Remove Slaves from the Cluster** option from the **Topology Actions** menu. A **Remove Slaves** dialog appears; this dialog lists all current slave nodes

You are warned that prior to removal, services on slave nodes must be stopped. To stop services on a slave node, return to the Dashboard and navigate to the role defined for that node, then select **Action > Stop** <service name>, for example **Action > Stop Hive**.

Enter the slave nodes you want to remove either individually, or as ranges, for example Node[1-9], then click Remove Slaves from Cluster.



- Prior to removal, services on slave nodes must be stopped.
- You do not need to provide a password to remove slave nodes.
- This text field auto-completes with nodes from the cluster.
- If you try and enter a node that does not exist, the text field does not become active and you are not able to perform the remove operation.
- If you close the dialog before the operation is finished, it continues in the background. Refresh the Topology page to see if it has successfully completed.



3.14 Logs

This screen displays all system logs.

3.14.1 Filtering Logs

You can use the dropdown menus to filter what logs are displayed based on:

- Logs Levels: You can select to display all logs (the default), or filter by one of the following log levels: debug, errors, fatal, info, trace, warnings
- Hosts: You can select to display logs for all hosts (the default), or can filter by hostname
- Roles: You can select to display logs for all roles (the default), or can filter by role name
- Time: You can select the time period over which you want to view logs; either for the Past 1 hour, or for one of the other time periods available (up to Past 8 weeks)

Search by Keyword

You can also further filter the logs displayed by searching the logs by keyword.

Enter any single keyword (any single string, including numbers) in the search box, then click **Search**.

Note that your keyword search is only applied to the set of logs displayed based on any filters you chose from the dropdown menus, described above.

3.14.2 Viewing Logs

The results of your filters, if any, are displayed on the screen. Often there will be too many to display on one page, use the Prev/Next/Page numbers at the top of the screen to navigate through them.

Each entry displays the time stamp, a log message, a time stamp, the host and the role, for example:

2014-02-25 23:05:51,137 WARN org.apache.hadoop.yarn.server.nodemanager.containermanager.AuxServices: The Auxilurary Service named 'mapreduce_shuffle' in the configuration is for class class org.apache.hadoop.mapred.ShuffleHandler which has a name of 'httpshuffle'. Because these are not the same tool...

yarn-nodemanager centos62-4 Show Logs | Show Details

2014-02-25 23:05:47,585 - WARN [NIOServerCxn.Factory:0.0.0.0/0.0.0.0:2181:NIOServerCnxn@347] - caught end of stream exception EndOfStreamException: Unable to read additional data from client sessionid 0x1446b2ad66b0001, likely client has closed socket at org.apache.zookeeper.server.NIOServerCnxn.d...

zookeeper-server centos62-2

If the log message is too long to display on the screen it is truncated and a **Show Details** link appears below the message; click to display the entire message.

Click **Show Logs** to display all contents of that log file based on the time stamp. You can see more log messages from before and after that time stamp by clicking the **Fetch More Logs Above/Fetch More Logs Below** links at the top and bottom of the page, respectively.



4 Creating a YUM EPEL Repository

Pivotal Command Center and Pivotal HD Enterprise expect some prerequisite packages to be pre-installed on each host, depending on the software that gets deployed on a particular host. In order to have a smoother installation it is recommended that each host would have yum access to an EPEL yum repository. If you have access to the Internet, then you can configure your hosts to have access to the external EPEL repositories. However, if your hosts do not have Internet access (or you are deploying onto a large cluster), then having a local yum EPEL repo would be highly recommended. This will also give you some control on the package versions you want deployed on your cluster.

Following are the steps to create a local yum repo:

- 1. Mount the RHEL/CentOS DVD on a machine that will act as the local yum repo.
- 2. Install a webserver on that machine (for example. httpd), making sure that HTTP traffic can reach this machine.
- 3. Install the following packages on the machine:

```
yum-utilscreaterep
```

4. Go to the directory where the DVD is mounted and run the following command:

```
createrepo
```

5. Create a repo file on each host with a descriptive filename in the

```
/etc/yum.repos.d/ directory of each host (for example, CentOS-6.1.repo) with the following
contents:[CentOS-6.1]
name=CentOS 6.1 local repo for OS RPMS
baseurl=http://172.254.51.221/centos/$releasever/os/
$basearch/
enabled=1
gpgcheck=1
gpgkey=http://172.254.51.221/centos/$releasever/os/$basearch/RPM-GPG-KEY-Centos/$releasever/os/$pasearch/RPM-GPG-KEY-Centos/$releasever/os/$pasearch/RPM-GPG-KEY-Centos/$releasever/os/$pasearch/RPM-GPG-KEY-Centos/$releasever/os/$pasearch/RPM-GPG-KEY-Centos/$releasever/os/$pasearch/RPM-GPG-KEY-Centos/$releasever/os/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/RPM-GPG-KEY-Centos/$pasearch/
```

6. Validate that you can access the local yum repos by running the following command:

```
Yum list
```



5 Command Line Reference

This section provides descriptions and syntax for the command line operations you can perform.



You need to perform command line operations on the Admin node.

5.1 Backup and Restore

You can backup data on the admin node where PCC is installed. Having the backup allows you to restore the admin node and PCC to a given state in case of failures or data corruption.



Backup and restore operations should be performed as root.

5.1.1 Backup

Run the backup script to copy all configuration files and all data in the PCC database into a backup file on a local disk. Once the tar file is created you should copy it off the admin node to a different storage to prevent it from being lost if the admin node fails.



You should backup your data each time you make a configuration or topology change for your cluster or if you add or delete clusters. Trying to restore old backup on the admin node when the cluster topology or configuration has been changed since the backup will result in inconsistent configuration and potentially unusable cluster.

To perform the backup run the following script:

/usr/lib/gphd/gphdmgr/bin/gphdmgr_backup.sh

The script does not require any input and produces the backup file phdmgr_backup_[timestamp].tar.gz in the current working directory, where timestamp is the number corresponding to the current system time.

5.1.2 Restore

Run the restore script to replace all configuration files and all data in the PCC database from a backup file on a local disk. There are two restore scenarios:

- Restore data on a new node. If you had a backup from the old admin node and you need to provision a new admin node, install PCC on the new admin node, then restore data from the backup produced on the old admin node. Hostnames, IP addresses and system configuration of the old and new nodes have to be identical.
- Restore data on the same node. You can restore data from the backup file on the same admin node where the backup was produced. Make sure the PCC version has not been updated since the backup was made.



The restore procedure replaces all your configuration files and database data with the data from the backup file, so use with caution.

To perform the restore run the following script:

```
# /usr/lib/gphd/gphdmgr/bin/gphdmgr_restore.sh phdmgr_backup_[timestamp].tar.gz
Your data will be permanently removed. Do you want to continue? y
```

The script takes the name of the backup file as an argument and asks the user for the confirmation before proceeding with the restore. The script takes care of stopping the PCC services, performing the restore, and restarting the services.