

Pivotal HAWQ

Version 1.2

Release Notes

Rev: 01

Table of Contents

1	Welcome to HAWQ 1.2	4
1.1	About HAWQ Components	4
1.1.1	PXF	4
1.1.2	MADlib	5
2	New Features	6
3	Supported Platforms	8
4	Installation options	9
5	Resolved Issues in HAWQ 1.2	10
6	Known Issues in HAWQ 1.2	14
7	Known Issues in PXF 2.x.x	16
8	Pivotal and HAWQ Interoperability	17
9	HAWQ 1.2 and Pivotal HD Documentation	19
10	Upgrading HAWQ	20

Copyright © 2013 GoPivotal, Inc. All rights reserved.

GoPivotal, Inc. believes the information in this publication is accurate as of its publication date. The information is subject to change without notice. THE INFORMATION IN THIS PUBLICATION IS PROVIDED "AS IS." GOPIVOTAL, INC. ("Pivotal") MAKES NO REPRESENTATIONS OR WARRANTIES OF ANY KIND WITH RESPECT TO THE INFORMATION IN THIS PUBLICATION, AND SPECIFICALLY DISCLAIMS IMPLIED WARRANTIES OF MERCHANTABILITY OR FITNESS FOR A PARTICULAR PURPOSE.

Use, copying, and distribution of any Pivotal software described in this publication requires an applicable software license.

All trademarks used herein are the property of Pivotal or their respective owners.

Use of Open Source

This product may be distributed with open source code, licensed to you in accordance with the applicable open source license. If you would like a copy of any such source code, Pivotal will provide a copy of the source code that is required to be made available in accordance with the applicable open source license. Pivotal may charge reasonable shipping and handling charges for such distribution.

1 Welcome to HAWQ 1.2

HAWQ extends the functionality of Pivotal Hadoop (HD) Enterprise, adding rich, proven parallel SQL processing facilities. These SQL processing facilities enhance productivity, rendering Hadoop queries faster than any Hadoop-based query interface on the market. HAWQ enables data analysis for a variety of Hadoop-based data formats using the Pivotal Extension Framework (PXF), without duplicating or converting source files.

HAWQ is a parallel SQL query engine that combines the key technological advantages of the industry-leading Greenplum Database with the scalability and convenience of Hadoop. Using HAWQ functionality, you can interact with petabyte range data sets. HAWQ provides users with a complete, standards compliant SQL interface.

Leveraging Greenplum Database's parallel database technology, HAWQ consistently performs tens to hundreds of times faster than all Hadoop query engines in the market.

1.1 About HAWQ Components

HAWQ comprises the following components:

- PXF
- MADlib

1.1.1 PXF

PXF enables SQL querying on data in the Hadoop components such as HBase, Hive, and any other distributed data file types. These queries execute in a single, zero materialization and fully-parallel workflow. PXF also uses the HAWQ advanced query optimizer and executor to run analytics on these external data sources. PXF connects Hadoop-based components to facilitate data joins, such as between HAWQ tables and HBase table. Additionally, the framework is designed for extensibility, so that user-defined connectors can provide parallel access to other data storage mechanisms and file types.

PXF Interoperability

PXF operates as an integral part of HAWQ, and as a light add-on to Pivotal HD. On the database side, PXF leverages the external table custom protocol system. The PXF component physically lives on the Namenode and each or some Datanodes. It operates mostly as a separate service and does not interfere with Hadoop components internals.

1.1.2 MADlib

MADlib is an open-source library for scalable in-database analytics. It provides data-parallel implementations of mathematical, statistical and machine learning methods for structured and unstructured data. MADlib combines the efforts used in commercial practice, academic research, and open-source development. You can find more information at <http://madlib.net>.

2 New Features



Note

For specific information about a previous release, please refer to the associated release notes.

HAWQ 1.2 supports the following features:

- **PL/Java:** HAWQ 1.2 supports PL/Java
- **HAWQ expand:** You can add resources to an existing HAWQ system using gpexpand.
- **HDFS Namenode High Availability (HA):** You can initialize HAWQ on HA HDFS. Initializing on HD HDFS helps you avoid query failure that can occur if you have a single Namenode.
- **Error tables:** The most common use of readable external tables is selecting data from them to load into regular database tables. This is typically done by issuing a CREATE TABLE AS SELECT or INSERT INTO SELECT command, where the SELECT statement queries external table data. By default, if the external table data contains an error, the entire command fails and no data is loaded into the target database table. To isolate data errors in external table data while still loading correctly formatted rows, you can define a readable external table with a SEGMENT REJECT LIMIT clause in the CREATE EXTERNAL TABLE command. HAWQ 1.2. supports single row error isolation where any rows with formatting errors are logged into an error table.
- **User-Defined Functions:** HAWQ 1.1.3 provided support for User-Defined Functions (UDF). The UDF feature extends the functionality of the HAWQ database by providing functions that can be evaluated in SQL statements. With each release, Pivotal extends UDF feature support. This section lists the UDF functionality supported in each release, and explicitly calls out features that are not supported.

UDF features supported in HAWQ 1.2

- User defined composite types
- Set returning functions
- PL/Java
- Windows Derived Aggregate Functions

UDF features not supported in HAWQ 1.2

- SECURITY DEFINER when creating functions
- Nested functions error out during execution
- User-Defined base type
- SORTOP is not supported for UDAs
- ALTER set encoding, set schema, and rename
- Window functions
- Enhanced table functions
- PL/Java Type Maps

See the *Pivotal HAWQ Installation Guide* for information about requirements and installation.

See the *Pivotal HAWQ Administrator Guide* for detailed information about features and usage.

- **MADlib 1.5:** When upgrading to HAWQ 1.2, you must upgrade to MADlib version 1.5 to achieve the functionality of the previous version.
- **Automated NameNode Failover**
- **PXF:** Automated High Availability, Hive 0.12 connectivity.

3 Supported Platforms

HAWQ 1.2 supports the Pivotal HD 2.0.1 platform.

4 Installation options

There are two ways to install HAWQ.

- Stand alone install – You can install HAWQ without using the PivotalHD tools ICM and PCC. For more information, please see *HAWQ 1.2 Installation and Upgrade Guide*
- Pivotal Command Center Command Line Interface – Please see *Pivotal HD Enterprise 2.0 Installation and Administrator Guide*.

5 Resolved Issues in HAWQ 1.2

The table below lists issues that are now resolved in HAWQ 1.2.



Note

For issues resolved in prior releases, refer to the corresponding release notes available from Support Zone.

Table: Resolved Issues in HAWQ 1.2

Issue	Category	Resolved in	Description
HAWQ-1834	Build and Installer	HAWQ 1.2.0.0	The plr_install.sh script failed with the error - <i>Platform not supported</i> . This issue has been resolved.
HAWQ-1721	Query Optimizer	HAWQ 1.2.0.0	The optimizer failed to process a query with join aliases. This issue has been resolved in the optimizer.
HAWQ-1706	Query Optimizer	HAWQ 1.2.0.0	For certain queries that have inner and outer joins, the optimizer failed while exploring alternative plans leading to a crash. This issue is now fixed in the optimizer.
HAWQ-1702	Query Optimizer	HAWQ 1.2.0.0	For some queries containing built-in functions such as: pg_stat_get_backend_pid, pg_stat_get_backend_activity_start, or pg_stat_get_backend_userid; the optimizer might generate incorrect plans. This was caused by function properties being mislabeled in the catalog. This issue is now fixed in the optimizer.
HAWQ-1694	HDFS Access Layer, Query Execution	HAWQ 1.2.0.0	In a kerberized cluster with a race condition, the master released the file system credentials before the segments reached the HDFS name node. This caused the entire query to fail. This issue has been resolved.
HAWQ-1692	Query Optimizer PXF	HAWQ 1.2.0.0	PXF Predicate Push-down did not work if Orca was enabled. This issue has been resolved.

Issue	Category	Resolved in	Description
HAWQ-1618	Infrastructure	HAWQ 1.2.0.0	YARN failed to load in SingleCluster This issue has been resolved.
HAWQ-1610	Build and Installer	HAWQ 1.2.0.0	PL/R package changes. Check the name of your plr package. If it is plr-1.1.4.0-5152.x86_64.tgz,download the latest version plr-1.1.4.0-5664.x86_64.tgz for HAWQ 1.1.4.0 from Pivotal. The new package contains the file plr.sql with the necessary PL/R helper functions.
HAWQ-1527	Build and Installer	HAWQ 1.2.0.0	HAWQ and PXF version strings are now 4 digits.
HAWQ-1491	AO tables Column Store	HAWQ 1.2.0.0	After truncating a table, the HAWQ input format did not work with the truncated table. This issue has been resolved.
HAWQ-1490	AO tables Column Store	HAWQ 1.2.0.0	The function HAWQConvertUtil.bytesToDecimal was not thread safe. This is because decimalCharArray is a public static variable. This issue has been resolved.
HAWQ-1489	AO tables Column Store	HAWQ 1.2.0.0	After truncating a table, gpextract did not work. This issue has been resolved.
HAWQ-1488	AO tables Column Store	HAWQ 1.2.0.0	If the HAWQAORRecord.getBoolean function encountered a column with boolean data type, it returned the incorrect result, false. This issue has been resolved.
HAWQ-1482	PXF	HAWQ 1.2.0.0	gphdfilters created a filter in the reverse order This issue has been resolved.
HAWQ-1455	Dispatch	HAWQ 1.2.0.0	Signal re-entrant during session idle. QD crashes. This issue has been resolved.
HAWQ-1451	Query Exexecution	HAWQ 1.2.0.0	Explain analyze statistics are not correct for work files . This issue has been resolved.

Issue	Category	Resolved in	Description
HAWQ-1450	Infrastructure	HAWQ 1.2.0.0	SingleCluster hdfs tool was not working with Hadoop 2.2 This issue has been resolved.
HAWQ-1429	Default	HAWQ 1.2.0.0	Unable to start HAWQ master because recovery failed. The master failed to start during recovery mode because some files existed locally and were missing on the HDFS layer. This issue has been resolved.
HAWQ-1418	Catalog and Metadata	HAWQ 1.2.0.0	HAWQ 1.1.4.0 did not support aggregate derived functions. This issue has been resolved.
HAWQ-1379	Management Tools	HAWQ 1.2.0.0	hawq_toolkit cannot be used directly after upgrading from an old version. This is because toolkit related objects are not created in the old version. Workaround: for each existing database instance where a user wants to use hawq_toolkit, perform following steps as superuser: <ol style="list-style-type: none"> 1. create a new schema named hawq_toolkit: <code>psql -q -c "CREATE SCHEMA hawq_toolkit" \$DATABASE_NAME</code> 2. create toolkit related objects: <code>psql -q -f \$INSTALL_DIR/share/postgresql/gp_toolkit.sql \$DATABASE_NAME</code> After performing the above operations on template1, every newly created database using template1 as template database, will have hawq_toolkit automatically, meaning no need to perform the above operation.
HAWQ-1364	PXF	HAWQ 1.2.0.0	While copying data to a writable interface HDFS table, showed remote component error 500. This issue has been resolved.
HAWQ-1358 1257	DDL Object	HAWQ 1.2.0.0	Received a confusing error when creating a table that distributes by text data type. This issue has been resolved.

Issue	Category	Resolved in	Description
HAWQ-1260	Query Execution	HAWQ 1.2.0.0	<p>A certain class of uncorrelated subqueries are known to fail. The subquery should have a user defined object and a distributed table. For example:</p> <pre>SELECT * FROM t1 WHERE t1.a < (SELECT foo(t2.b) FROM t2 LIMIT 1);</pre> <p>In this example, the subquery "SELECT foo(t2.b) FROM t2 LIMIT 1" has no correlation with the outer query. The subquery also invokes the UDF "foo()" and queries a distributed table "t2". Another example:</p> <pre>SELECT array(SELECT foo(t1.a) FROM t1);</pre> <p>Such type of queries fail with the following error:</p> <pre>ERROR cache lookup failed for ...</pre> <p>This issue has been resolved.</p>
HAWQ-1184	DDL Object	HAWQ 1.2.0.0	<p>ALTER TABLE ADD COLUMN with default NULL was not supported for append-only tables.</p> <p>This syntax is now supported.</p>
HAWQ-872	DDL Object	HAWQ 1.2.0.0	<p>In certain cases, INSERT INTO SELECT from the same table might insert an incorrect number of tuples. This happens if the table is altered prior to the insert.</p> <p>This issue has been resolved.</p>

6 Known Issues in HAWQ 1.2

Issue	Category	Description
HAWQ-1920	Query Optimizer	In some cases, the system was getting stuck in recovery mode because segments continued to run plans with motion nodes during the recovery process. Such plans are invalid during recovery, and are no longer being generated.
HAWQ-1918	Catalog and Metadata	Nested functions in any language are not supported in HAWQ 1.2.
HAWQ-1900	Management Tools Documentation	HAWQ supports sending email alert notification for.: <ul style="list-style-type: none"> • All PANIC-level error conditions • All FATAL-level error conditions • ERROR-level conditions that are "internal errors" (for example, SIGSEGV) • Database system shutdown and restart • Segment failure and recovery • Standby master out-of-sync conditions • Master host manual shutdown or other software problem (in certain configurations) Greenplum Database will not be able to send an alert or notification
HAWQ-1868	DML	If a query does not have a FROM clause, and contains the random() function in the SELECT clause along with another function that returns multiple rows, then the query will generate the same random number rather than generating different random numbers.
HAWQ-1859	Build and Installer	Run <code>plr_install.sh</code> to copy the <code>pgcrypto.so</code> on the master and segments. To install <code>pgcrypto</code> functions for another database, run the following: <pre>psql -d<TARGET_DATABASE> -f \$GPHOME/share/postgresql/contrib/pgcrypto/pgcrypto.sql</pre>
HAWQ-1728	Query Optimizer	If ORCA is on, the INSERT command fails, but works fine with ORCA off.
HAWQ-1543	Upgrade	In a single node setting, <code>gpmigrator</code> tries to create temporary directories twice with the same name under <code>DATA_DIRECTORY</code> and <code>MASTER_DIRECTORY</code> , set <code>debug_level</code> to <code>gpinitssystem</code> . The second time will fail.
HAWQ-1456	Transaction	Running DROP SCHEMA and CREATE TABLE on the same table makes the table inaccessible.
HAWQ-1453	HDFS Access Layer	Executing concurrent INSERT and ALTER TABLE statements, generates the error: <pre>ERROR: read beyond eof in table "tbl_isolation" in file "hdfs://smdw:9000/hawq/gpdb20131226t190718-885441423/releng4/16385.tbl_isolation" (cdbbufferedread.c:199) (seg4 slice1 sdw2:31100 pid=316232) (cdbdisp.c:199)</pre>

Important: Pivotal Hadoop (PHD) 1.1. has a new High Availability (HA) feature. This feature has the following known issue with HAWQ: If NameNode HA is configured in a cluster, HAWQ is unable to take advantage of HA capability when the primary configure HA for HDFS. In the case of the primary NameNode failing, HAWQ will not be able to failover to the secondary NameNode. Manual re-direction to secondary NameNode is not currently supported in this release. Therefore we highly recommend that HA should be disabled when used with HAWQ.

7 Known Issues in PXF 2.x.x

Issue	Category	Description
HAWQ- 1739	PXF	PXF does not filter UTF8 encoded parameters correctly.
HAWQ- 1720	PXF	Error table has one extra error reported if the last row has an error
HAWQ- 1649	PXF	Intermittent failures when using pxf_profile.
HAWQ- 1481	PXF	PXF Filter pushdown handles badly constants values with embedded quotes .

8 Pivotal and HAWQ Interoperability

Pivotal releases a number of client tool packages on various platforms that can be used to connect to HAWQ. The following table describes the client tool package compatibility with HAWQ. Client tool packages are available at the EMC Download Center.

Table: Interoperability Matrix

Client package	Description	Operating system	Client version	HAWQ version
Connectivity	Standard PostgreSQL Database Drivers (ODBC, JDBC)	Windows 2008 RedHat 6.4 and 6.2, 64 bit	4.2.6SP	1.2.0.0
HAWQ Client	Command Line Interface	Windows 2008 RedHat 6.4 and 6.2, 64 bit	4.2.6SP	1.2.0.0
Pivotal Command Center	A web-based tool for managing and monitoring your Pivotal HD cluster. Note: Pivotal Command Center 2.0.x does not support DCA V1, DCA V2 or Greenplum Database.	Windows 2008 RedHat 6.4 and 6.2, 64 bit CentOS 6.4 and 6.2, 64 bit	2.2	1.2.0.0
PXF	Extensibility layer to provide support for external data formats such as HBase and Hive.	Windows 2008 RedHat 6.4 and 6.2, 64 bit CentOS 6.4 and 6.2, 64 bit	2.2	1.2.0.0
Pivotal HD	Pivotal Hadoop	Windows 2008 RedHat 6.4 and 6.2, 64 bit CentOS 6.4 and 6.2, 64 bit	2.0	1.2.0.0
pgcrypto	A library of cryptographic functions	Windows 2008 RedHat 6.4 and 6.2, 64 bit CentOS 6.4 and 6.2, 64 bit	1.2.0.0 1.1.3.0-4609	1.2.0.0 1.1.3.x and 1.1.4.x

Client package	Description	Operating system	Client version	HAWQ version
PL/R	Ability to create and invoke user defined functions in R	Windows 2008 RedHat 6.4 and 6.2, 64 bit CentOS 6.4 and 6.2, 64 bit	1.2.0.0 1.1.4.0-5664	1.2.0.0 1.1.4.x
PL/Java	Ability to create and invoke user defined functions in Java	Windows 2008 RedHat 6.4 and 6.2, 64 bit CentOS 6.4 and 6.2, 64 bit	1.2.0.0	1.2.0.0

9 HAWQ 1.2 and Pivotal HD Documentation

The following HAWQ and related documentation is available in PDF format on our website at www.gopivotal.com.

HTML versions of our documentation is available here: docs.gopivotal.com/pivotalhd/

Additionally, you can still access product documentation from EMC's [Support Zone](#).

Table: HAWQ Documentation

Title	Revision
Pivotal HAWQ 1.2 Release Notes (This document)	A01
Pivotal HAWQ 1.2 Installation and Upgrade Guide	A01
Pivotal HAWQ 1.2 Administrator Guide	A01
Pivotal HD Enterprise 2.0 Installation and Administrator Guide	A01
Pivotal HD 2.0 Stack and Tool Reference Guide	A01
Pivotal Command Center 2.2 User Guide	A01
Pivotal Extension Framework 2.2 Installation and User Guide	A01

10 Upgrading HAWQ

For more information about upgrading HAWQ and other components, see the *HAWQ Installation and Upgrade Guide*.